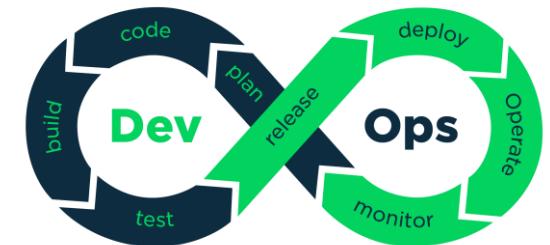


DevOps - Let the Journey Begin



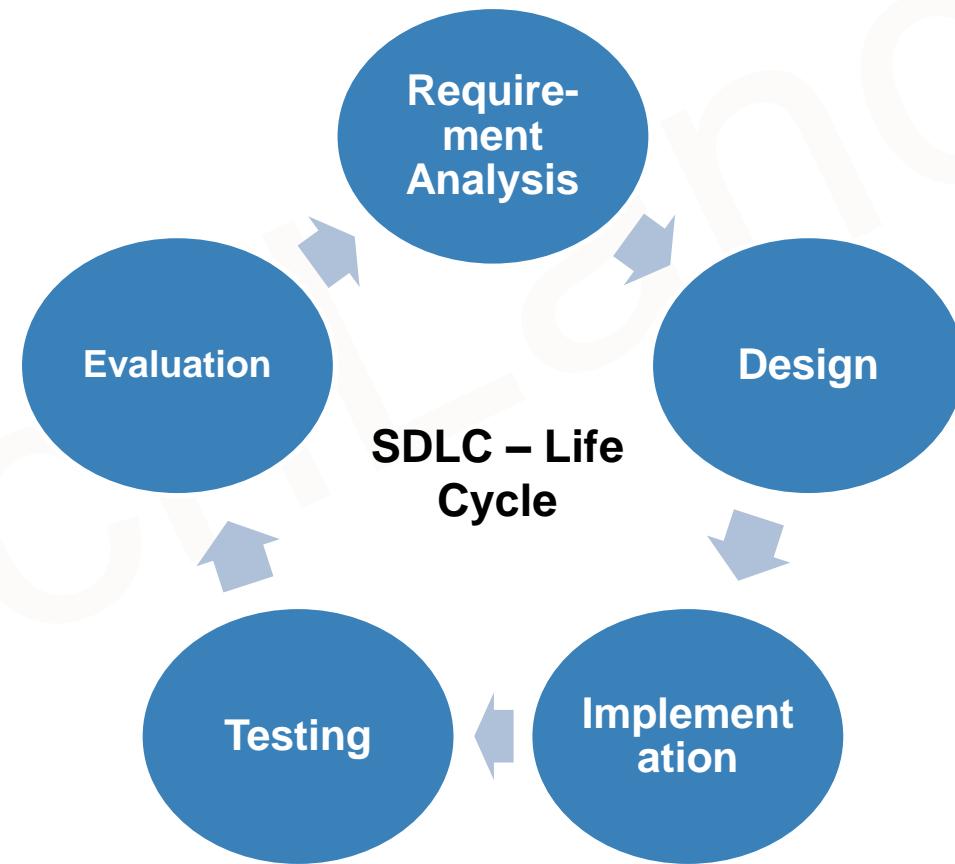
Raman Khanna

DevOps

What is DevOps?

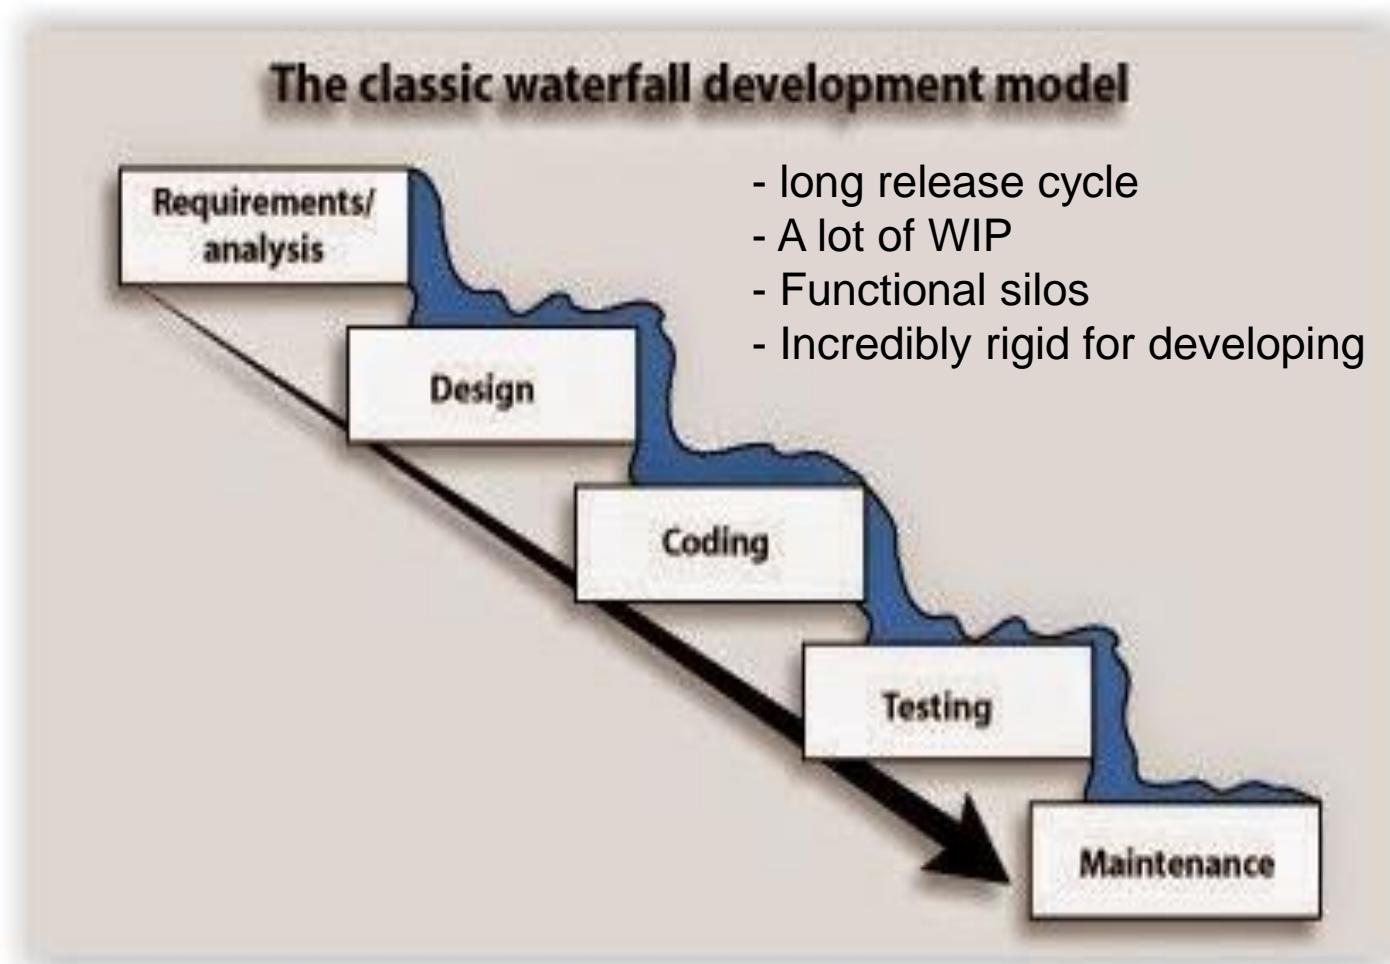
SDLC Model

- A systems development life cycle is composed of **several clearly defined and distinct work phases** which are used by systems engineers and systems developers to plan for, design, build, test, and deliver information systems



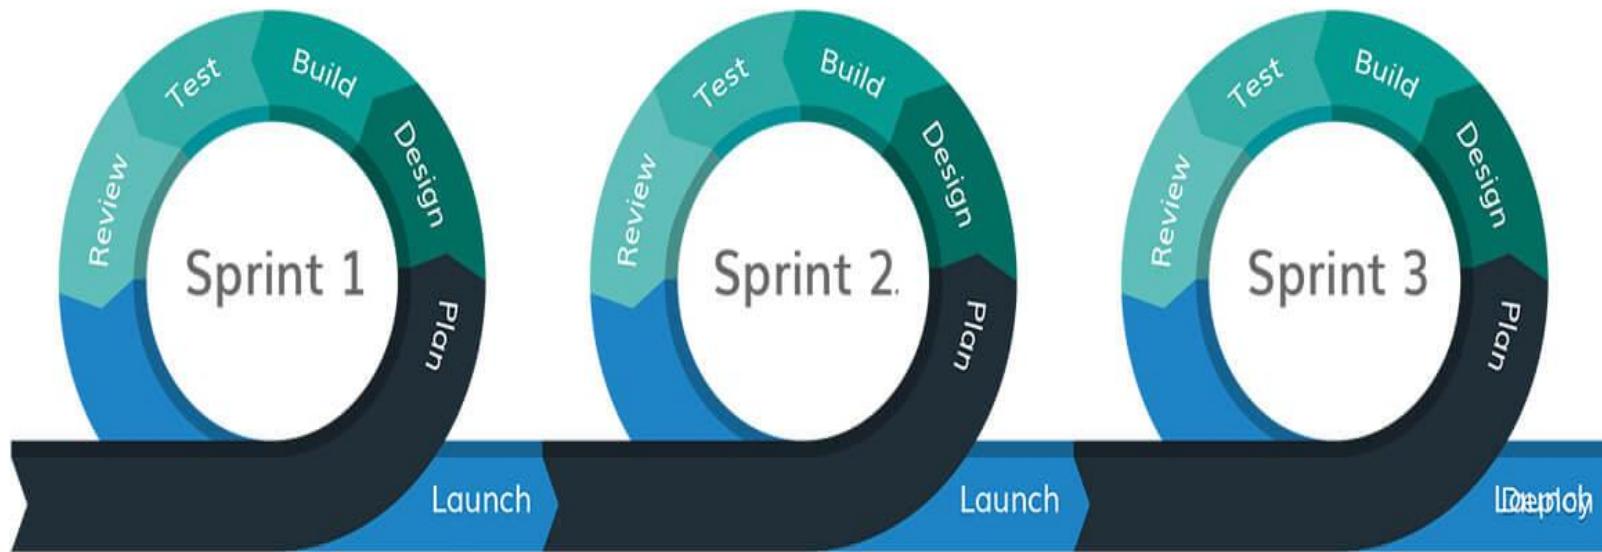
Waterfall Model

1. Determine the Requirements
2. Complete the design
3. Do the coding and testing
(unit tests)
4. Perform other tests
(functional tests, non-functional tests, Performance testing, bug fixes etc.)
5. At last deploy and maintain



Agile

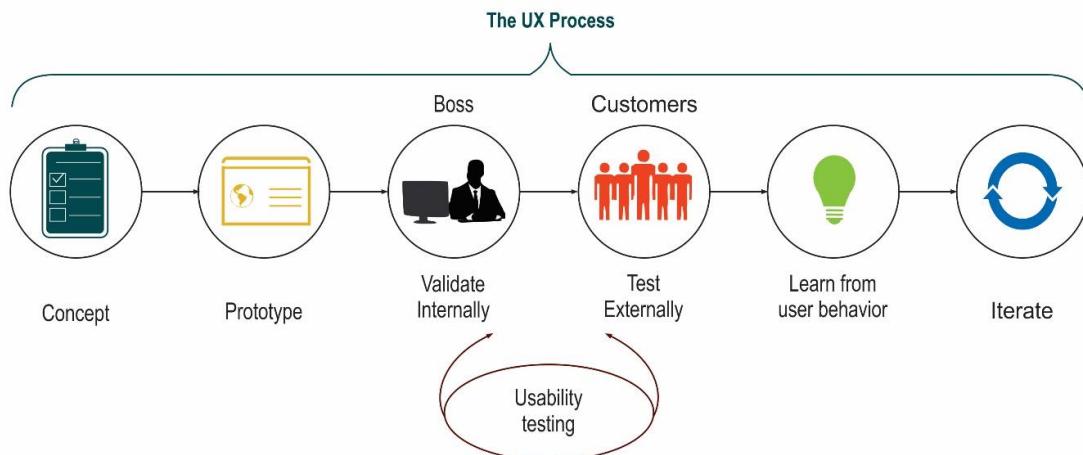
Agile Methodology



- Shorter release cycle
- Small batch sizes (MVP)
- Cross-functional teams
- Incredibly agile

Lean Development

Lean Development (LD)



Not like this...



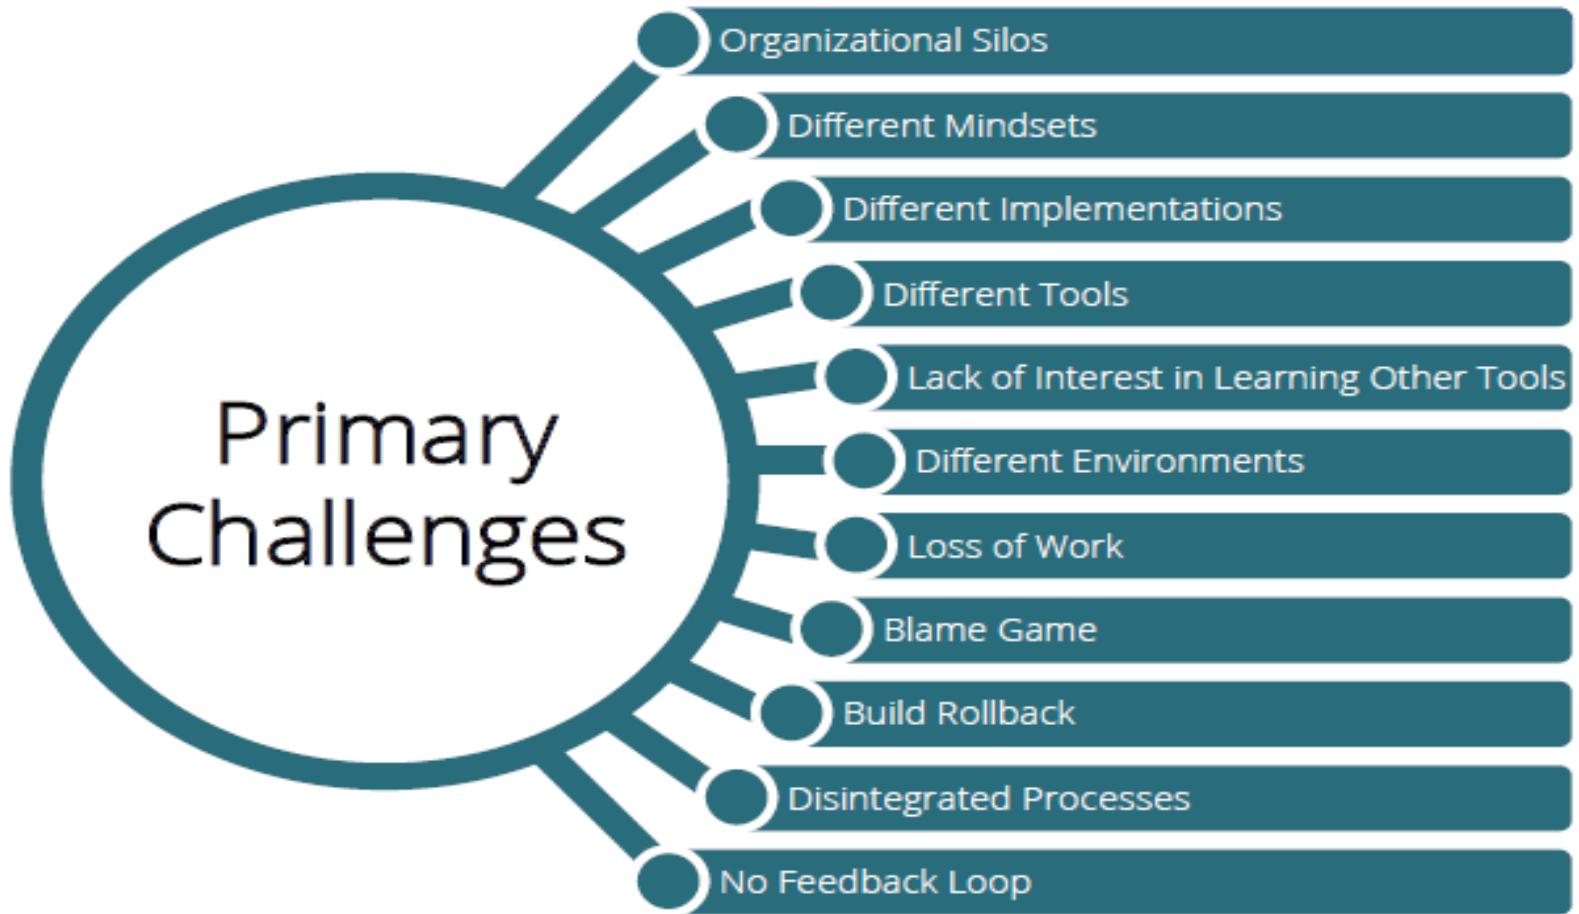
...instead like this!



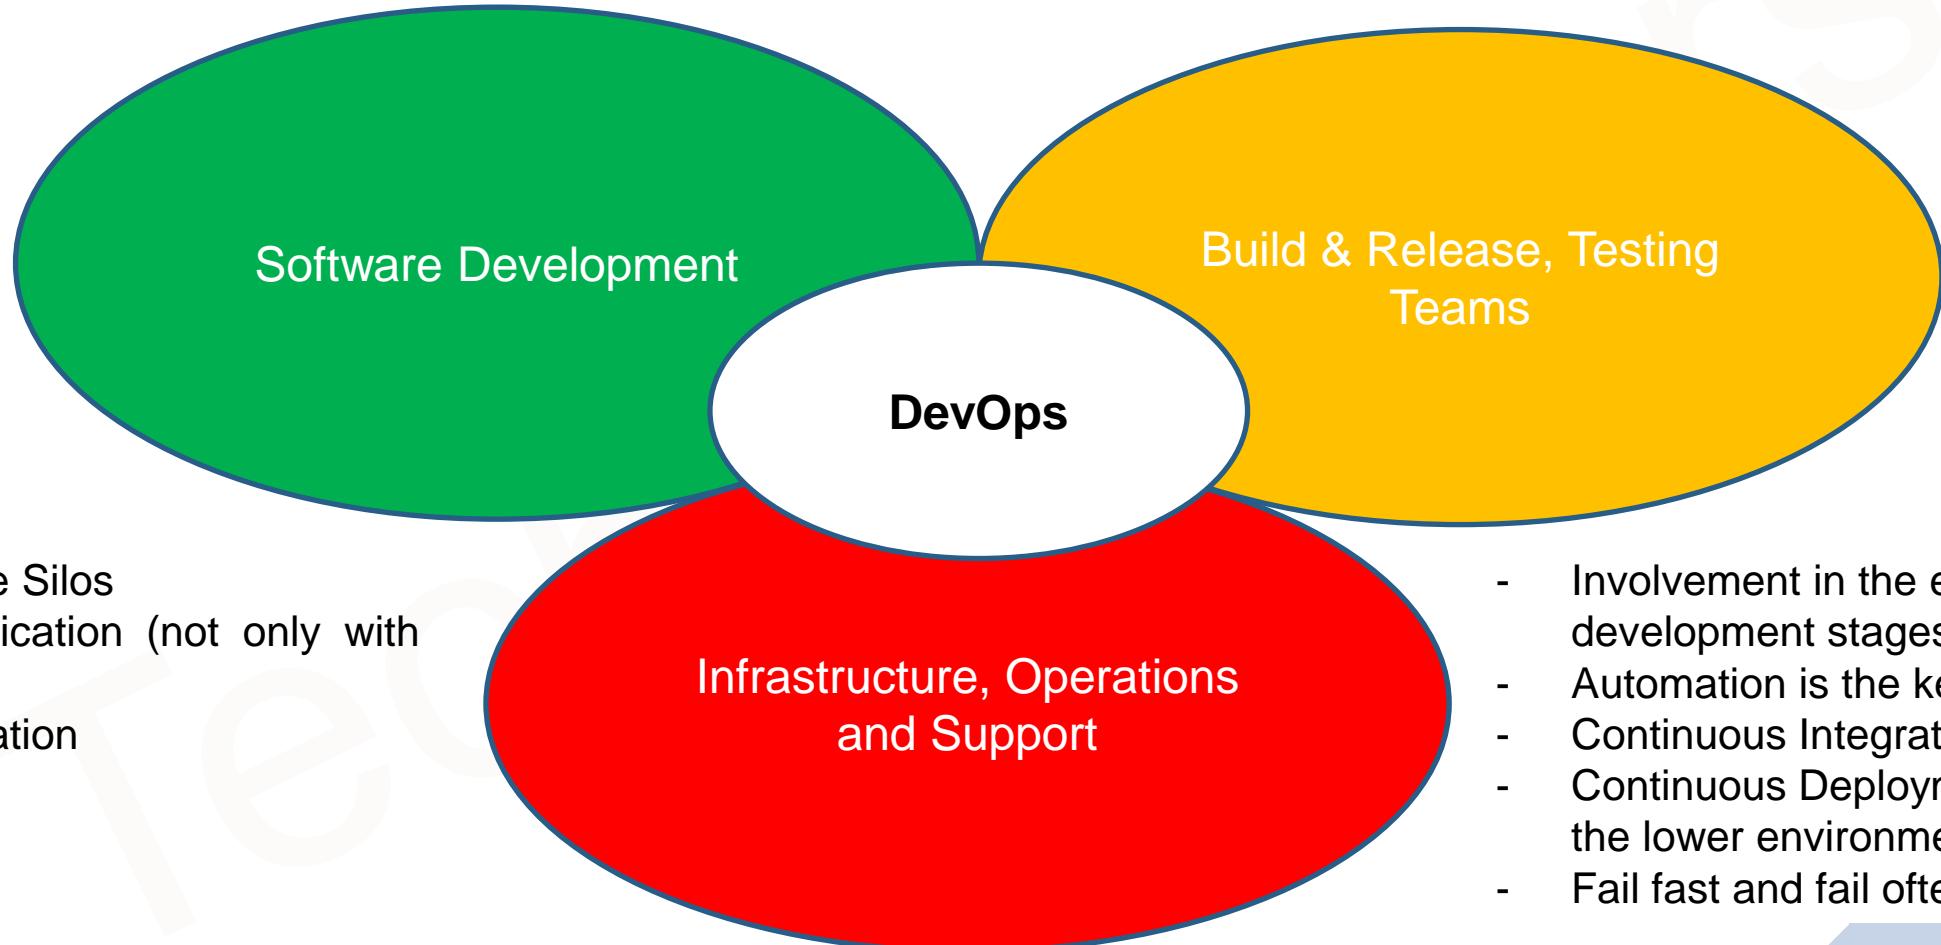
- Suddenly ops was the bottleneck (more release less people), again WIP is more!

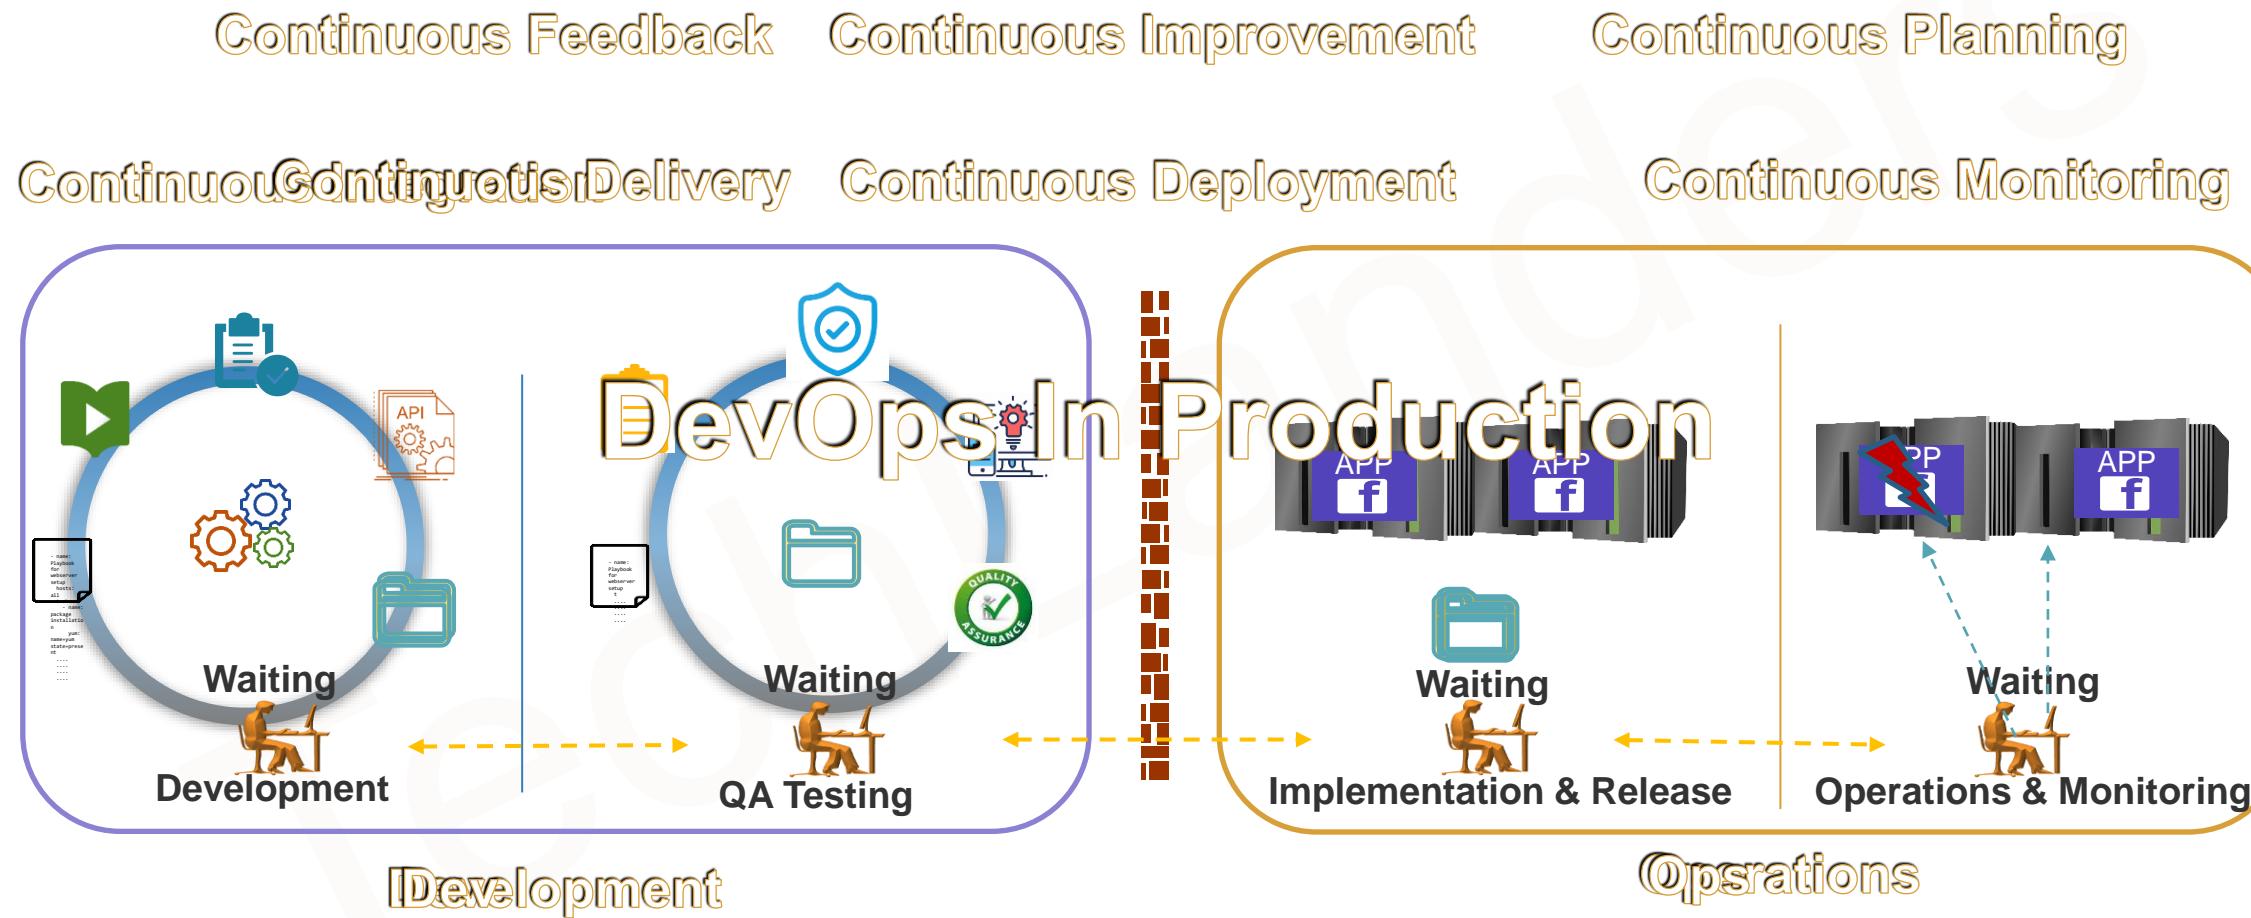
Challenges

Some of the challenges with the traditional teams of Development and Operations are:



DevOps





DevOps Essence

Efficiency - Faster time to market

Predictability - Lower failure rate of new releases

Reproducibility – Version everything

Maintainability - Faster time to recovery in the event of a new release crashing or otherwise disabling the current system

DevOps Core Principles

1. Customer-Centric Action



2. Create with the End in Mind



3. End-to-End Responsibility



4. Cross-Functional Autonomous Teams



5. Continuous Improvement



6. Automate Everything You Can



How to Build DevOps Organization Culture

Retention is as important as recruitment

Establish Cross-functional team structure

Small teams are better

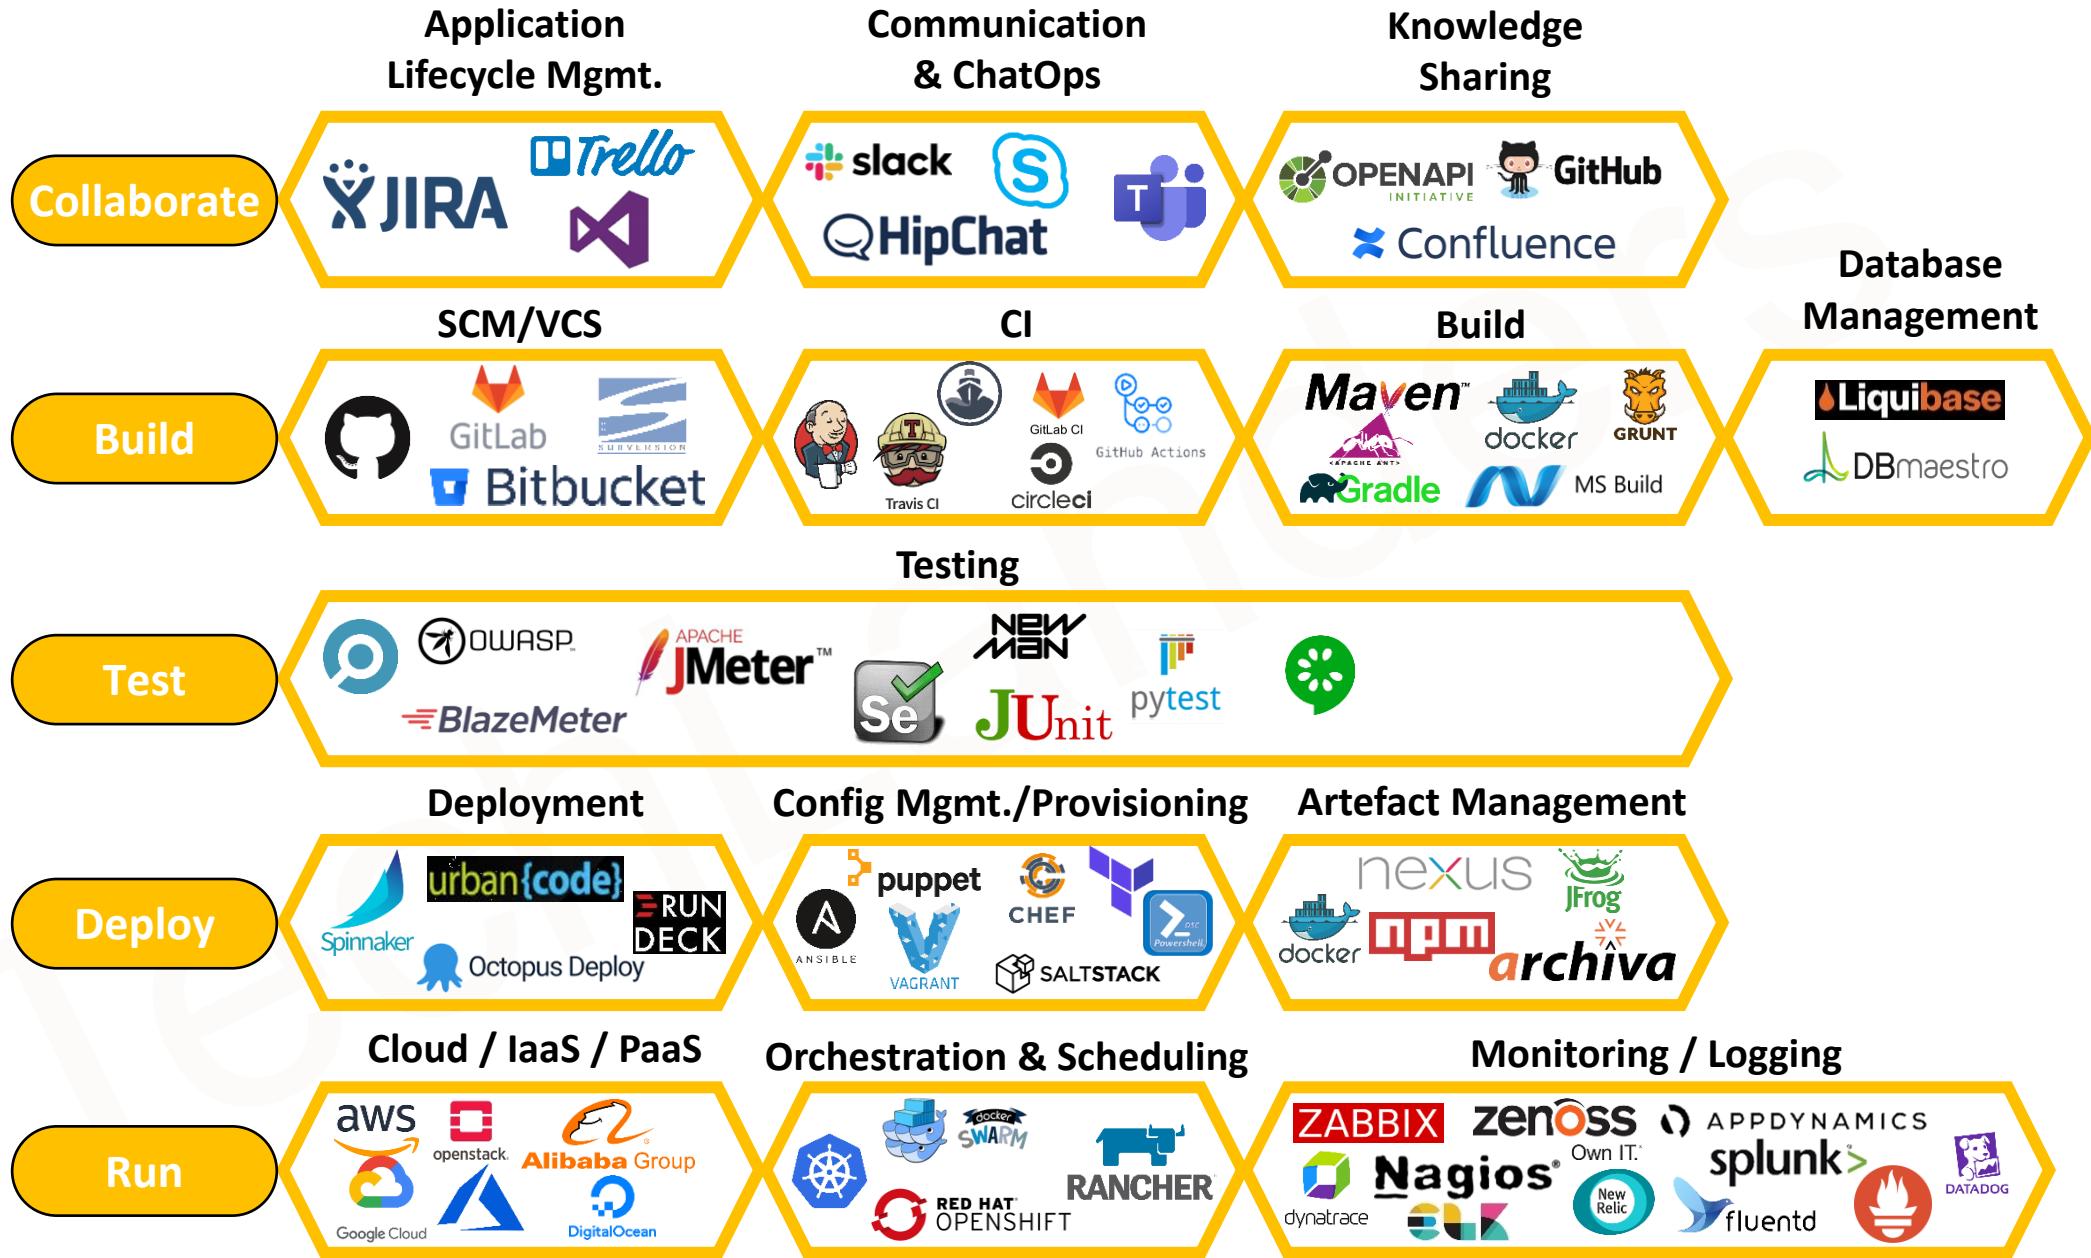
Cool tools can attract and retain

Give autonomy

Automate with existing staffs and give them a chance to learn

Take out few resources from each team, build a new virtual team for automation.

DevOps Toolsets



Version Control Systems with GIT

What is Version Control System

As name states Version Control System is the “**Management of changes to anything**”.

Version Control is way of storing files in central location accessible to all team members and enabling them to keep track of changes being done in the source code by whom, when & why. It also help teams to recover from some inevitable circumstances.

Think about traditional versioning of file with names – Login001.java, Login002.java, Login_final.java.

Its not just for code, it also helps in

- Backups & Restoration

- Synchronization

- Reverts

- Track Changes

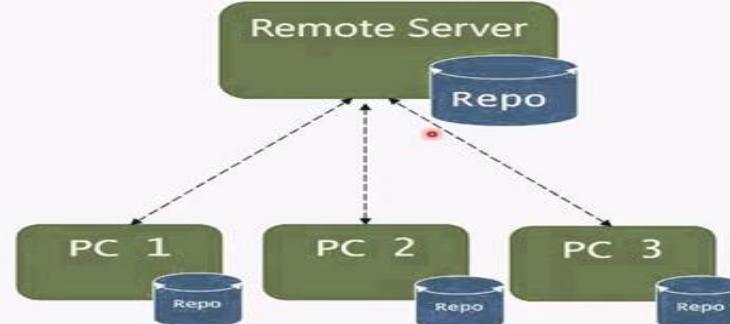
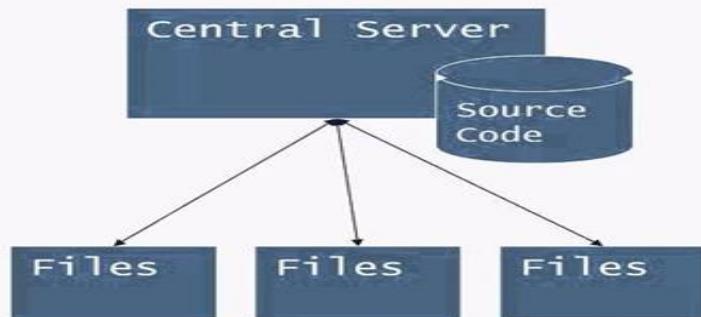
- Most importantly in Parallel Development

Types of VCS

Majorly VCS is divided into two parts:

- Centralized Version Control System – CVS, Subversion, Visual Source Safe
- Distributed Version Control System – Mercurial, Bitkeeper, Git

Centralized vs. Distributed



Git History

Linus uses BitKeeper to manage Linux code

Ran into BitKeeper licensing issue

Liked functionality

Looked at CVS as how not to do things

April 5, 2005 - Linus sends out email showing first version

June 15, 2005 - Git used for Linux version control

Why Git?

- **Branching:** gives developers a great flexibility to work on a replica of master branch.
- **Distributed Architecture:** The main advantage of DVCS is “**no requirement of network connections to central repository**” while development of a product.
- **Open-Source:** Free to use.
- **Integration with CI:** Gives faster product life cycle with even faster minor changes.

Git Installation

- yum install autoconf libcurl-devel expat-devel gcc gettext-devel kernel-headers openssl-devel perl-devel zlib-devel -y
- Visit git release page - <https://github.com/git/git/releases> and pick desired version.
- curl -O -L <https://github.com/git/git/archive/v2.14.0.tar.gz>
- tar -zxvf v2.14.0.tar.gz
- cd git-v2.14.0
- make clean
- make configure
- ./configure --prefix=/usr/local
- make
- make install
- ln -s /usr/local/bin/git /usr/bin/git

Git Commands

- git -version
- [root@techlanders ~]# git config --global user.email "Gagandeep.singh@techlanders.com"
- [root@techlanders ~]# git config --global user.name "Gagandeep Singh"
- [root@techlanders ~]# git config --global -l
- user.name=Gagandeep Singh
- user.email=Gagandeep.singh@techlanders.com
- [root@techlanders ~]#
- //Initializing a repo
- [root@master git]# mkdir /Repo1
- [root@master git]# cd /Repo1/
- [root@master Repo1]# git init
- Initialized empty Git repository in /Repo1/.git/
- [root@master Repo1]#

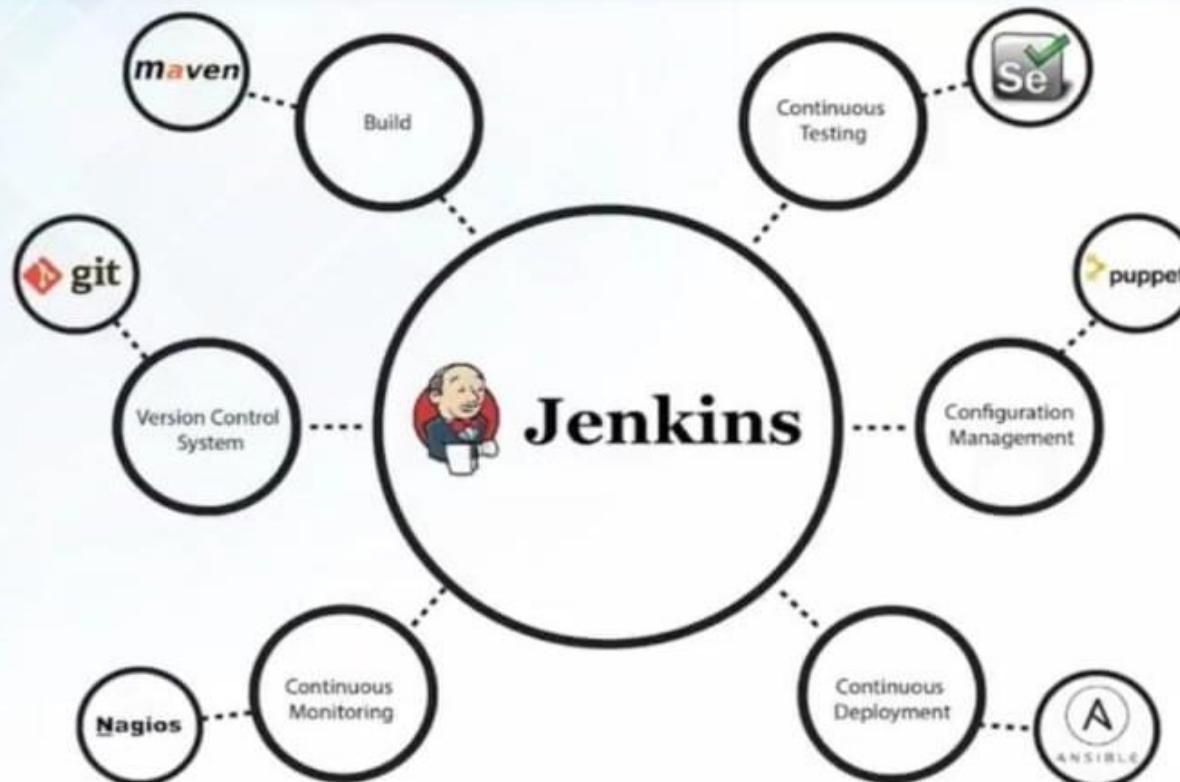
Github/Bitbucket

- **What is Bitbucket /GitHub/Gitlab?**
- Bitbucket is a Git solution for professional teams. In simple layman language its a UI for Git, offered by Atlassian, similarly we have different available UI solutions from Github (most famous) and Gitlab.
- GitHub is a code hosting platform for version control and collaboration. It lets you and others work together on projects from anywhere.
- **Host in the cloud:** free for small teams (till 5 users) and paid for larger teams.
- **Host on Your server:** One-Time pay for most solutions.
- Visit "<https://bitbucket.org/>" and click "Get Started" to sign up for free account.
- Visit "<https://github.com/>" for Github details

Continuous Integration with Jenkins

What is Jenkins?

Jenkins is an open source automation tool written in Java with plugins built for Continuous Integration purpose. Plugins allows integration of various DevOps stages.



What is Continuous Integration?

Before Continuous Integration

The entire source code was built and then tested.

Developers have to wait for test results

No Feedback

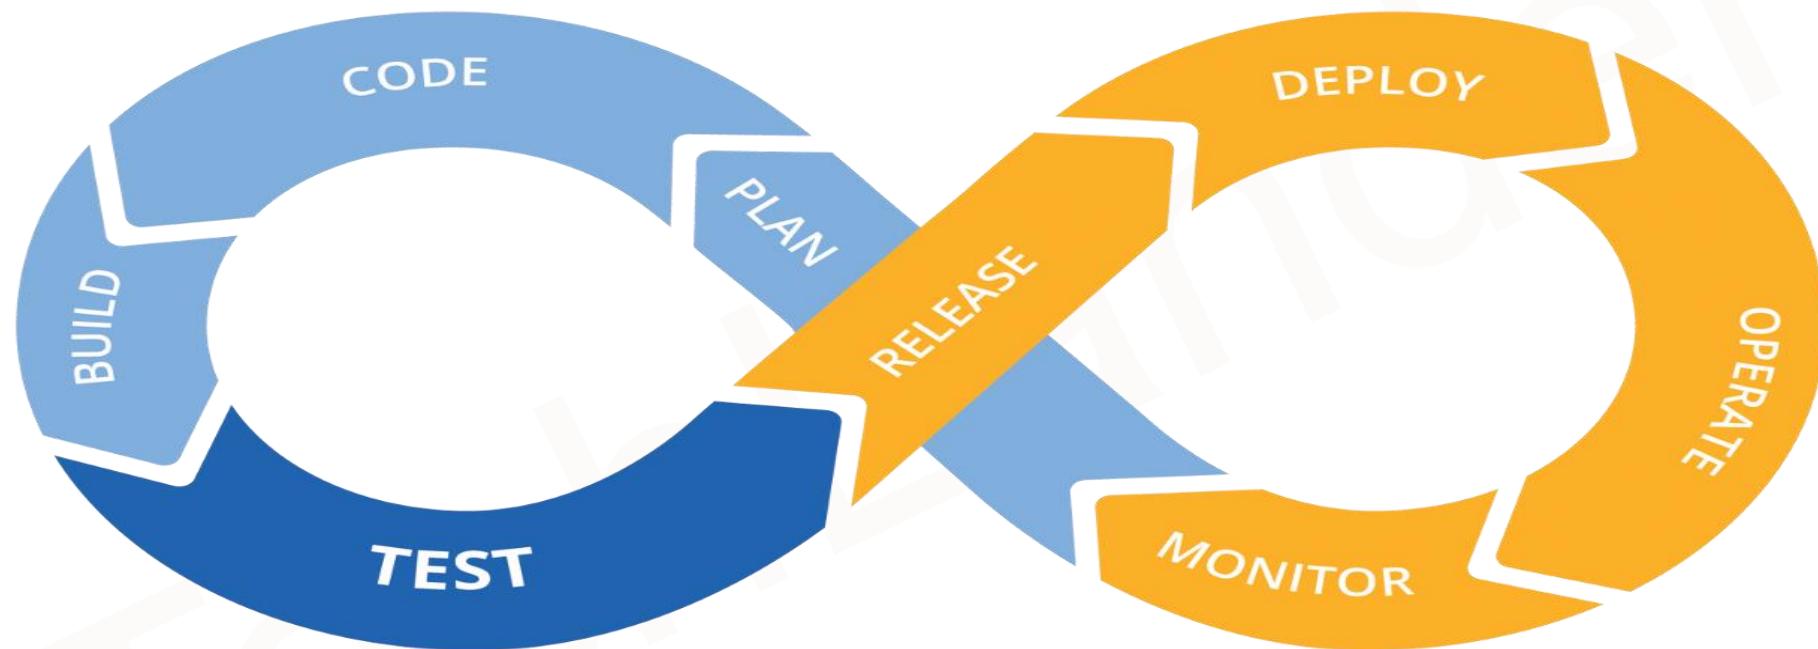
After Continuous Integration

Every commit made in the source code is built and tested.

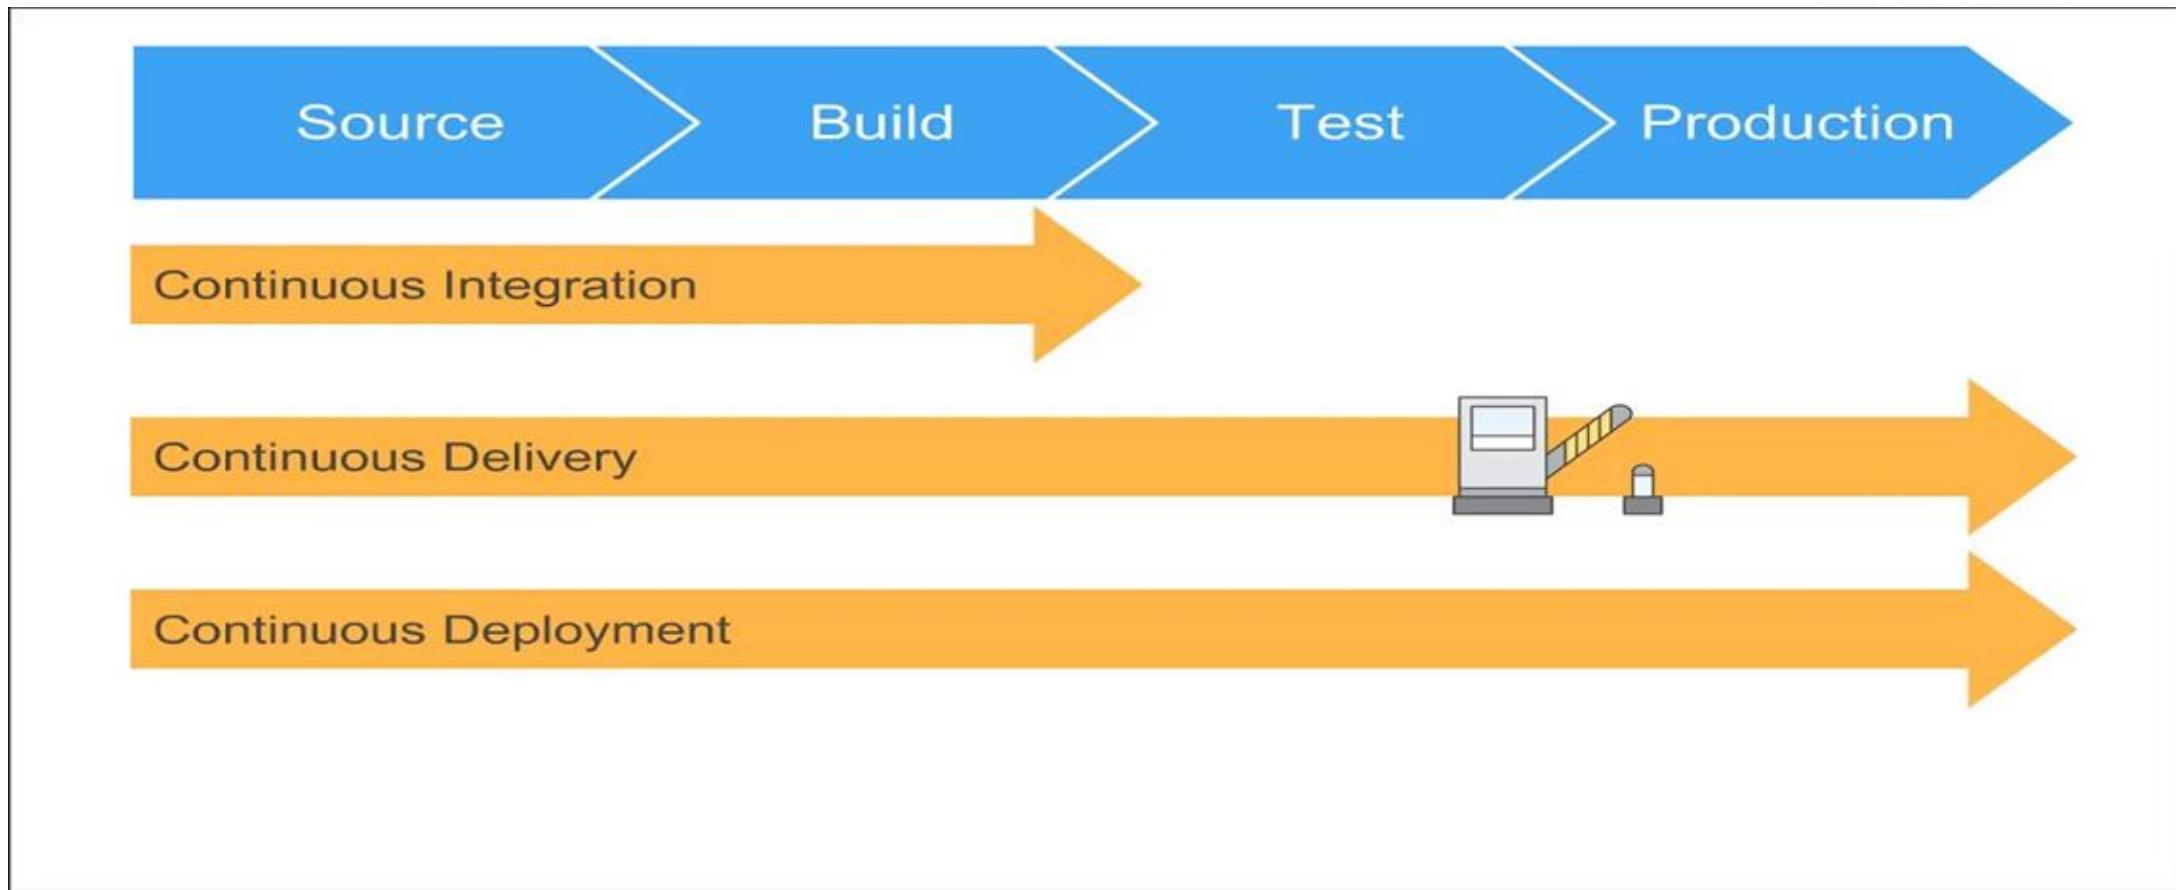
Developers know the test result of every commit made in the source code on the run

Feedback is present

CI/CD pipeline overview



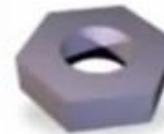
CI/CD pipeline overview



CI/CD pipeline overview



Jenkins



Buildbot



Travis CI

Atlassian
 Bamboo

What is Continuous Integration

- In Simple term Continuous Integration is a term which means “checking the compatibility of a change you have made with the remaining code base or modules of an application or to check the impact of that change on the application functionality”.
- Continuous Integration is not helping us in fixing the bugs but it helps us to Identify those bugs in the early phases of a development lifecycle and with much faster rate.
- Continuous Integration is not the sole responsibility of the developers but every individual working in the team is responsible for the Healthy Continuous Integration.

Why Continuous Integration

- CI helps in reduction of efforts which will increase exponentially with the increase in:
 - ▷ Number of components
 - ▷ Number of Bugs
 - ▷ Parallel development
 - ▷ Time from last Integration

Benefits

Project Management

- ▷ Detect system development problems in earlier stages of development
- ▷ Reduce risks of increase in cost or budget
- ▷ Reduce the risk of missing delivery timelines

Code Quality

- ▷ Measurable and visible code quality
- ▷ Continuous automatic regression and unit test
- ▷ Visible code coverage

Automation

- ▷ Automated pipelines
- ▷ Early time to reach market

Monitoring

- ▷ Complete Visibility to projects
- ▷ Time to react to changes

And much more ..

Jenkins History

- Jenkins was originally developed as the Hudson project. Hudson's creation started in summer of 2004 at Sun Microsystems. Kohsuke Kawaguchi, the current CTO of CloudBees, "A Company providing Enterprise Level Support for Jenkins, when he was working with Sun Micro Systems"
- Certain clashes between Sun Micro Systems and Hudson Community for the Management of the Project



- Today Jenkins is the most used Continuous Integration tool in the market, which provide variety of functions to run with the help of large Jenkins Community.

Brief about Jenkins

- An Open Source CI Server with MIT (Massachusetts Institute of Technology) Licensing
Large Community Base
- More than 300,000 installations on record
- More than 1300 plug-in support which make it compatible with almost every delivery tool available in market
- Annual meet-up with the Jenkins World
- CloudBees is the biggest promoter and Contributor of Jenkins Community and also provide Enterprise Support for Jenkins
- *Jenkins is a self-contained, open source automation server which can be used to automate all sorts of tasks related to building, testing, and deploying software.*

Create our first project

- Install Jenkins
- Run Jenkins as root user by modifying /etc/sysconfig/Jenkins file
- Create your first project in Jenkins
- Understand plugins and global configuration
- Install git and integrate github with Jenkins
- Install Docker and integrate docker with Jenkins
- Create a pipeline to build a docker image, run containers from the image and all code to be pulled for github repo

SonarQube – Software testing - SAST

What is Software Testing?

What is Software Testing?

Software Testing is a part of Software Development Lifecycle, which is aimed to ensure that code to be deployed is of high quality and standards, with no bugs, logical errors and issues.

It is always better to spend a little in early stages on testing than spending a lot at a later stage.

Code should be clean and reusable and have security in place.

SAST vs DAST

Sr No	Static Application Security Testing (SAST)	Dynamic Application Security Testing (DAST)
1	White Box Security Testing	Black Box Security Testing
2	Source Code Is Required	A Runtime Application is required
3	Vulnerabilities Found Earlier in SDLC	Vulnerabilities Found Later in SDLC
4	Less Expensive to Fix	More Expensive to Fix
5	Unable to Identify Timing- and Environment-Related Issues	Can Identify Run-Time and Environment-Related Issues
6	Generally Supports all Kinds of Software	Typically only scans Web applications and Services

Few Other benefits of SAST

Detects Overcomplexity in the code

Smells the code for future failures/issues

Enforces Best Coding Practices

Project Specific Rules can be created

Avoid Technical Debts

SonarQube

SonarQube (formerly Sonar) is an open-source platform developed by SonarSource for continuous inspection of code quality to perform automatic reviews with static analysis of code to detect bugs, code smells, and security vulnerabilities on multiple programming languages.

It offers reports on duplicated code, coding standards, unit tests, code coverage, code complexity, comments, bugs, and security vulnerabilities.

27+ Languages supported.

SonarQube can record metrics history and provides evolution graphs.

Fully automated analysis and integration with Maven, Ant, Gradle, MSBuild and continuous integration tools (Bamboo, Jenkins, Azure DevOps, Hudson, etc.).

Open to all tools via WebAPI and Webhooks.

SonarQube

SonarQube is available for free under the GNU Lesser General Public License. An enterprise version for paid licensing also exists, as well as a data center edition that supports high availability.

SonarQube integrates with Eclipse, Visual Studio, and IntelliJ IDEA development environments through the SonarLint plug-ins

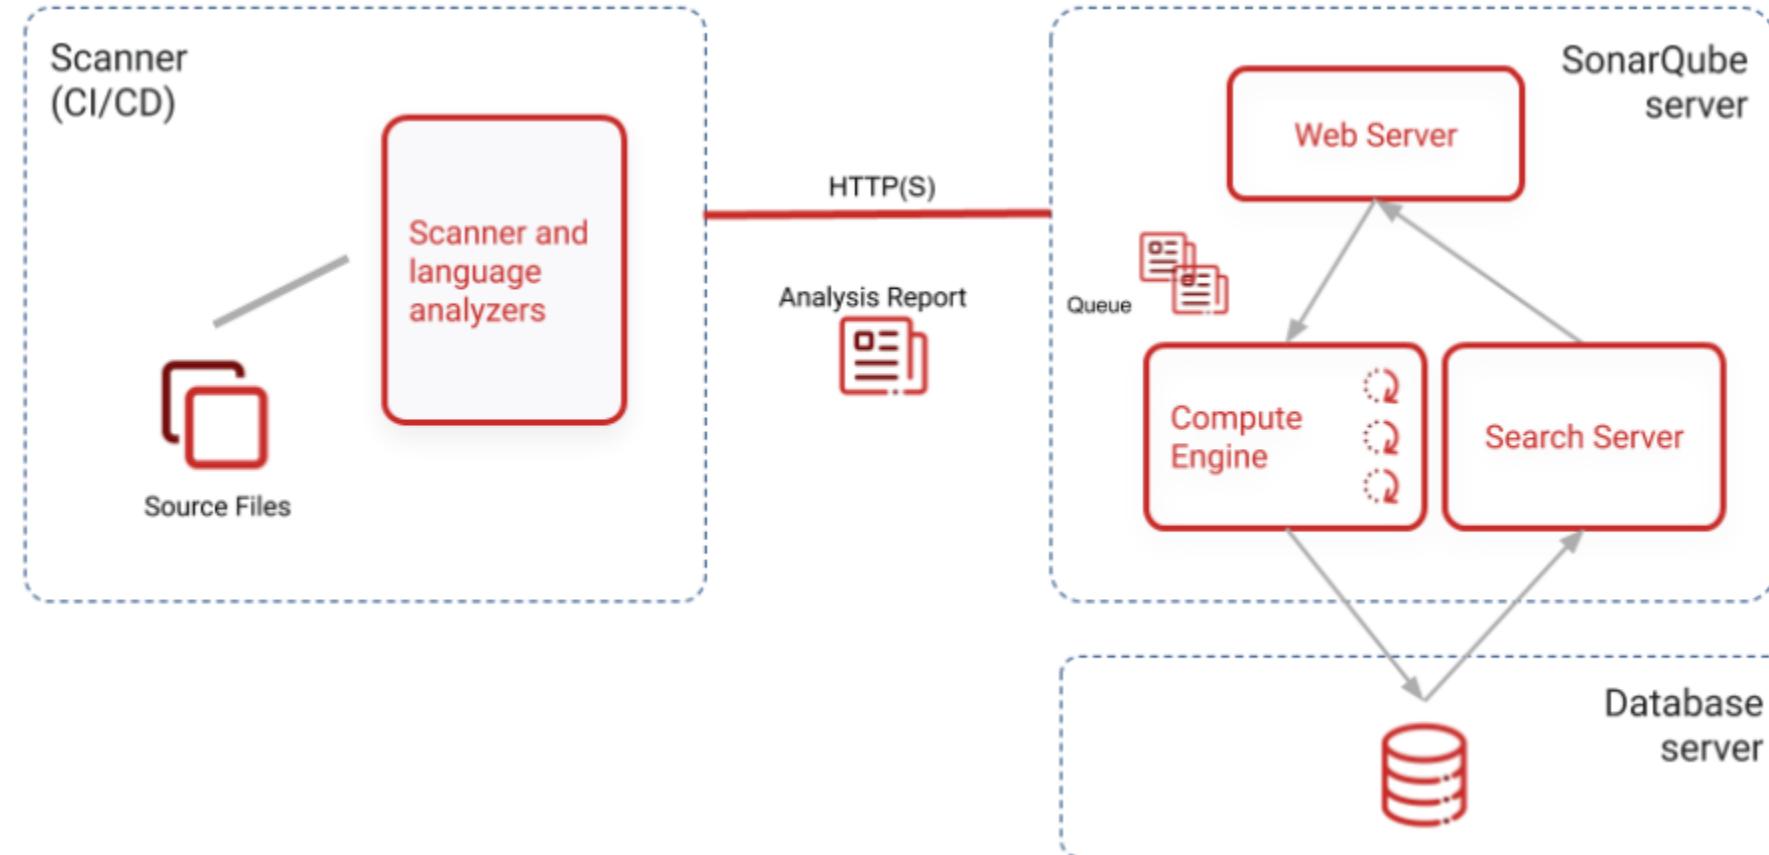
Analyzes all branches of a VCS

Discover Memory Leaks

Quality gates and Quality profiles for customized requirements

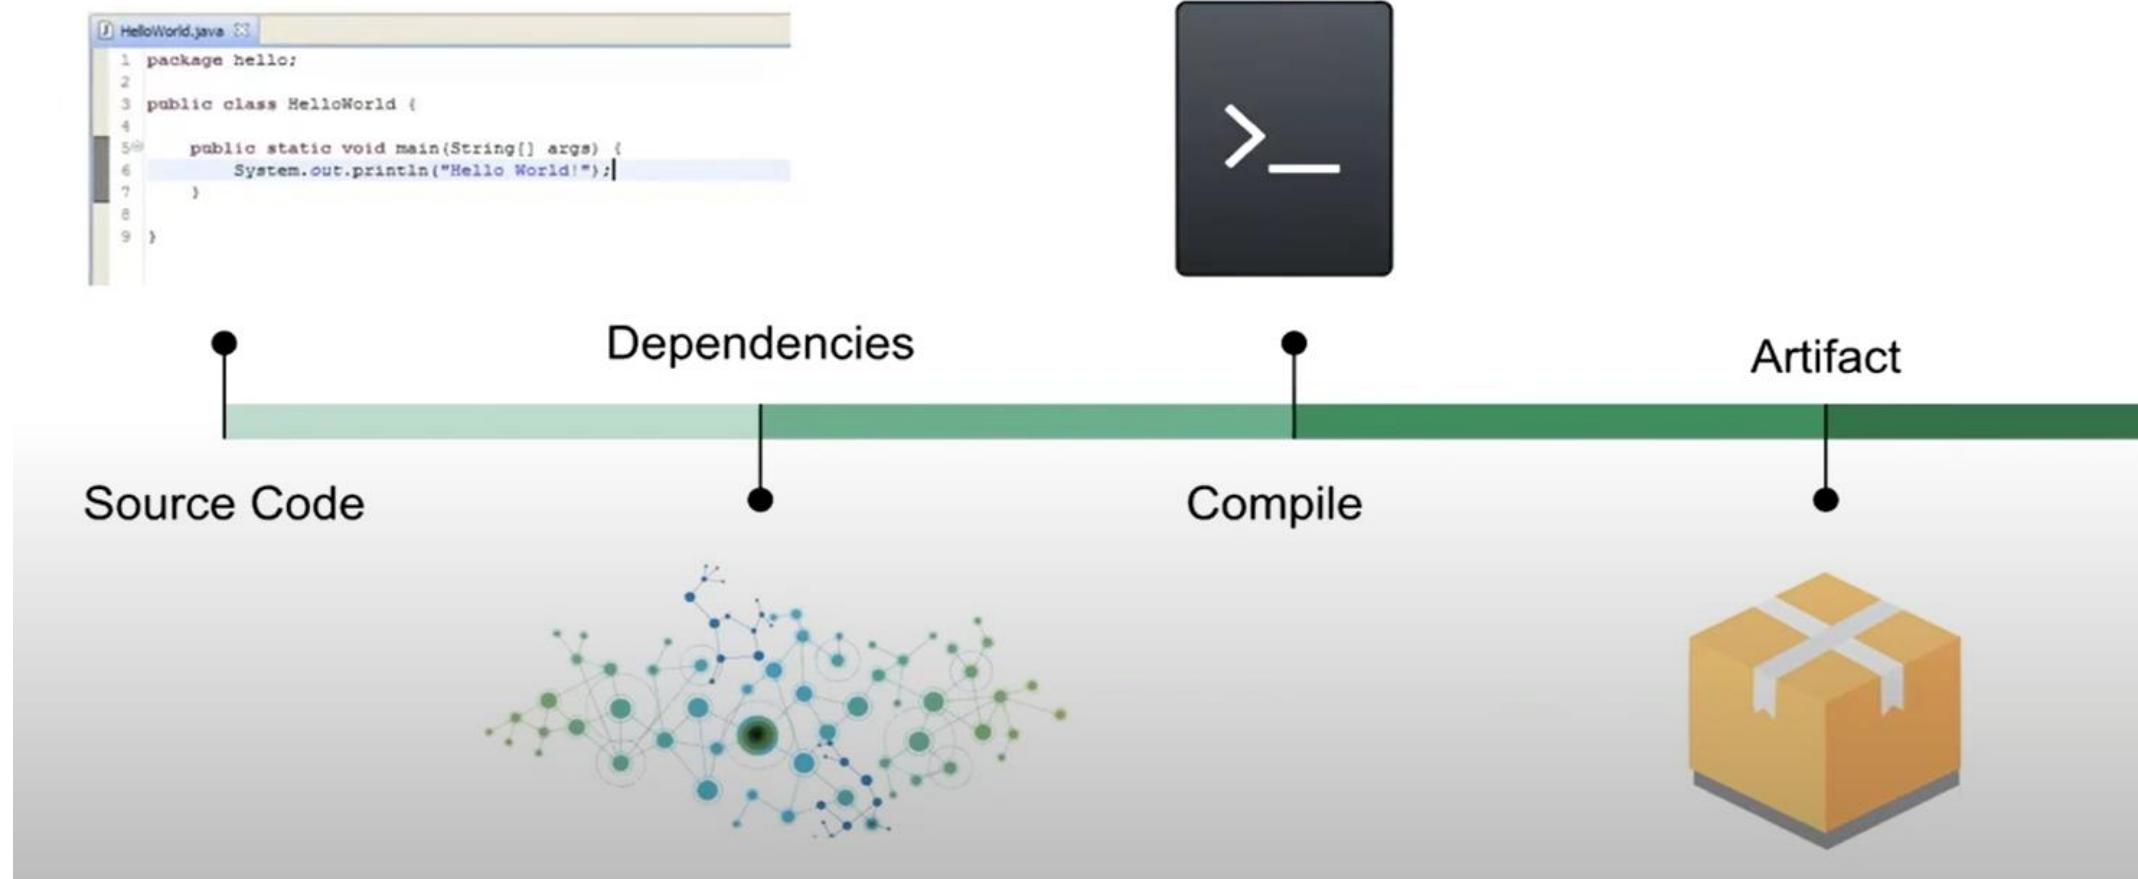
Good Visualization.

SonarQube Components



Nexus -repo manager| A Binary Artifacts manager

Code to Release



Artifact Management

To Version control your artifacts

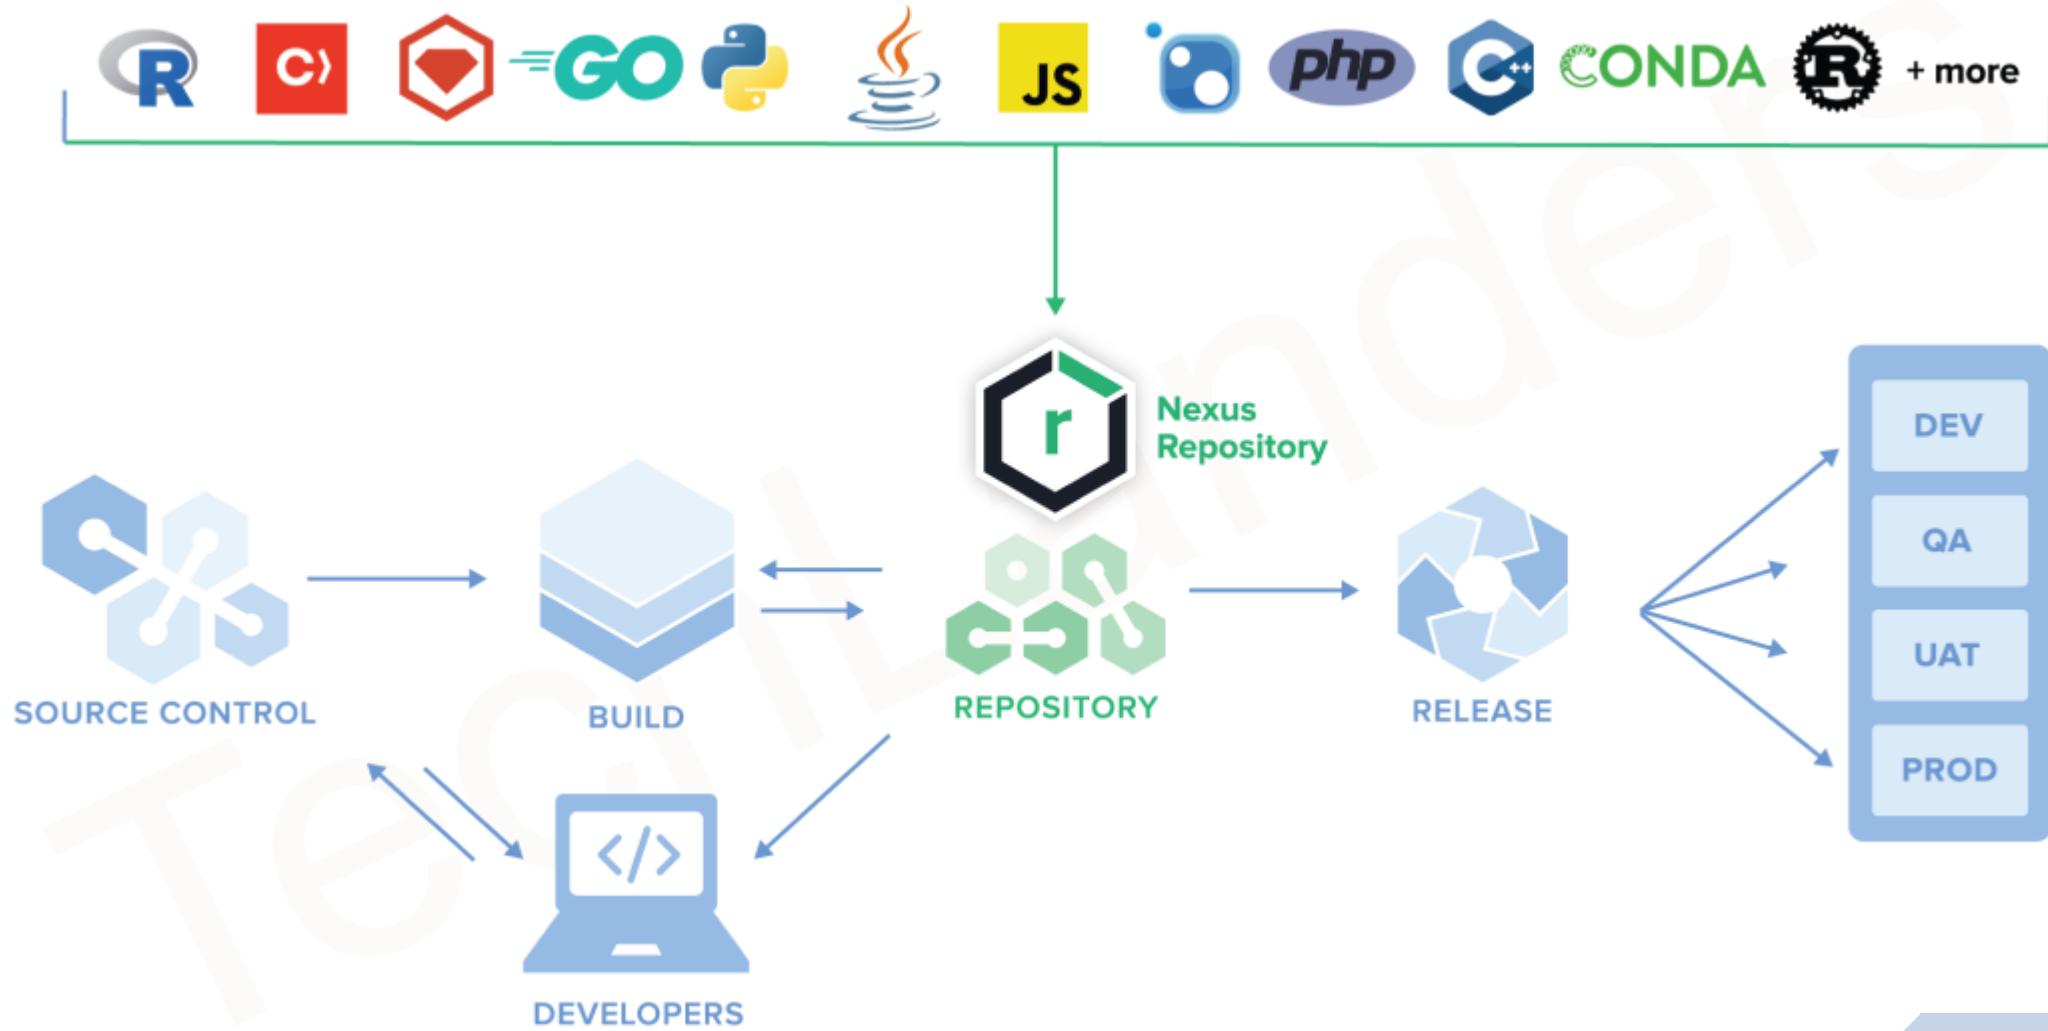
Manage them at remote Central Location

Manage Dependencies and Artifacts together

Binary Repository Manager: A collection of binary software artifacts and metadata, stored in a way which can be used by clients, package managers and CI Servers, to retrieve/store binaries during the development and build process.

Why not VCS?

Source Code Managers (VCS)	Binary/Artifacts Repository Management
Text	Binary
Diffable	Not Diffable
Versioned By content	Versioned by Name
Mutable	Immutable



Containerization

Raman

Agenda

- Introduction
- Docker Components
- Classroom Environment
- Containers
- Docker – Images
- Docker - Building Images
- Deep Dive – Images
- Deep Dive – Containers
- MicroServices Example
- Container Network Model
- Docker Volumes

Session: 1

Introduction

Before proceeding with the new terms and technologies lets have a look on the history/traditional approaches used for the application since ages.

What is Container?

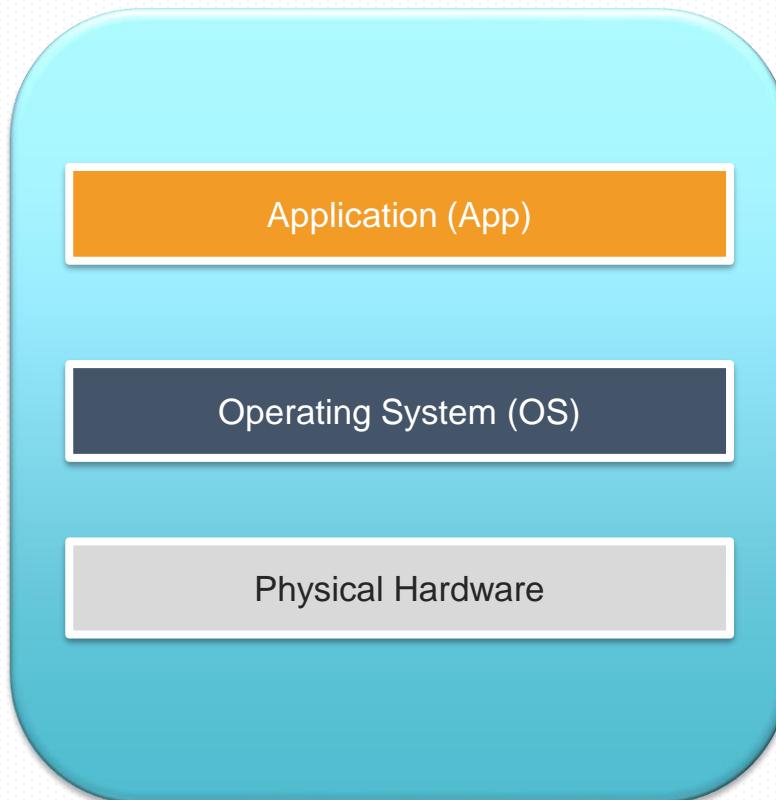
A History Lesson

- The problem got the solution by a technology called “Hypervisor-Based Virtualization”.
- One physical server can contain multiple running applications.
- Each application need a VM to run the application binaries.

A History Lesson

- In the traditional ages of development and deployment of application

Unused Resources
Difficult Migrations
Vendor dependency



Slow Deployment time
Difficult to Scale
Huge cost in Infrastructure

Virtualization

OverHead

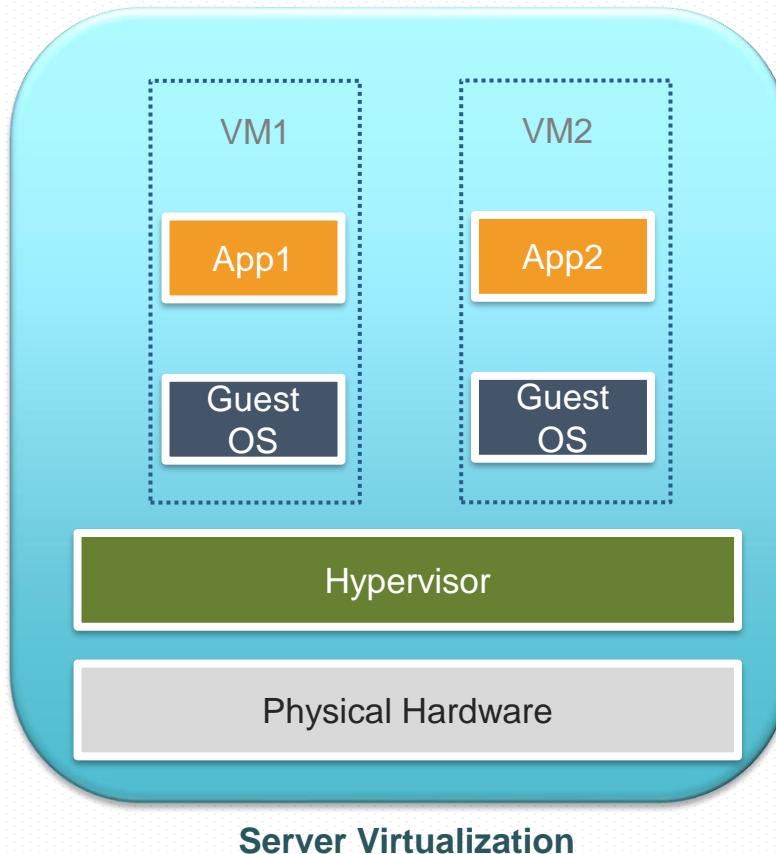
Portability issues

Boot-time still in Minutes

Scaling issue in Hybrid Env

Migrations still failing

Costly Solution



Better Resource Pooling

Easier to Scale

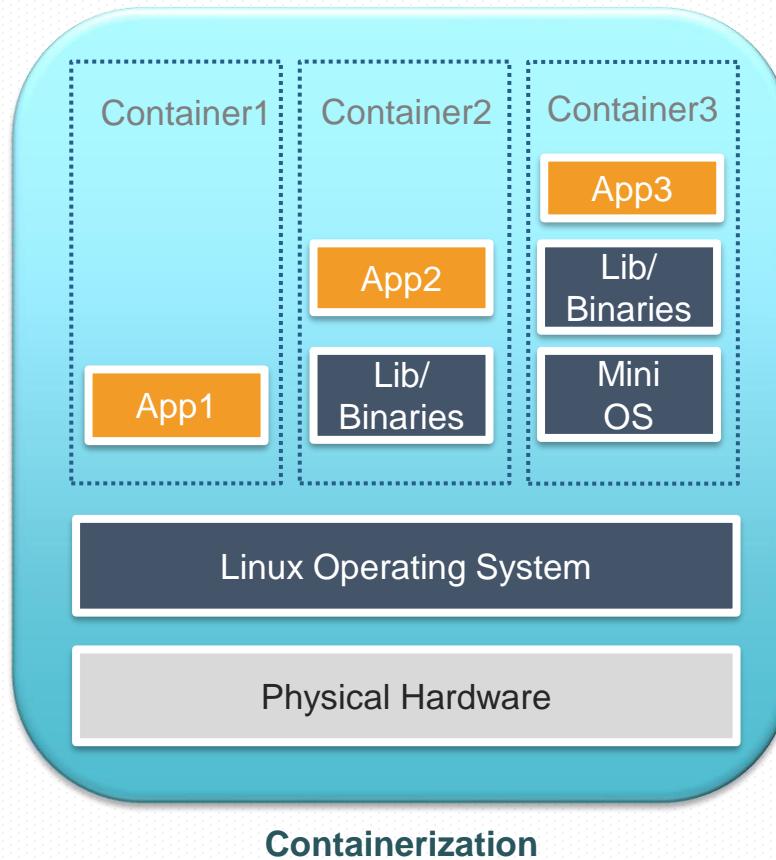
Flexibility & Easy Migration

Faster Deployments

Faster Boot time

Containerization

- Less OverHead
- Highly Portable
- Scaling in Hybrid Env
- High Migrations success ratio
- Cost Effective Solution

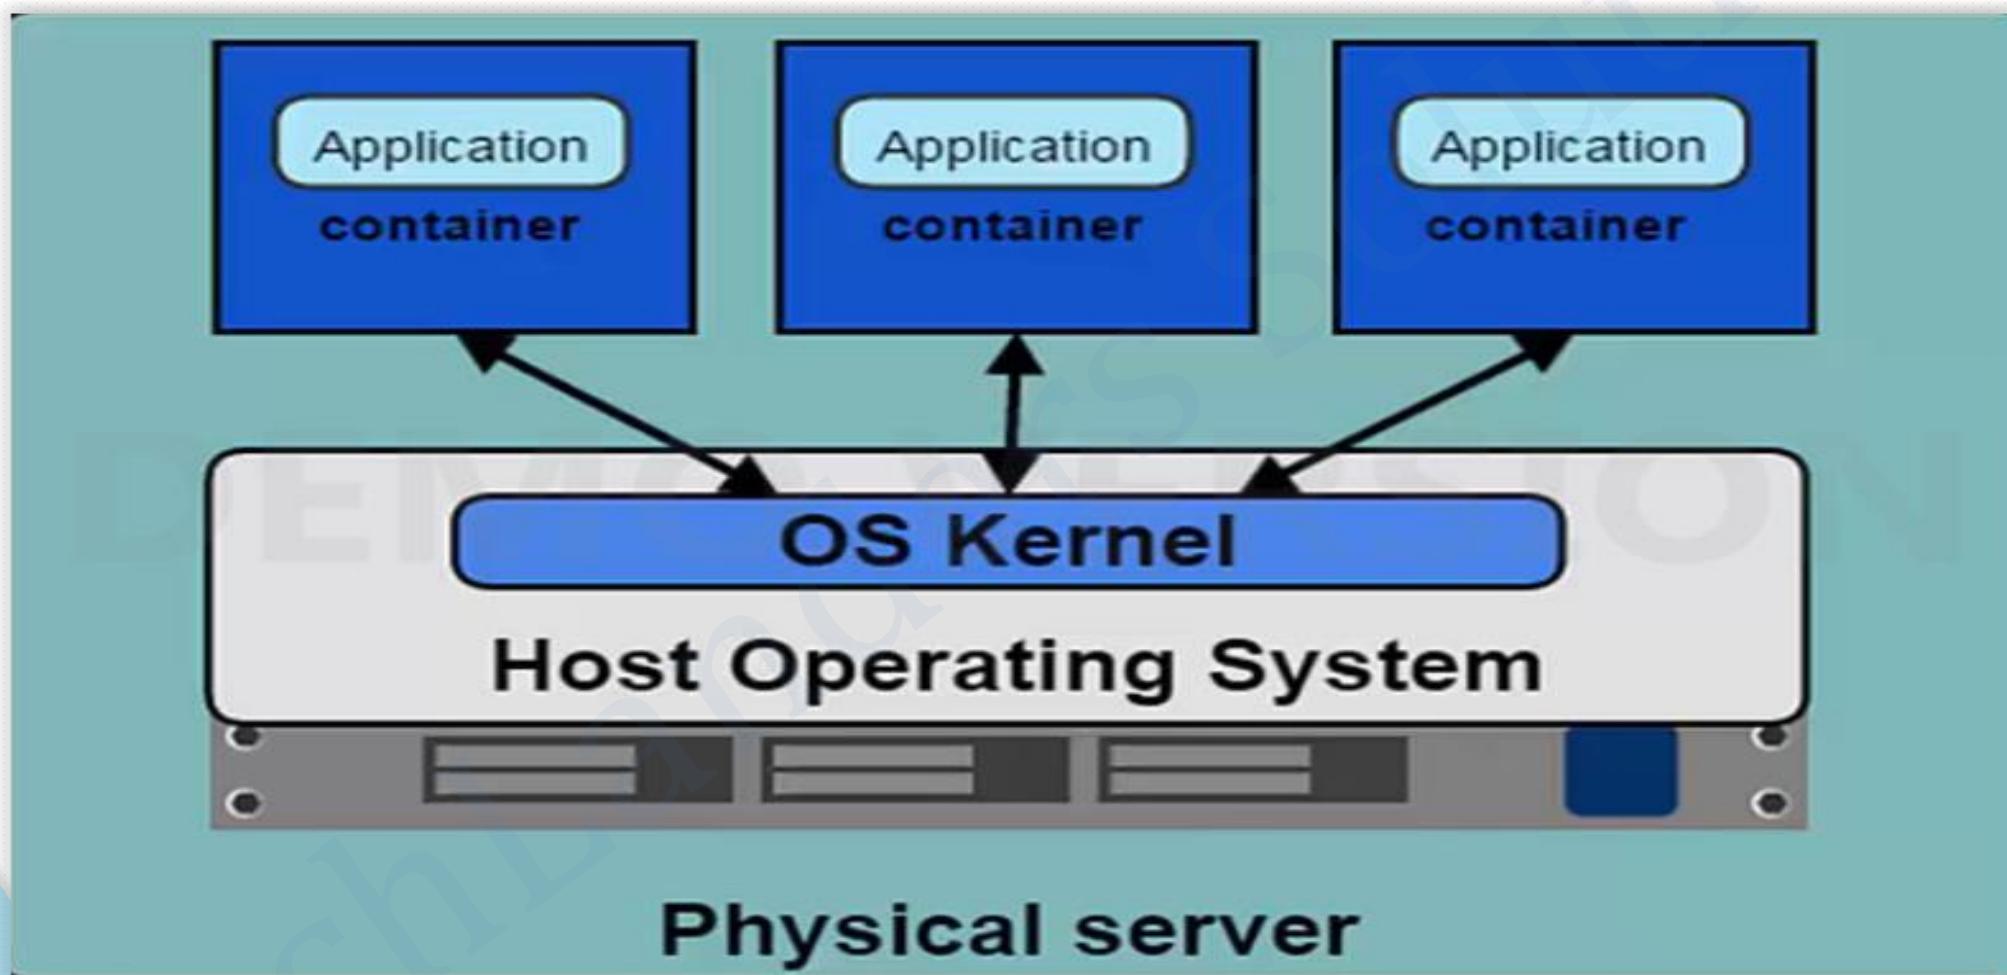


- Much Better Resource Pooling
- Extended Scaling
- Flexibility & Easy Migration
- Faster Deployments (Seconds)
- Faster Boot time (Seconds)

Introducing Containers

- Container based virtualization uses the kernel on the host's operating system to run multiple guest instances
- Each guest instance is called a “Container”
- Each container has its own
 - Root Filesystem
 - Processes
 - Memory
 - Devices
 - Network Ports
- From outside it looks like a VM but it's not a VM

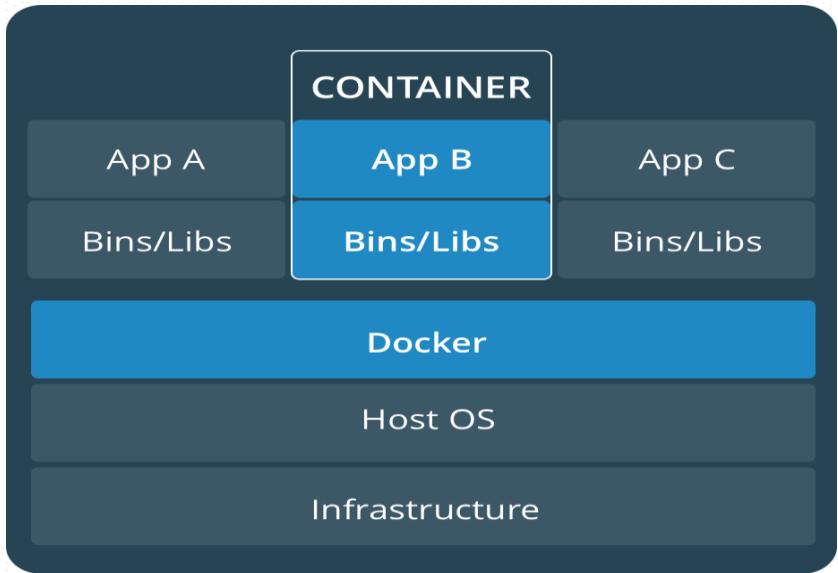
Overview of Containers



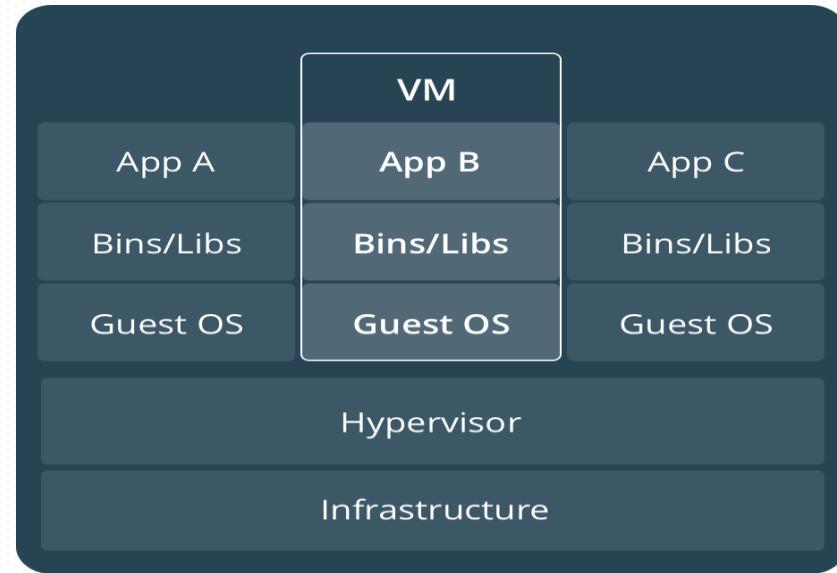
Containers VS VM's

- Container are more light weight
- No need to install dedicated guest OS, no virtualization like VM is required
- Stop/Start time is very fast
- Less CPU, RAM, Storage Space required
- More containers per machine than VM's
- Great Portability

Containers VS VM's

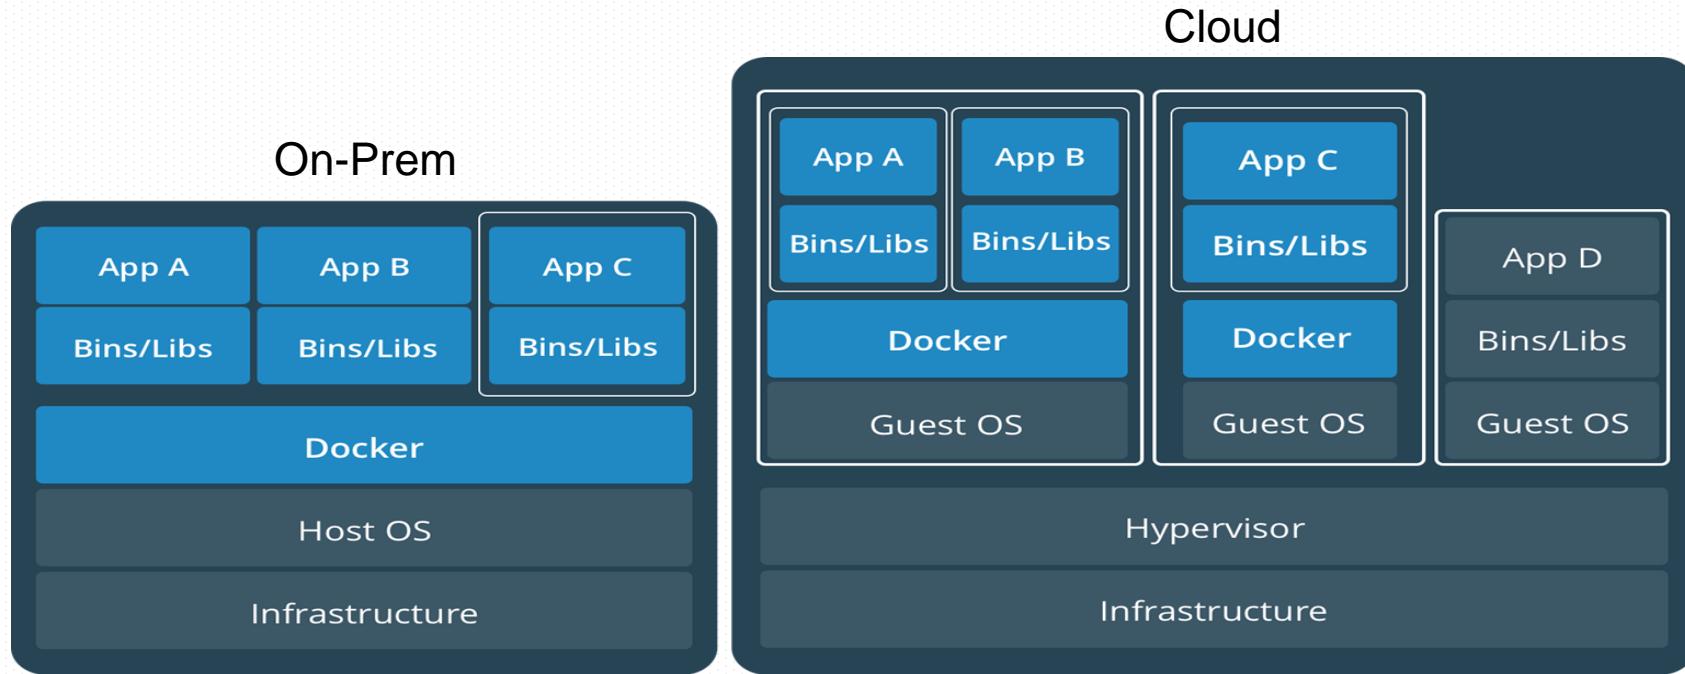


Containers are an app level construct



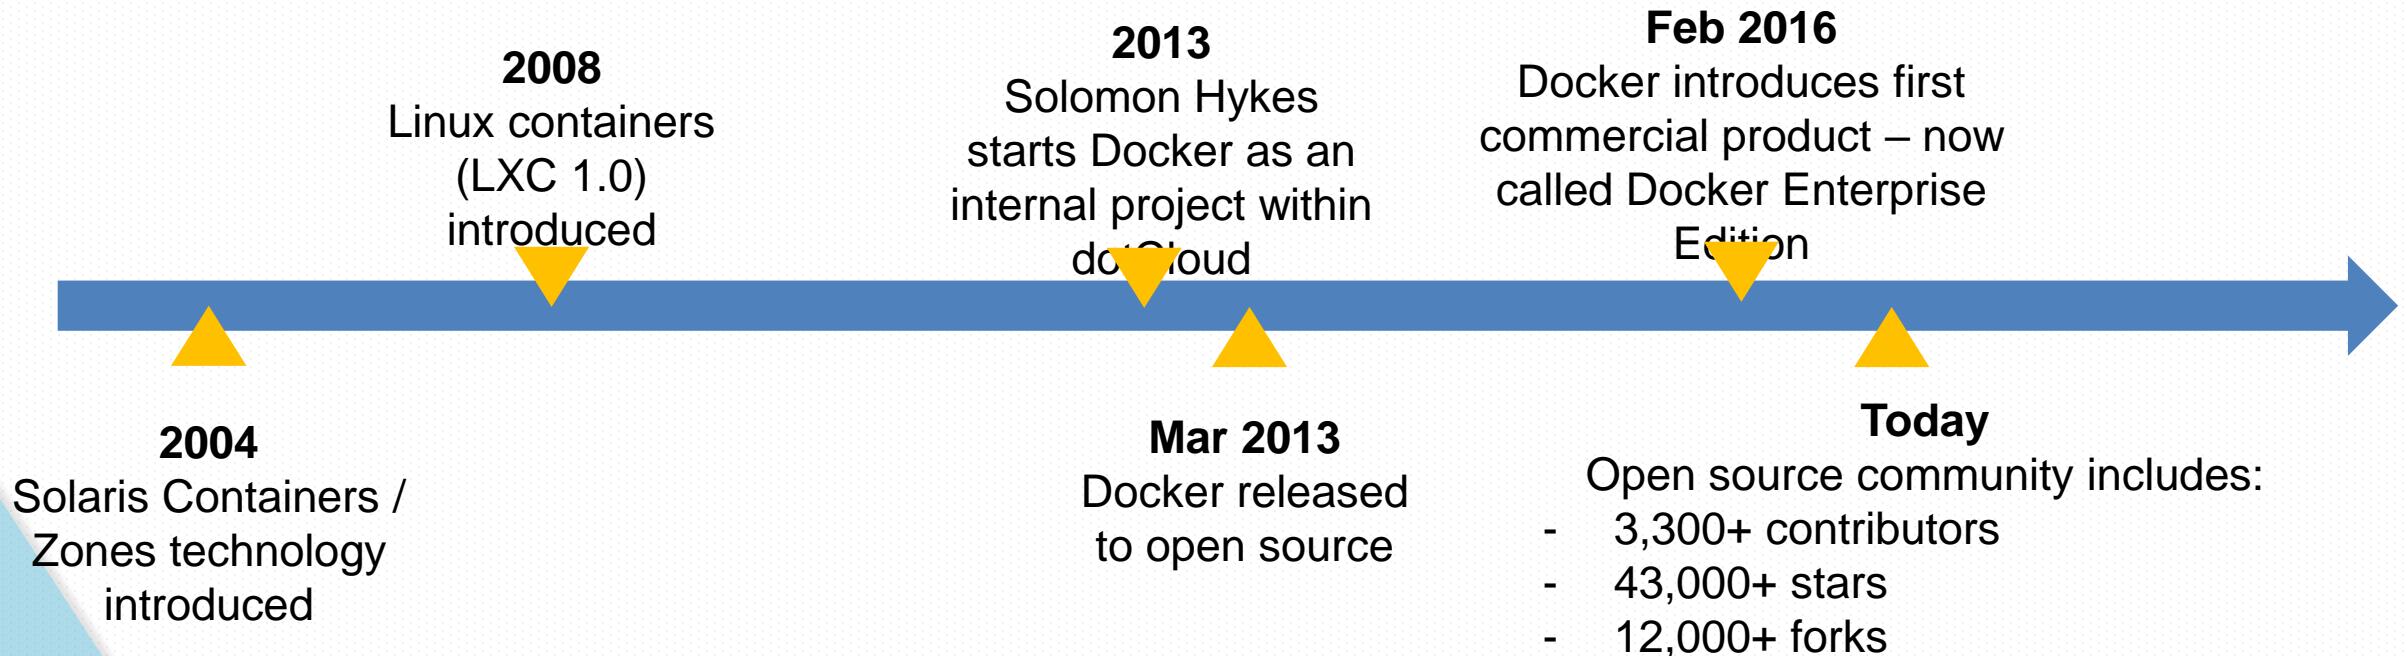
VMs are an infrastructure level construct to turn one machine into many servers

Containers and VM's together



Containers and VMs together provide a tremendous amount of flexibility for IT to optimally deploy and manage apps.

Origins of Docker Project



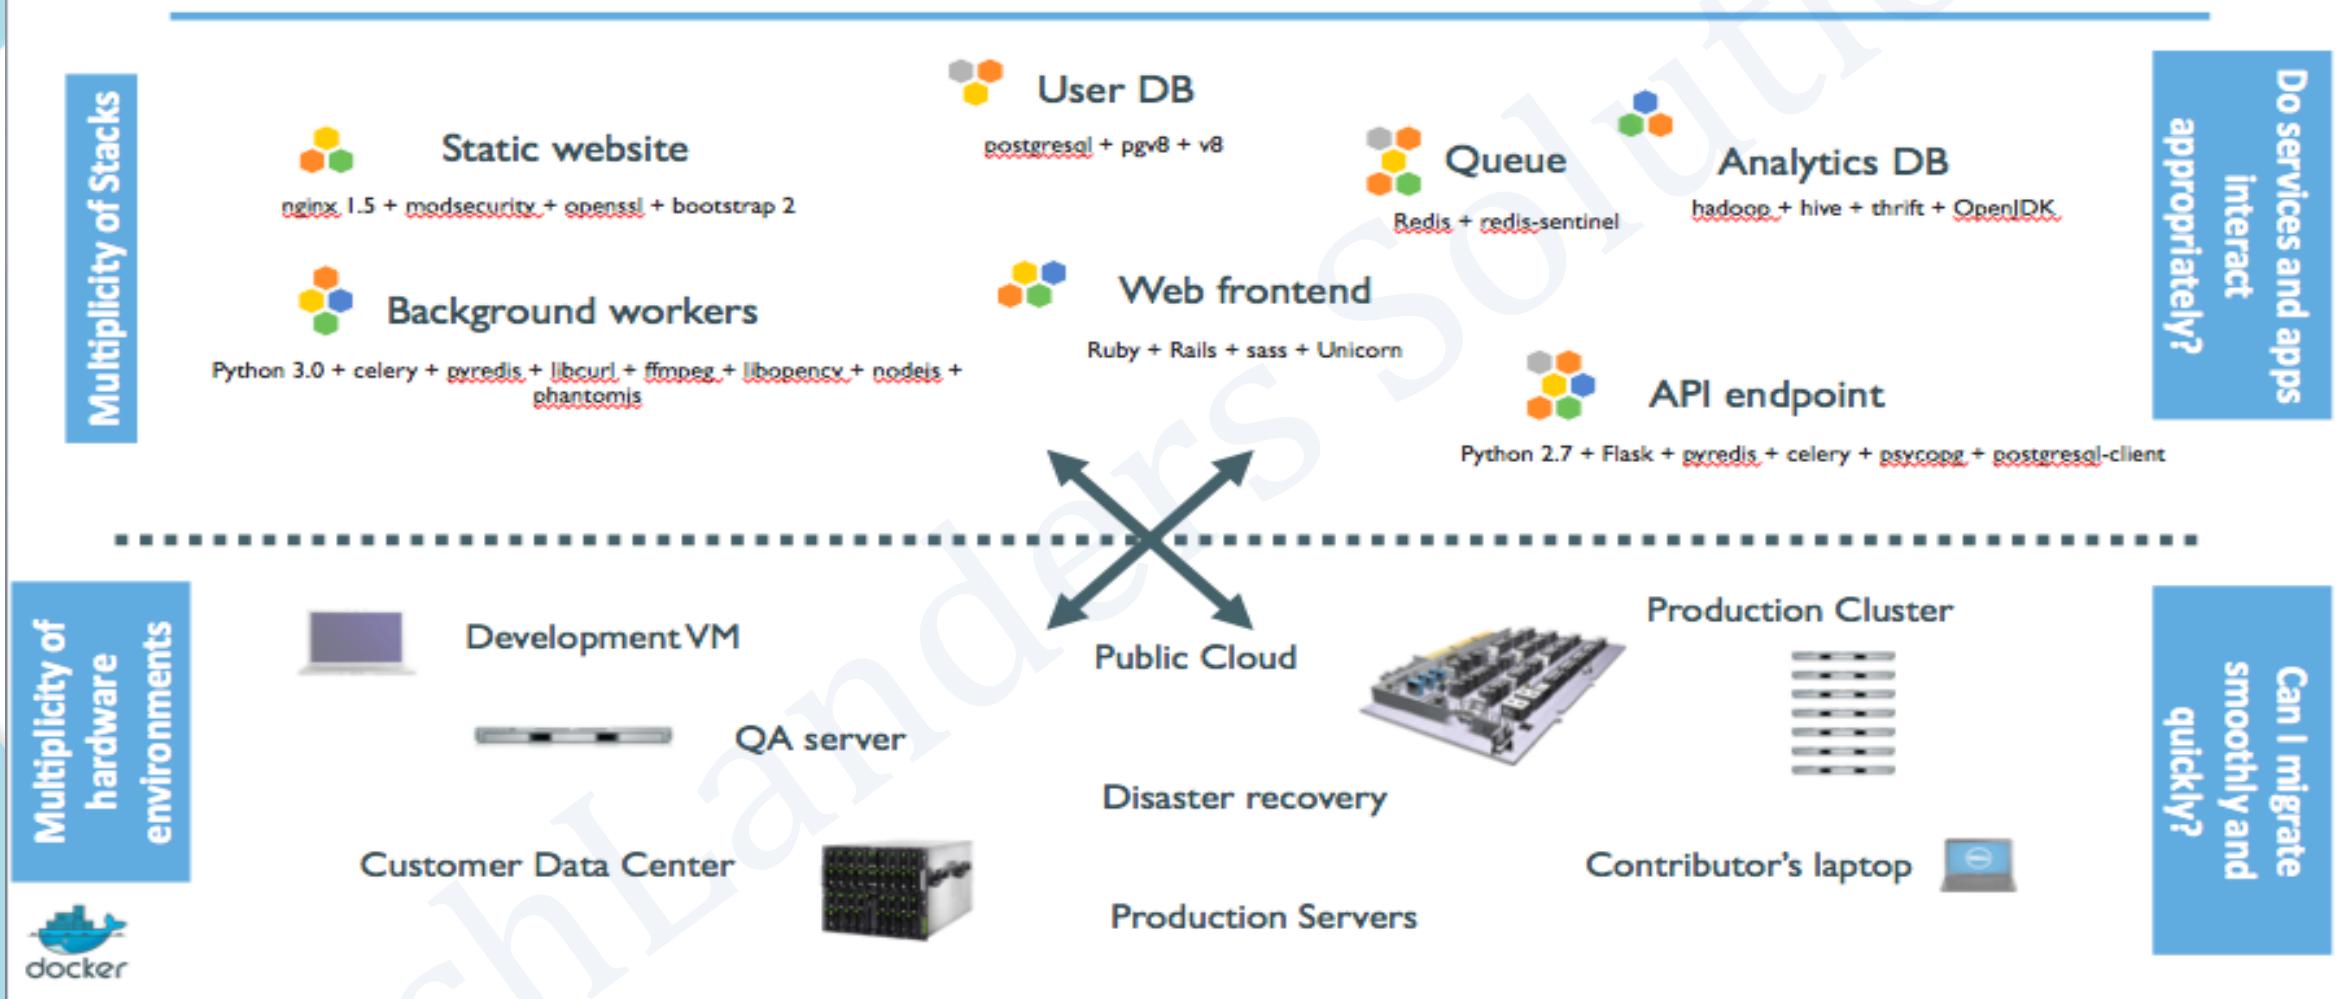
Origins of Docker Project

- 'dotCloud' was operating a PaaS platform, using a custom container engine.
- This engine was based on 'OpenVZ' (and later, LXC) and AUFS.
- It started (circa 2008) as a single Python script.
- By 2012, the engine had multiple (~10) Python components. (and ~100 other micro-services!)
- End of 2012, 'dotCloud' refractors this container engine.
- The codename for this project is "Docker."

About Docker Inc.

- Docker Inc. Formerly 'dotCloud' Inc, used to be a French company
- Docker Inc. is the primary sponsor and contributor to the Docker Project:
 - Hires maintainers and contributors.
 - Provides infrastructure for the project.
 - Runs the Docker Hub.
- HQ in San Francisco.
- Backed by more than 100M in venture capital.

Deployment Problem



Matrix Checks

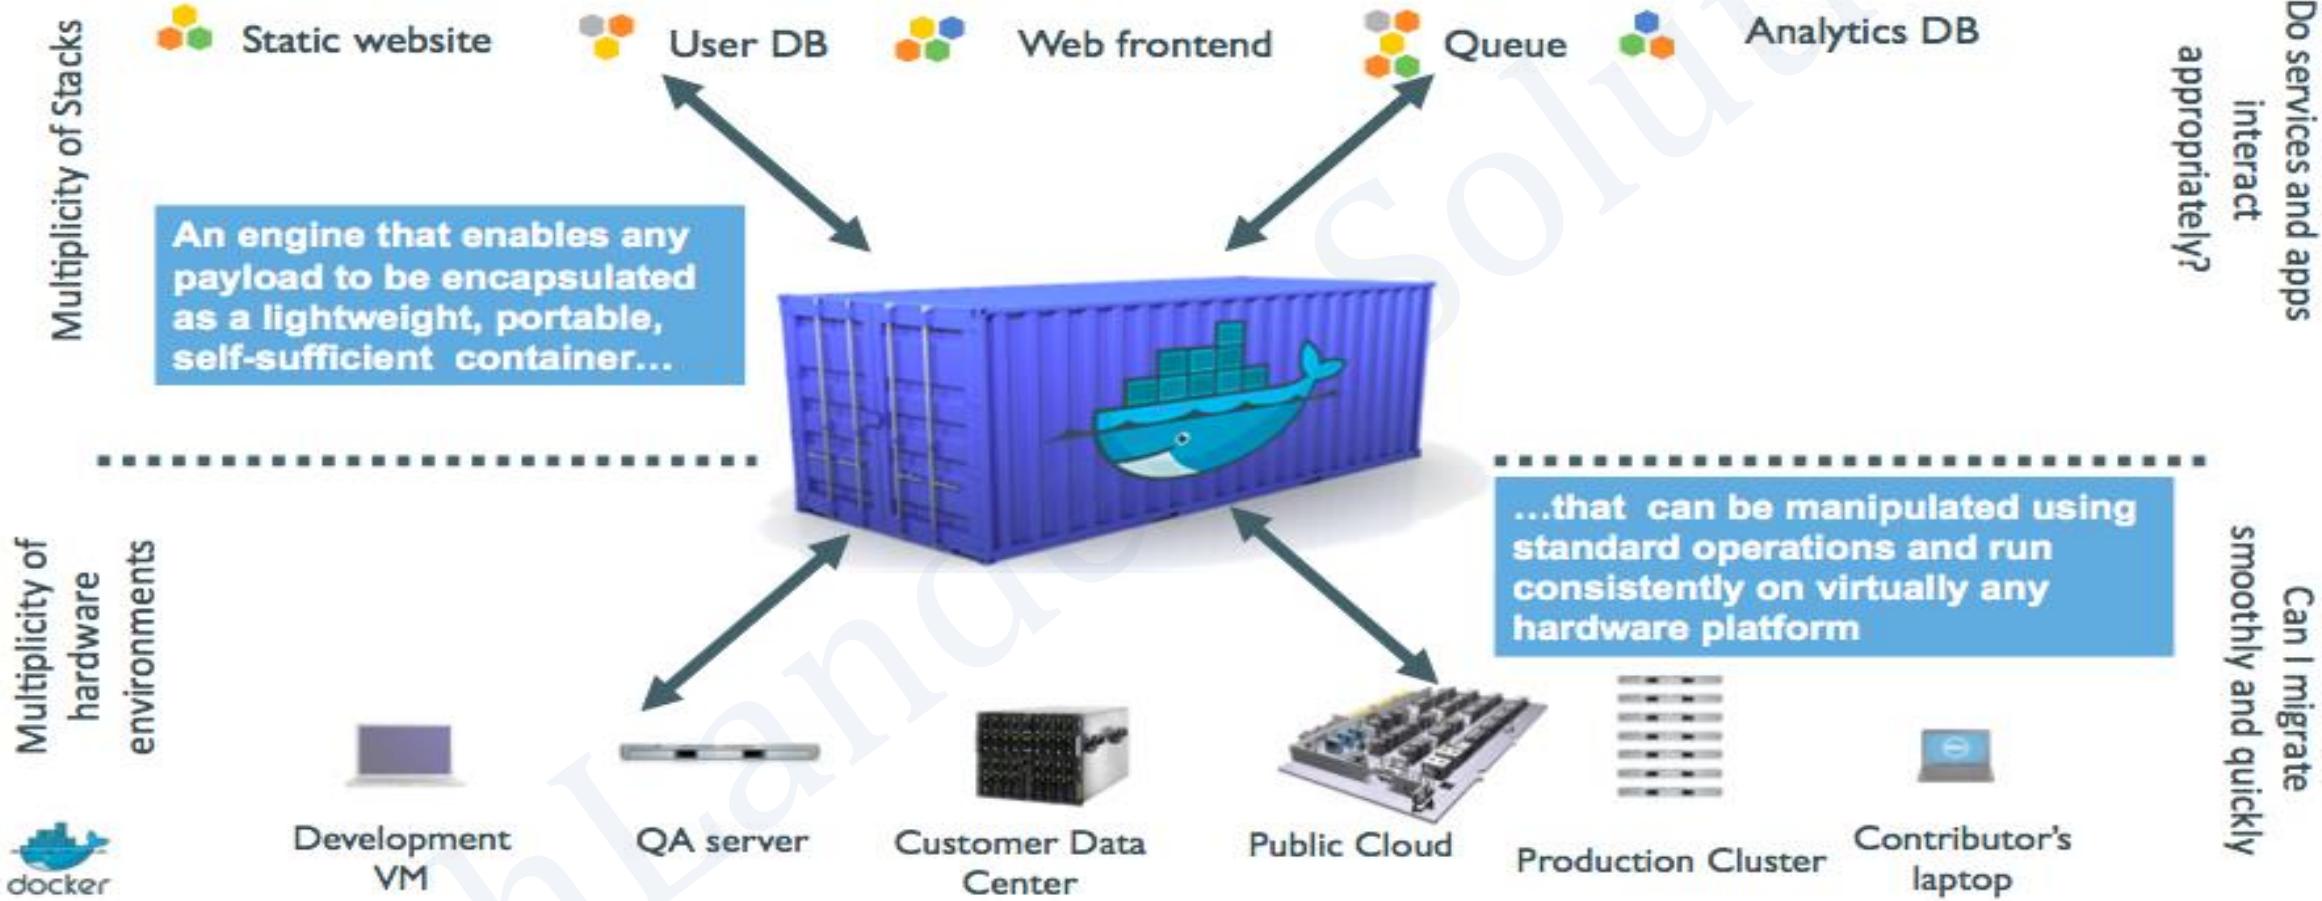
	Static website	?	?	?	?	?	?	?						
	Web frontend	?	?	?	?	?	?	?						
	Background workers	?	?	?	?	?	?	?						
	User DB	?	?	?	?	?	?	?						
	Analytics DB	?	?	?	?	?	?	?						
	Queue	?	?	?	?	?	?	?						
	Development VM		QA Server		Single Prod Server		Onsite Cluster		Public Cloud		Contributor's laptop		Customer Servers	



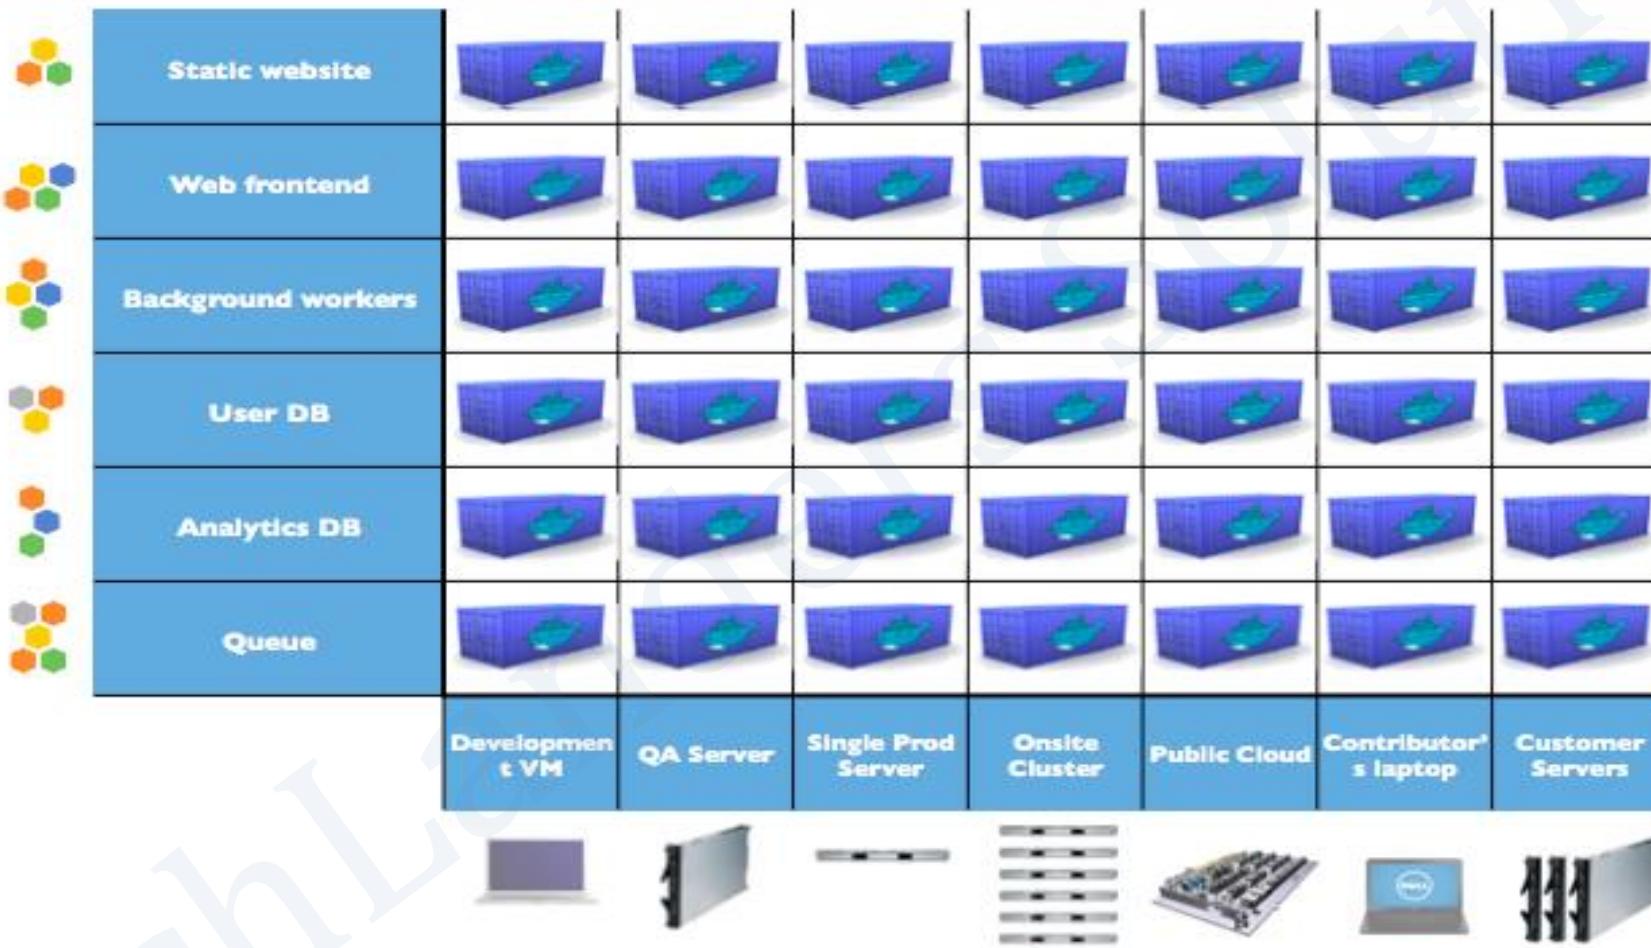
Intermodal Shipping Containers



Shipping Container for Applications



Eliminate the Matrix



Results

Speed

- No OS to boot = applications online in seconds

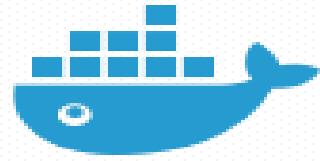
Portability

- Less dependencies between process layers = ability to move between infrastructure

Efficiency

- Less OS overhead
- Improved VM density

Adoption in Just 4 years



14M

Docker
Hosts



900K

Docker
apps



77K%

Growth in
Docker job
listings



12B

Image pulls
Over 390K%
Growth



3300

Project
Contributors

Why Docker?



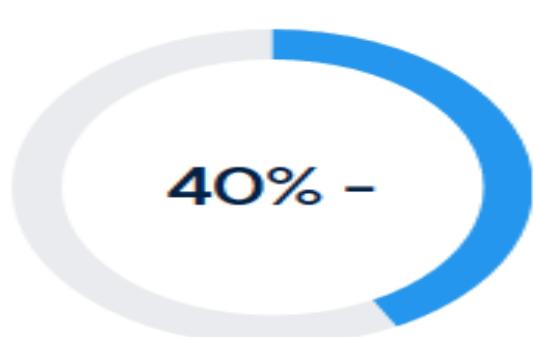
Faster Time to Market



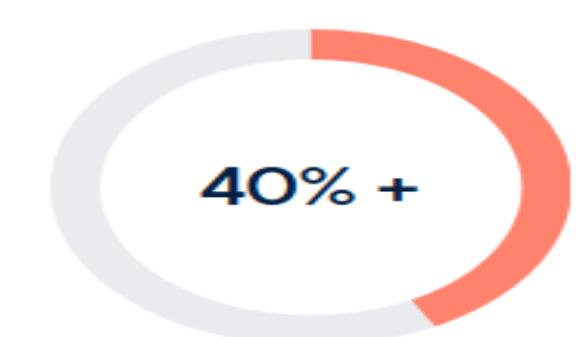
Developer Productivity



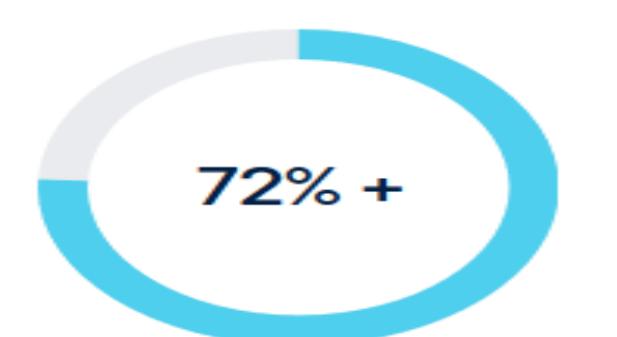
Deployment Velocity



IT Infrastructure Reduction



IT Operational Efficiency



Faster Issue Resolution

Session: 2

Docker Components

Docker Overview

- ▶ Docker is an open platform for developing, shipping, and running applications.
- ▶ Docker enables you to separate your applications from your infrastructure so you can deliver software quickly.
- ▶ With Docker, you can manage your infrastructure in the same ways you manage your applications.
- ▶ By using Docker's methodologies for shipping, testing, and deploying, you can reduce time of customer delivery.

Docker Platform

- ▶ Docker provides the ability to package and run an application in a loosely isolated environment called a container. The **isolation** and **security** allow you to run many containers simultaneously on a given host.
- ▶ Because of the **lightweight** nature of containers, which run without the extra load of a hypervisor, you can run more containers on a given hardware combination than if you were using virtual machines.

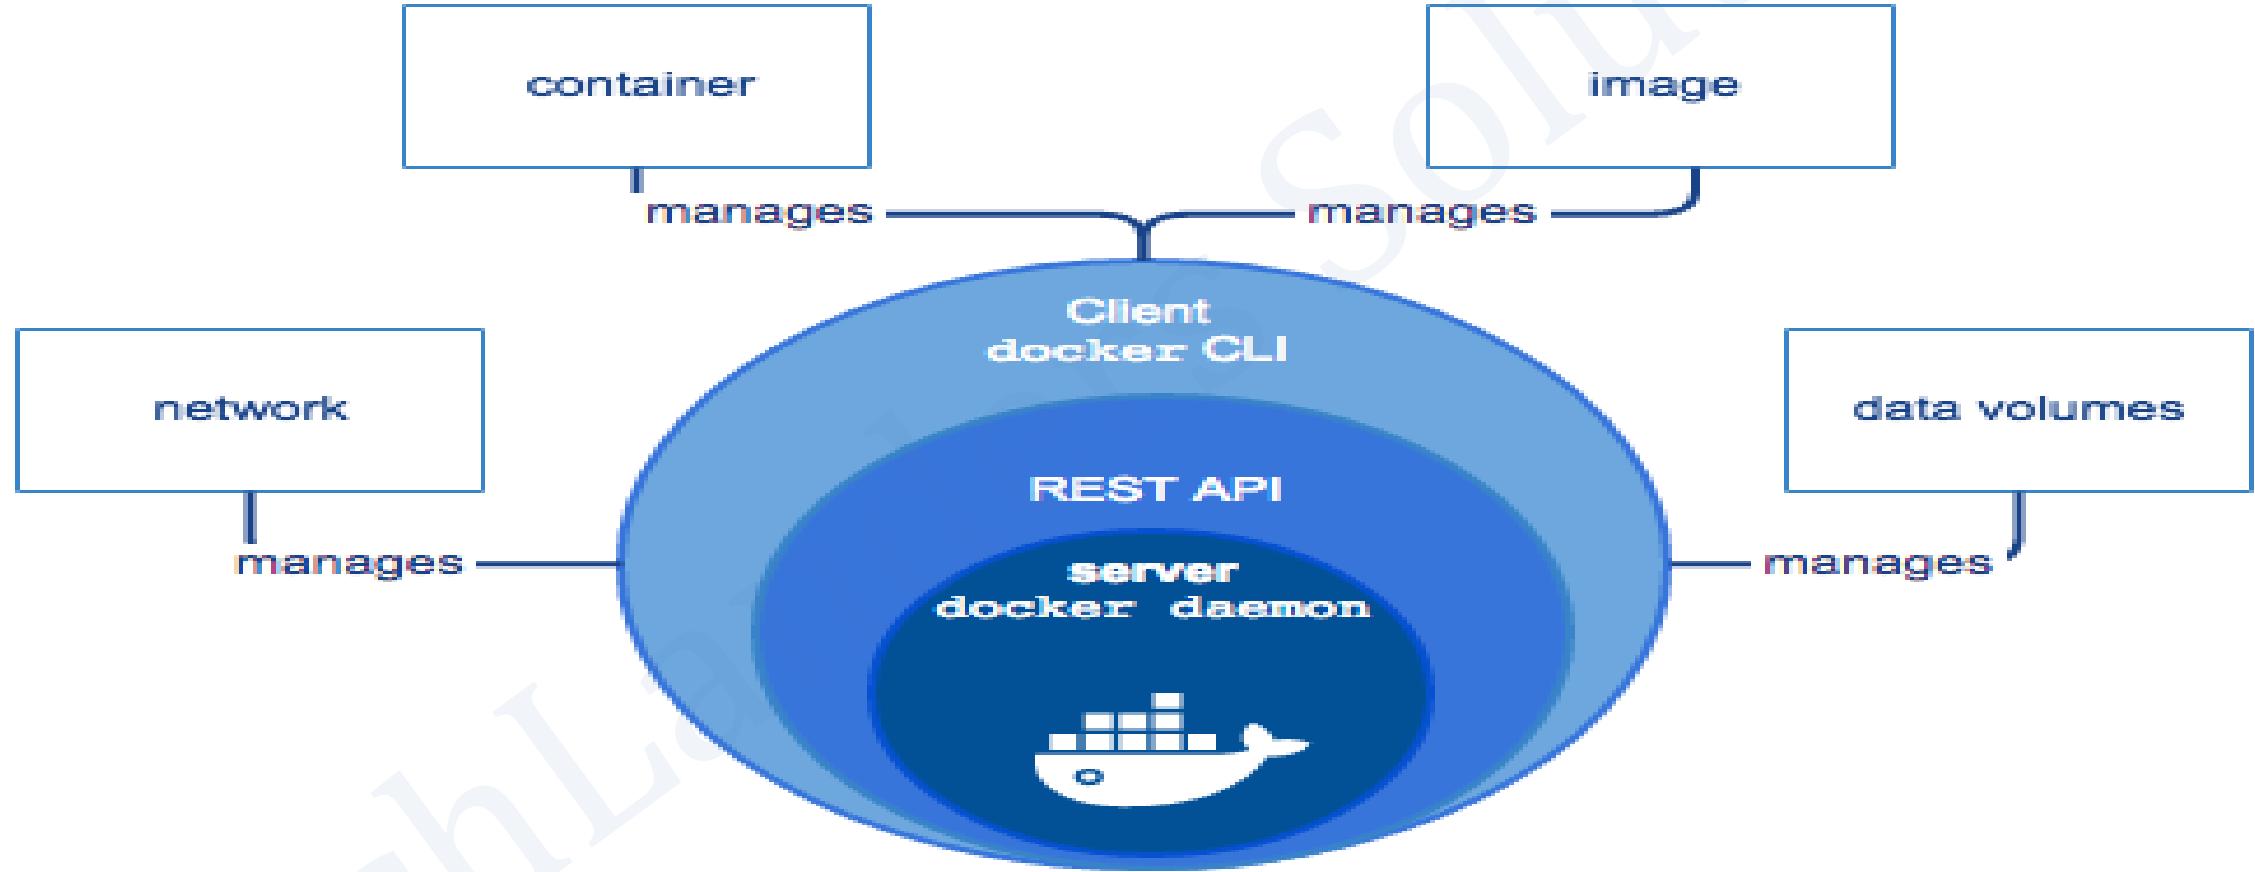
Docker Test

- ▶ docker --version
- ▶ Test the docker functioning: docker run hello-world

Docker Platform

- ▶ Docker provides tooling and a platform to manage the lifecycle of your containers:
 - Encapsulate your applications (and supporting components) into Docker containers
 - Distribute and ship those containers to your teams for further development and testing
 - Deploy those applications to your production environment, whether it is in a local data center or the Cloud

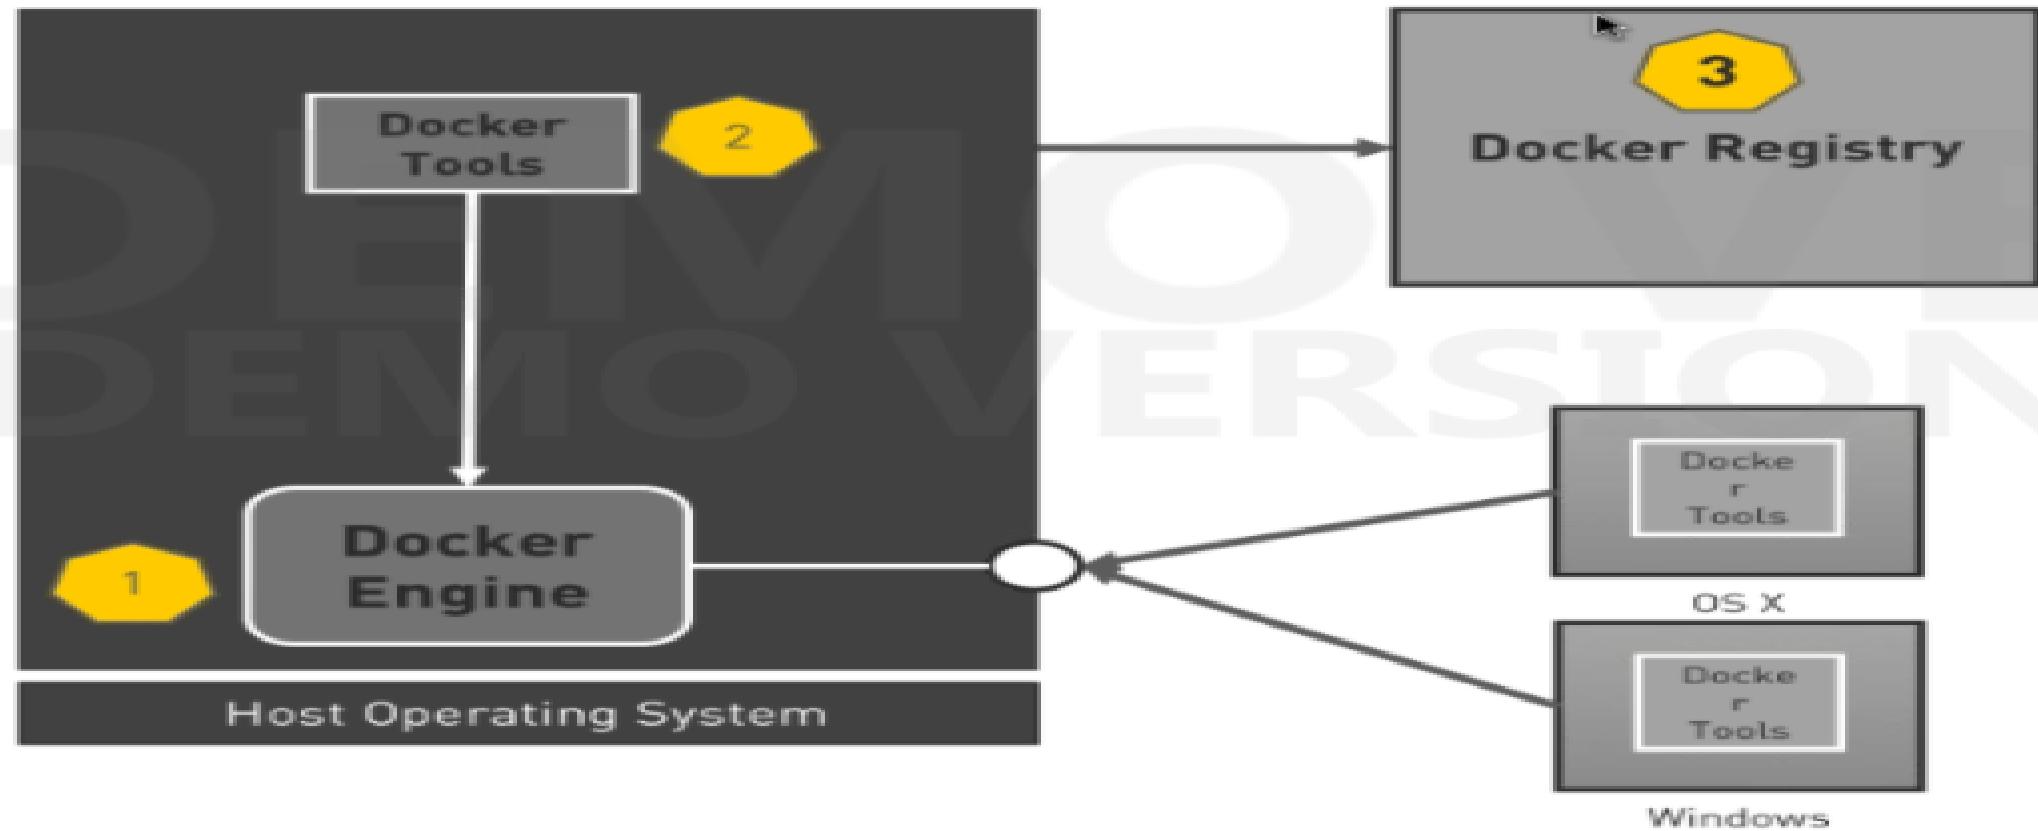
Docker Engine



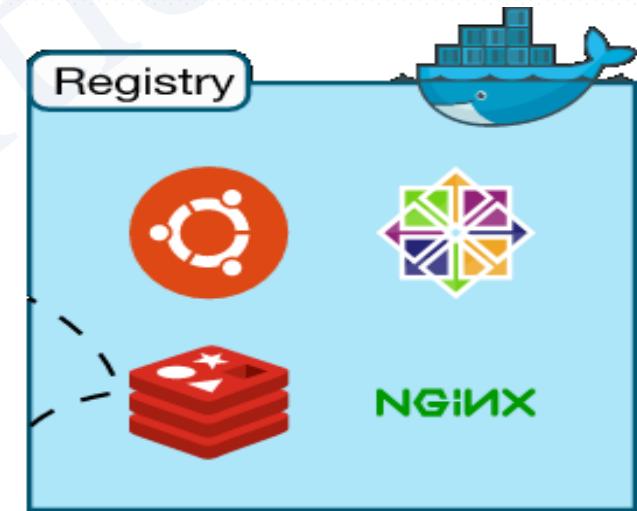
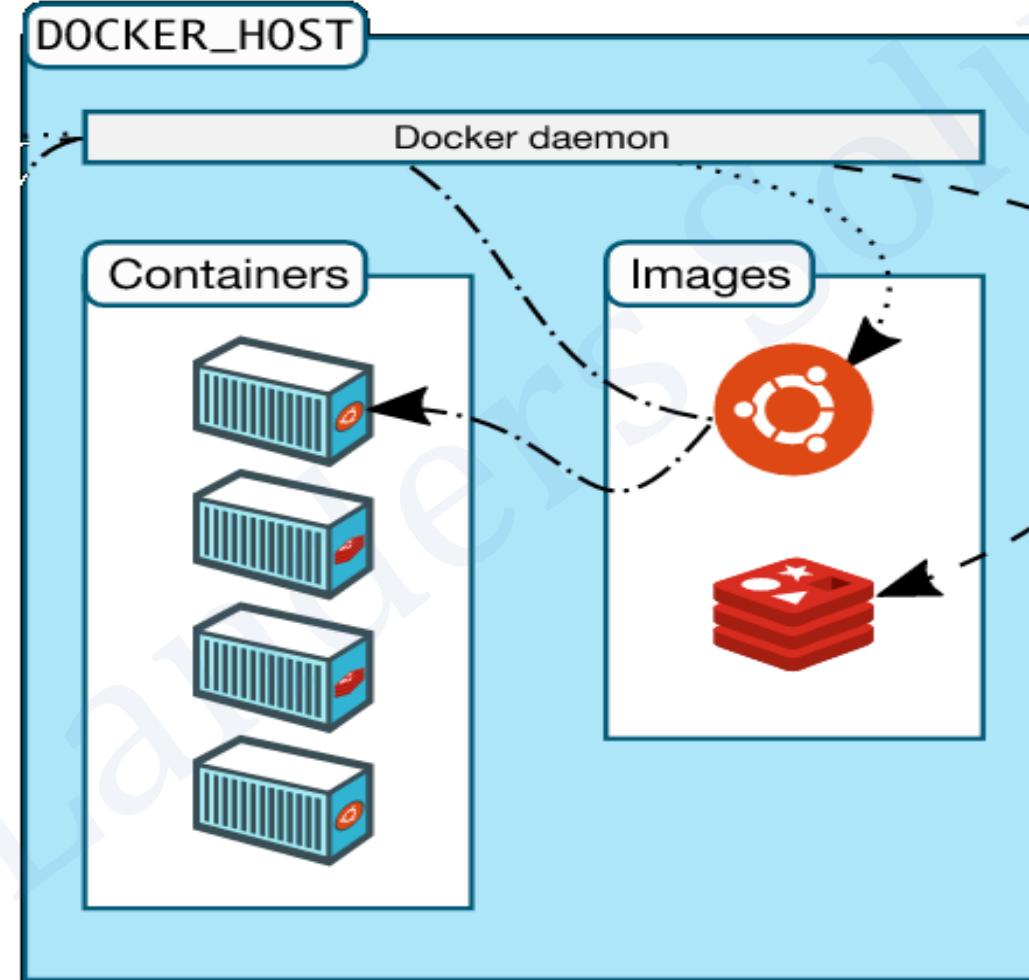
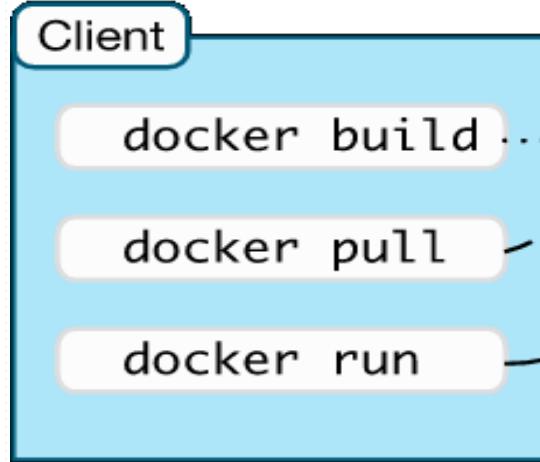
Docker Engine

- ▶ Docker Engine is a client-server application with these major components:
 - A server which is a type of long-running program called a daemon process.
 - A REST API which specifies interfaces that programs can use to talk to the daemon and instruct it what to do.
 - A command line interface (CLI) client.
- ▶ The CLI uses the Docker REST API to control or interact with the Docker daemon through scripting or direct CLI commands
- ▶ The daemon creates and manages Docker objects, such as images, containers, networks, and data volumes

Docker Architecture



Docker Architecture



Docker Architecture

- ▶ Docker uses a client-server architecture.
- ▶ The Docker *client* talks to the Docker *daemon*, which does the heavy lifting of building, running, and distributing your Docker containers.
- ▶ The Docker client and daemon *can* run on the same system, or you can connect a Docker client to a remote Docker daemon.
- ▶ The Docker client and daemon communicate using a REST API, over UNIX sockets or a network interface.

Docker Architecture

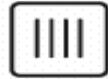
- ▶ The Docker Daemon
 - The Docker daemon runs on a host machine. The user uses the Docker client to interact with the daemon.
- ▶ The Docker Client
 - The Docker client, in the form of the docker binary, is the primary user interface to Docker.
 - It accepts commands and configuration flags from the user and communicates with a Docker daemon.

Docker Architecture



Image

The basis of a Docker container. The content at rest.



Container

The image when it is ‘running.’ The standard unit for app service



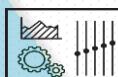
Engine

The software that executes commands for containers. Networking and volumes are part of Engine. Can be clustered together.



Registry

Stores, distributes and manages Docker images



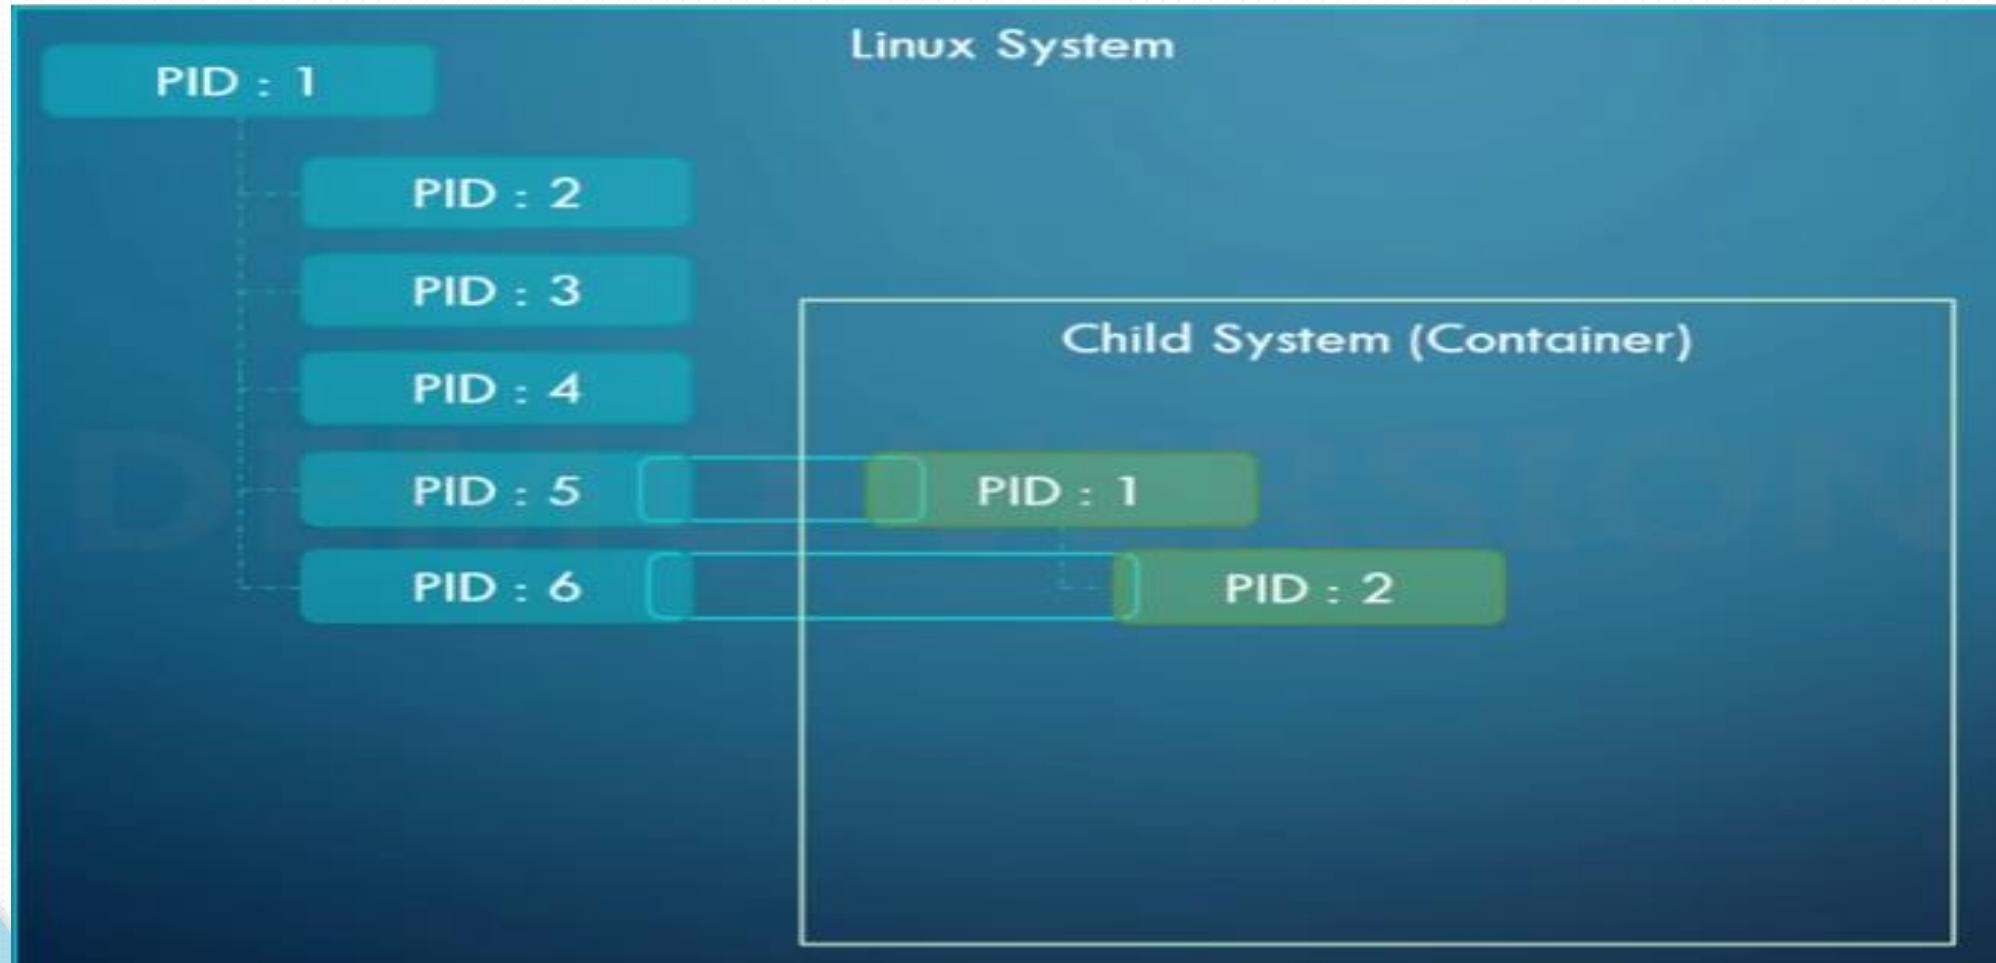
Docker Architecture

- Lets understand how does applications works in isolation under the hood.
- Docker uses namespace which is nothing but an isolated environment like VM, but on top of VM called container.



Docker Architecture

- The demonstration will be given once we understand the container operations.



Docker Images

- ▶ A Docker image is a read-only template with instructions for creating a Docker container.
- ▶ For example, an image might contain an Ubuntu operating system with Apache web server and your web application installed. You can build or update images from scratch or download and use images created by others.
- ▶ A docker image is described in text file called a Dockerfile, which has a simple, well-defined syntax.
- ▶ Docker images are the build component of Docker.

Docker Containers

- ▶ A Docker container is a running instance of a Docker image.
- ▶ You can run, start, stop, move, or delete a container using Docker API or CLI commands.
- ▶ When you run a container, you can provide configuration metadata such as networking information or environment variables.
- ▶ Each container is an isolated and secure application platform, but can be given access to resources running in a different host or container, as well as persistent storage or databases.
- ▶ Docker containers are the run component of Docker.

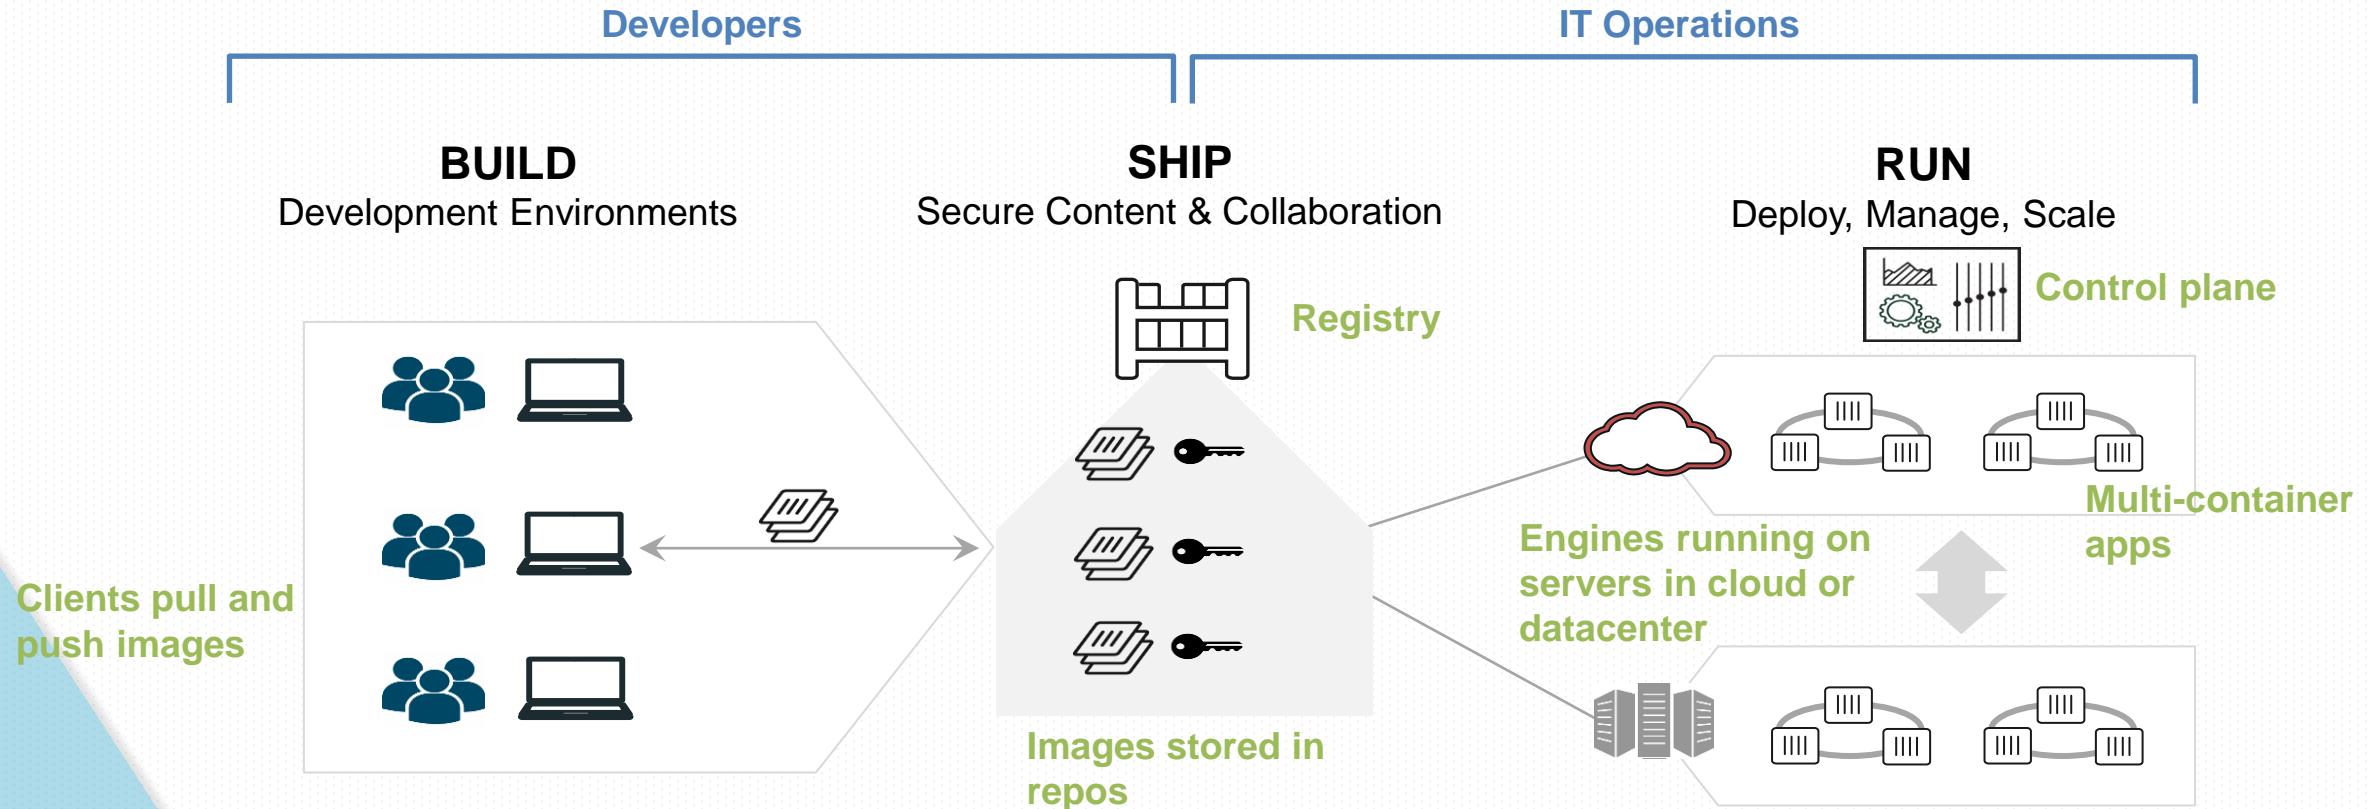
Docker Registries

- ▶ A docker registry is a library of images.
- ▶ A registry can be public or private, and can be on the same server as the Docker daemon or Docker client, or on a totally separate server.
- ▶ Docker registries are the distribution component of Docker.
- ▶ “Docker Hub” is known as global registry.

Docker Features

- Lightweight
 - Containers running on a single machine all share the same operating system kernel so they start instantly and make more efficient use of RAM. Images are constructed from layered filesystems so they can share common files, making disk usage and image downloads much more efficient.
- Open
 - Docker containers are based on open standards allowing containers to run on all major Linux distributions and Microsoft operating systems with support for every infrastructure.
- Secure
 - Containers isolate applications from each other and the underlying infrastructure while providing an added layer of protection for the application.

Container as a Service



Session: 3

Environment

Docker Engine Install Demo

- ▶ Docker Engine/Client would be installed on Training Environment as demo LAB.
- ▶ <https://docs.docker.com/engine/installation/linux/centos/>
- ▶ <https://www.cyberciti.biz/faq/how-to-install-docker-on-amazon-linux-2/>

Docker Installation Key Points

- docker.service Systemd File:

```
[root@TechLanders lib]# more /usr/lib/systemd/system/docker.service
```

- Docker Socket file:

```
[root@TechLanders lib]# file /var/run/docker.sock  
/var/run/docker.sock: socket
```

- Docker PID file and Docker Container PID file (This is Docker Daemon which will have pid1)

```
root@TechLanders run]# cat /var/run/docker.pid  
3715
```

```
[root@TechLanders libcontainerd]# cat /var/run/docker/containerd/docker-containerd.pid  
4251
```

- Docker Detailed Information:

```
[root@TechLanders libnetwork]# docker info
```

Manage Docker as a non-root user

- ▶ The docker daemon binds to a Unix socket instead of a TCP port.
- ▶ By default that Unix socket is owned by the user "root" and other users can only access it using sudo.
- ▶ The docker daemon always runs as the root user.
- ▶ If you don't want to use sudo when you use the docker command, add users to Unix group called "docker"
example: `usermod -aG docker <user-name>`
- ▶ When the docker daemon starts, it makes the ownership of the Unix socket read/writable by the docker group.

Session: 4

Containers

How Container works

- ▶ A container uses the host machine's Linux kernel, and consists of any extra files you add when the image is created, along with metadata associated with the container at creation or when the container is started.
- ▶ Each container is built from an image.
- ▶ The image defines the container's contents, which process to run when the container is launched, and a variety of other configuration details.
- ▶ The Docker image is read-only. When Docker runs a container from an image, it adds a read-write layer on top of the image (using a UnionFS) in which your application runs.

How Container works

- When you use the "docker run" CLI command, the Docker Engine client instructs the Docker daemon to run a container.

```
docker run ubuntu ps ax
```

- This example tells the Docker daemon to run a container using the centos Docker image, to remain in the foreground in interactive mode (-i), provide a tty terminal (-t) and to run the /bin/bash command.

```
docker run -i -t centos /bin/bash
```

Because if you exit the current running process /bin/bash (pid 1), container will stop/exit. Use “Ctrl pq” to safe exit without stopping the container.

How Container works

```
[root@TechLanders libcontainerd]# docker run -i -t centos /bin/bash
Unable to find image 'centos:latest' locally
latest: Pulling from library/centos
d9aaaf4d82f24: Pull complete
Digest: sha256:eba772bac22c86d7d6e72421b4700c3f894ab6e35475a34014ff8de74c10872e
Status: Downloaded newer image for centos:latest
[root@9a06b1a61fc5 /]#
```

```
[root@TechLanders overlay]# ls -lrt /var/lib/docker/image/overlay2/repositories.json
-rw----- 1 root root 545 Sep 18 03:17 /var/lib/docker/image/overlay2/repositories.json
```

How Container works

```
[root@TechLanders overlay]# cat repositories.json
```

```
{"Repositories":{"centos":{"centos:latest":"sha256:196e0ce0c9fb31da595b893dd39bc9fd4aa78a474bbdc21459a3ebe855b7768","centos@sha256:eba772bac22c86d7d6e72421b4700c3f894ab6e35475a34014ff8de74c10872e":"sha256:196e0ce0c9fb31da595b893dd39bc9fd4aa78a474bbdc21459a3ebe855b7768"},"hello-world":{"hello-world:latest":"sha256:05a3bd381fc2470695a35f230afefd7bf978b566253199c4ae5cc96fafa29b37","hello-world@sha256:1f19634d26995c320618d94e6f29c09c6589d5df3c063287a00e6de8458f8242":"sha256:05a3bd381fc2470695a35f230afefd7bf978b566253199c4ae5cc96fafa29b37"}}}
```

```
[root@TechLanders overlay]# docker image list
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
centos	latest	196e0ce0c9fb	3 days ago	197MB
hello-world	latest	05a3bd381fc2	5 days ago	1.84kB

```
[root@TechLanders overlay]# docker image inspect centos
```

More Useful - Container

- ▶ Do something in our container:

Lets suppose we try to use “talk” for communication.

- Let's check how many packages are installed:

```
rpm -qa | wc -l
```

A non interactive - Container

- ▶ In your Docker environment, just run the following command:

```
docker run jpetazzo/clock
```

- This container just displays the time every second.
- This container will run forever.
- To stop it, press ^C.
- Docker has automatically downloaded the image jpetazzo/clock.

Run in background - Container

- ▶ Containers can be started in the background, with the -d flag (daemon mode):

```
docker run -d jpetazzo/clock
```

- We don't see the output of the container.
- But don't worry: Docker collects that output and logs it!
- docker ps -a
- docker logs <container-id>
- Docker gives us the ID of the container.

List Running Containers

- ▶ With docker ps, just like the UNIX ps command, lists running processes.

docker ps

docker ps -l

docker ps -a

docker ps -q

- The (truncated) ID of our container.
- The image used to start the container.
- That our container has been running (Up) for a couple of minutes.
- Now, start multiple containers and use “docker ps” to list them.

Stop our Container

- There are two ways we can terminate our detached container.
 - Killing it using the docker “kill” command.
 - Stopping it using the docker “stop” command.
- The first one stops the container immediately, by using the KILL signal.
- The second one is more graceful. It sends a TERM signal, and after 10 seconds, if the container has not stopped, it sends KILL.

Removing Container

- Let's remove our container:

```
docker rm <yourContainerID>
```

Attaching to a Container

- You can attach to a container:

```
docker attach <containerID>
```

- The container must be running.
- There *can be multiple clients attached to the same container.*

Restarting a Container

- When a container has exited, it is in stopped state.
- It can then be restarted with the “start” command.

```
docker start <containerID>
```

- The container must be running.
- You can also use restart command

```
docker restart <container-id>
```

LAB1

- Create one Ubuntu Container in interactive and terminal mode
- Exit out of the container
- Check the status of container
- Restart the container
- Get attached to the container
- Get out of container without exiting „ stop the container
- Remove the container
- Create the container now in detached mode and follow the same steps ..

LAB2

- Create a new nginx container using (-d) option.
- Docker run -d nginx
- Now check the difference using docker ps -a command
- Try to access this container
- Docker attach <container-ID>
 cat /etc/os-release (THIS WILL NOT WORK) (BECAUSE WE ARE NOT USING INETRACTIVE IN THE CREATION)
- Use docker exec -it <CONTAINER-ID> /bin/bash
- Remove everything
- Docker stop
- Docker rm
- Docker run -dt nginx (YOU NEED TO PROVIDE A TERMINAL HERE) → THIS IS TRUE WITH BASE IMAGES
- Docker attach <container-ID>
 cat /etc/os-release

LAB3

- *On your host machine install httpd package*

#For centos

- Yum install httpd -y

#For ubuntu :

- apt-get install apache2

- *Verify installation:*

rpm -qa | grep -i httpd

#For ubuntu

find / -name apache2 or dpkg -list | grep -i apache2

- *Create a container :*

- #For ubuntu :

docker run -it -name c1 ubuntu

docker run -it centos

LAB3 cont..

Check the package inside it:

```
dpkg --list | grep -i apache
```

*Install the package if not there :

```
apt-get update -y
```

```
apt-get install apache2
```

```
dpkg --list | grep apache2
```

#For centos :

```
rpm -qa | grep -i httpd
```

•If facing an error while installing httpd in centos container run below :

```
sudo sed -i 's/mirrorlist/#mirrorlist/g' /etc/yum.repos.d/CentOS-*
```

```
sudo sed -i 's|#baseurl=http://mirror.centos.org|baseurl=http://vault.centos.org|g' /etc/yum.repos.d/CentOS-*
```

```
sudo yum update -y
```

Bridge

`docker run ubuntu`

Docker Host

Container

Container

172.17.0.2

docker0

172.17.0.3

Container

Container

172.17.0.4

172.17.0.1

172.17.0.5

None

```
docker run \  
--network=none  
ubuntu
```

Docker Host

Container

Host

```
docker run \  
--network=host  
ubuntu
```

Docker Host

5000

Container

Session: 5

Docker – Containers Advanced

Checking Info

Checking System usage:

```
$ docker system df // use -v for verbose run
```

Checking Docker Information:

```
$ docker system info
```

Checking Docker events:

```
$ docker system events //Checking runtime events
```

Checking Logs:

```
$ docker logs {containerid}
```

Cleaning up Unused resources:

```
$ docker system prune //use -a for deleting all unused images  
docker volume prune //to remove all unused volumes
```

Container

- *Set a Name to container*

```
docker run --name my-redis -d redis
```

- *Setting a hostname from outside:*

```
docker run -dit --name container1 --hostname c1 centos
```

- *Finding stats for container:*

```
docker top container-name  
docker stats
```

- *Inspecting Docker containers*

```
docker inspect container-name
```

- *Homedirectory*

/var/lib/docker

Resource binding to a Container

Limiting memory(Quota):

```
$ docker run -dit -m 300M --name c1 redis
```

Verify :

```
Docker stats c1
```

Limiting CPUs:

```
docker run -dit --cpus 0.02 --name c2 redis
```

```
Check the nanocpus
```

```
Docker inspect c2 | grep -I cpu
```

Session: 5

Ports and Volumes

Working with Volumes

- Containers have ephemeral storage
- Data shared within the container is not accessible after it is terminated
- Volumes provide persistence to containers
- Each volume is implicitly or explicitly mapped to a host directory

Working with Volumes

- Question is who is responsible for performing all these operations in Docker.
- And the answer is “Storage Drivers”, some of the common drivers are like “aufs, zfs, overlay, overlay2, device mapper etc.)
- The choice of device driver depends upon the underlying OS.
- “aufs” is the default Storage Driver for Ubuntu, however this is not present in Centos, in such cases “device mapper” is a better option to go with.
- The key part is “docker will pick the best storage driver for you automatically”.
- Execute docker info | more to know about the drivers used by the docket for your host.

Volume mapping

- cd /home/ubuntu/raman/
 - touch hostfile
 - #On the command-line, with the -v flag for “docker run”.
docker run -it --name c1 -v /home/ubuntu/raman:/appdata centos
- # Edit the file hostfile in appdata folder in container
- #Exit out the container safely
- # See the changes getting reflected in local as well.

Listing Volumes

- If a container is stopped, its volumes still exist and are available.
- Since Docker 1.9, we can see all existing volumes:
`docker volume ls`

Docker Ports and Volumes

- Two most important things in docker are Ports and Volumes.
- docker images
- docker run -d nginx:latest
- docker ps
- docker inspect <container-id>| grep -i ip
- docker ps
- docker stop <container-id>
- docker rm `docker ps -a -q`

Docker Ports and Volumes

- Lets expose the containers to a random port by docker itself
- `docker run -d --name=nginxserver -P nginx:latest`
- `docker ps` (you will get a port mapped by the docker for you)
- `docker port <conatainer-name> $CONTAINERPORT`
- `docker stop yogeshserver`

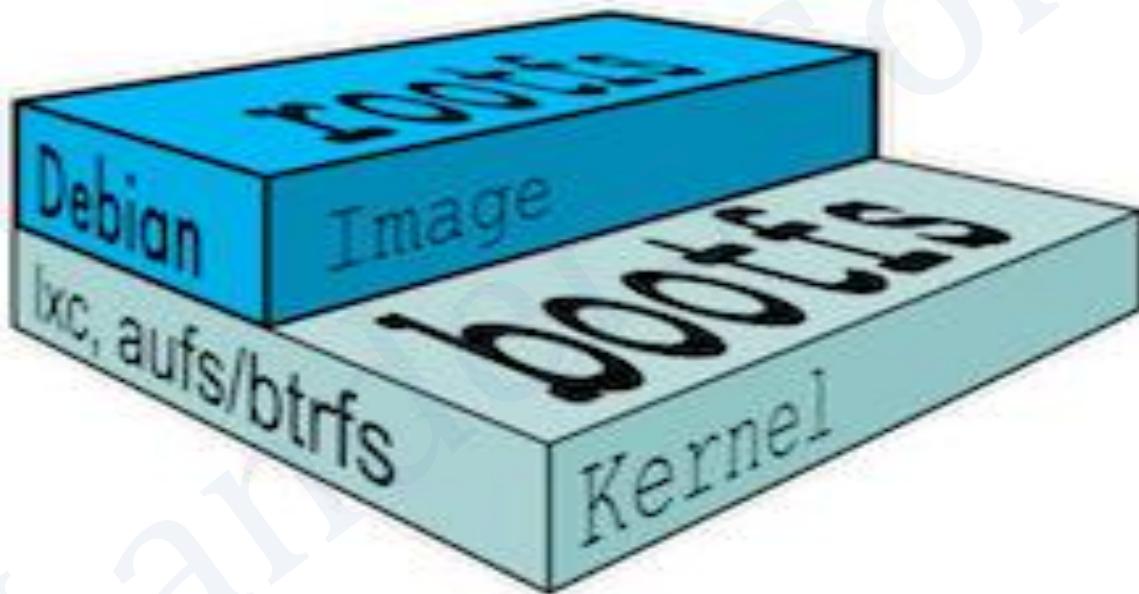
Docker Ports and Volumes

- But how to bound the container with a particular port:
- docker run -d -p 8080:80 nginx
- Even you can bound multiple ports:
- docker run -d -p 8080:80,8081:443 nginx
- Once done stop and remove the container

Session: 6

Docker - Images

Docker Images



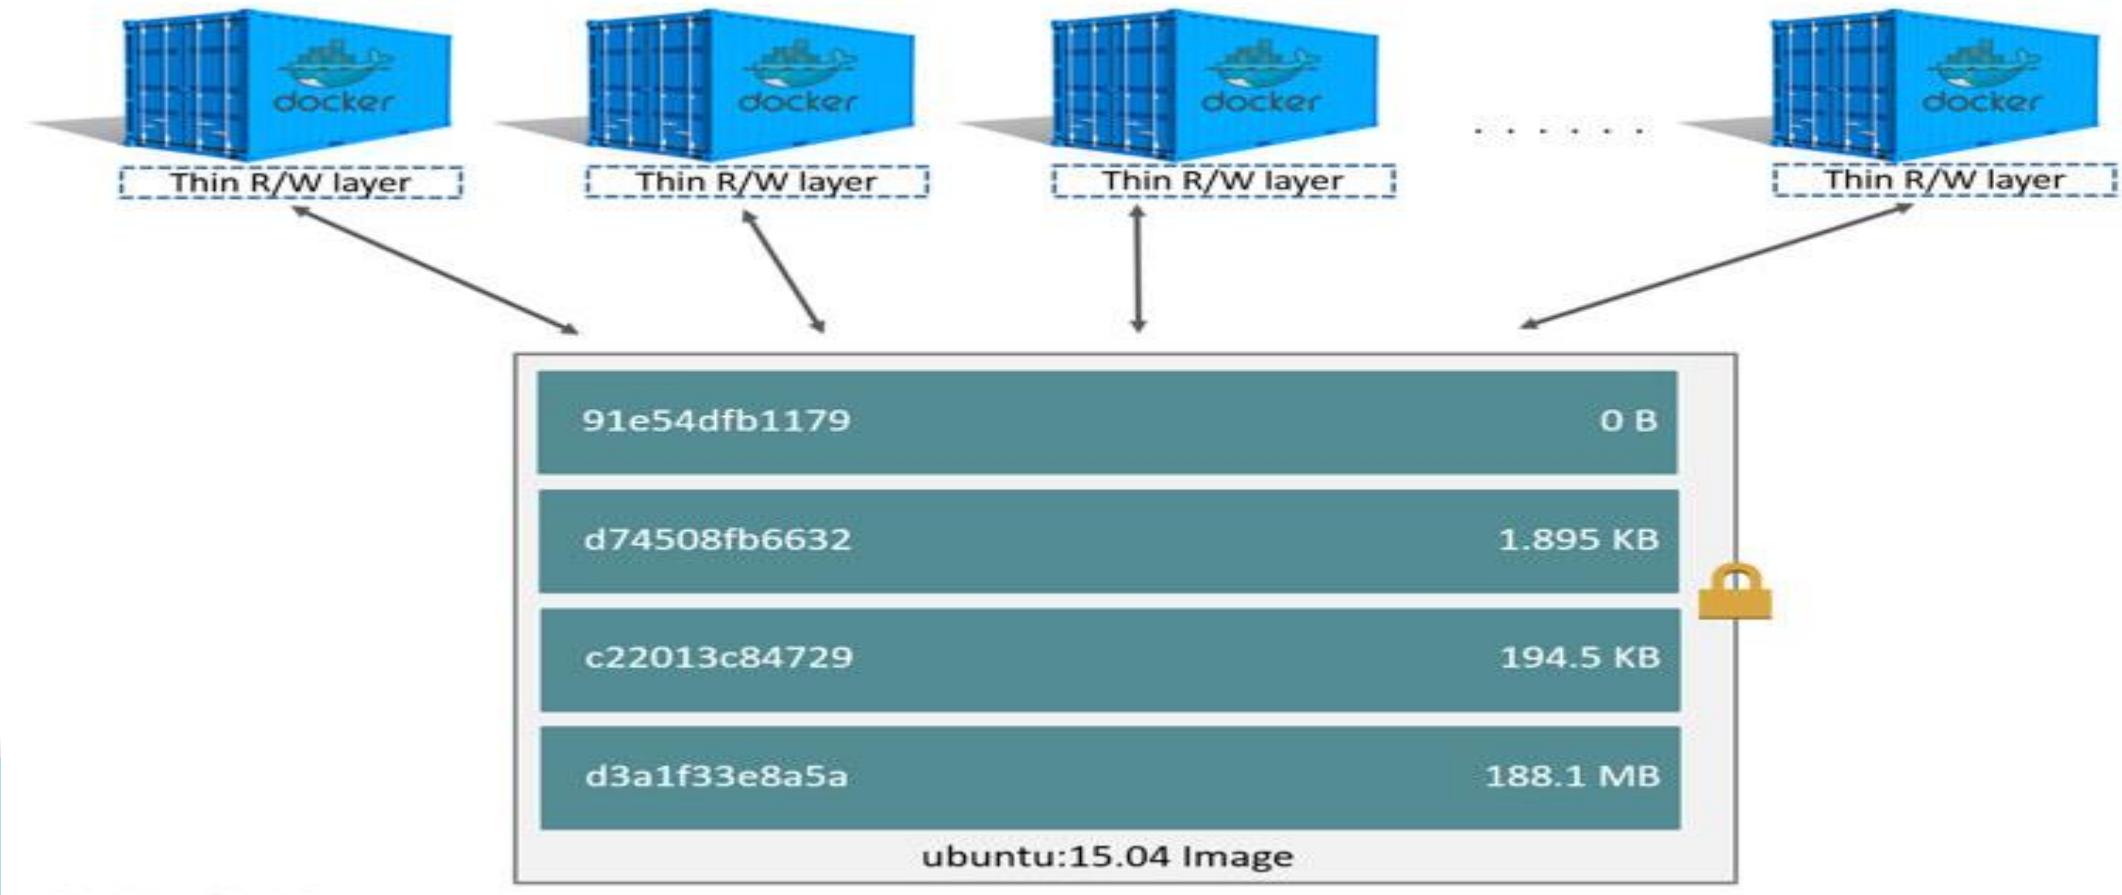
What is an image?

- An image is a collection of files + some meta data.
- Images are made of *layers*, *conceptually stacked on top of each other*.
- Each layer can add, change, and remove files.
- Images can share layers to optimize disk usage, transfer times, and memory use.

Container vs Image

- An image is a read-only filesystem.
- A container is an encapsulated set of processes running in a read-write copy of that filesystem.
- docker run starts a container from a given image.

Container vs Image



Store & manage images

- Images can be stored:
 - On your Docker host.
 - In a Docker registry.
- You can use the Docker client to download (pull) or upload (push) images.
- To be more accurate: you can use the Docker client to tell a Docker server to push and pull images to and from a registry.
- Lets explore docker public registry called “docker hub”

Showing current images

- Let's look at what images are on our host now.

docker images

docker image list

Searching for images

- We cannot list all images on a remote registry, but we can search for a specific keyword:

docker search zookeeper

- "Stars" indicate the popularity of the image.

Downloading images

- There are two ways to download images.
 - Explicitly, with “docker pull”.
 - Implicitly, when executing “docker run” and the image is not found locally.
- Pulling an image.

```
docker pull debian:jessie
```

- Images can have tags.
- Tags define image versions or variants.
- “docker pull ubuntu” will refer to “ubuntu:latest”.
- The :latest tag is generally updated often.

Docker Image LifeCycle

- To list docker images: docker images
- To remove docker images locally: docker rmi <image-name>
- Special Cases:
- Lets try to remove docker images used by current container: docker rmi <image-name>
- This time you will get an error, but the image can be removed forcefully using “-f” option.
- Docker rmi -f <image-name> (Note: It wont work with Image ID)

Docker Image Information

- Docker Image details:
- docker image list
- Detailed Infromation about docker image:
- ls -lrt /var/lib/docker/image/overlay2/imagedb/content/sha256
- High level certificate sign information (SHA):
- cat /var/lib/docker/image/overlay2/repositories.json
- All details with command line:
- docker image inspect <image-name>
- Even you can check all of the information about the docker image:
- docker history <image-id>

Docker Image and Container

LifeCycle

- To list docker: docker ps -a
- To list docker images: docker images
- To remove container: docker rm <container-id> or <Name>
- To remove multiple container: docker rm `docker ps -a -q`
- To remove docker images locally: docker rmi <image-name>

Docker Image and Container LifeCycle

- Special Cases:
- Lets try to remove docker images used by current container: `docker rmi <image-name>`
- This time you will get an error, but the image can be removed forcefully using “-f” option.
- `Docker rmi -f <image-name>` (Note: It wont work with Image ID)
- The best part is though you have removed the image forcefully, but this will not impact the current container as the container preserves the metadata in containers folder.
- Your containers are safe!.

Session: 7

Building Images

Building Images Interactively

- Let's have a Use Case:
 - We will build an image that has httpd.
 - First, we will do it manually with docker commit.
 - Then, we will use a Dockerfile and “docker build”.

Create a new container

- Let's start from base image "centos":

```
docker run -dit --name c1 centos:7
```

```
docker attach c1
```

```
yum update -y
```

```
yum install -y httpd
```

```
exit
```

- Inspect the changes:

```
docker diff <yourContainerId>
```

- Commit the changes:

```
docker commit -m "added httpd and updated" -a "raman khanna" c1 ramacentosimage:v1
```

Dockerfile overview

- A Dockerfile is a build recipe for a Docker image.
- It contains a series of instructions telling Docker how an image is constructed.
- The “docker build” command builds an image from a Dockerfile.

First Dockerfile

- Create a directory to hold our Dockerfile.

```
mkdir myimage
```

- Create a Dockerfile inside this directory.

```
cd myimage  
vi Dockerfile
```

- Write below in our Dockerfile

```
FROM centos:7  
RUN yum update -y  
RUN yum -y install httpd
```

First Dockerfile

- “FROM” indicates the base image for our build.
- Each “RUN” line will be executed by Docker during the build.
- No input can be provided to Docker during the build.

First Dockerfile

- Build the Dockerfile:

docker build -t httpd .

Or

docker build -t httpd /home/ubuntu/raman/myimage/

- -t indicates the tag to apply to the image.
- . indicates the location of the Directory of Dockerfile.

Run & Tag the image

- Let's run the new images:

```
docker run -it <newImageId>
```

```
rpm -qa | grep -i httpd
```

Using Image & viewing history

- The history command lists all the layers composing an image.
- For each layer, it shows its creation time, size, and creation command.
- When an image was built with a Dockerfile, each layer corresponds to a line of the Dockerfile.

```
docker history httpd
```

Dockerfile

- Dockerfile:

```
FROM centos:7
MAINTAINER Raman Khanna raman.khanna@TechLanders.com
RUN mkdir /data
RUN yum update -y
RUN yum -y install httpd php
RUN echo " TechLanders Solutions Deals in DevOps and Cloud" > /var/www/html/index.html
EXPOSE 80
VOLUME /data
RUN echo "httpd" >> /root/.bashrc
CMD ["/bin/bash"]
```

- Build the image:

```
docker build -t webapp:v1 .
```

```
docker run -dit --name c1 -p 8080:80 webapp:v1
```

```
curl 172.31.84.13:8080
```

Browse in browser as well

COPY Instruction

- For Use Case, let's build a container that copy file from localhost
echo " TechLanders Solutions Deals in DevOps and Cloud " > /home/ubuntu/image/index.html

Dockerfile content

```
FROM centos:7
RUN yum update -y
RUN yum install -y httpd
COPY ./index.html /var/www/html/index.html
EXPOSE 80
WORKDIR /var/www/html
CMD ["httpd","-D","FOREGROUND"]
```

COPY Instruction

- For Use Case, let's build a container that compiles a basic "Hello world" program in C.
- hello.c

```
[root@TechLanders yogesh]# cat hello.c
#include<stdio.h>
int main () {
    puts("Hello, TechLanders!");
    return 0;
}
```

- Dockerfile

```
[root@TechLanders yogesh]# cat Dockerfile
FROM ubuntu
RUN apt-get update
RUN apt-get install -y build-essential
COPY hello.c /
RUN make hello
CMD /hello
```

Note: Using COPY keyword we can copy the files from Docker Host to a container.

COPY Instruction

- docker build -t helloworld:v1 .
- docker run -it --name prodv4 helloworld:v1 /bin/bash
- root@519a9d815a29:/# ls -lrt /hello
- -rwxr-xr-x. 1 root root 8600 Sep 18 11:54 /hello
- root@519a9d815a29:/# ./hello
- Hello, world!

ENTRYPOINT vs. CMD

CMD

Specify the command to run at container startup

```
CMD ["somecommand.sh"]
```

```
> docker run myimage
```

Running somecommand.sh

...

ENTRYPOINT

```
ENTRYPOINT ["somecommand.sh"]
```

```
> docker run myimage
```

Running somecommand.sh

...



Dockerfile >

```
FROM debian:buster  
  
COPY . /myproject  
  
RUN apt-get update ...  
  
CMD ["./cmd1.sh"]
```

> "docker run my-image"

Exec ==> ./cmd1.sh

> "docker run my-image cmd2.sh"

Exec ==> cmd2.sh



Dockerfile >

```
FROM debian:buster
```

```
COPY . /myproject
```

```
RUN apt-get update ...
```

```
ENTRYPOINT ["entrypoint1.sh"]
```

```
CMD ["param1", "param2"] 
```

> "docker run my-container"

Exec ==>

```
entrypoint1.sh param1 param2
```

> "docker run my-container cmd2"

Exec ==>

```
entrypoint1.sh cmd2
```

Docker Restart Policy

```
$ docker run -dit --restart unless-stopped centos
```

Flag	Description
no	Do not automatically restart the container. (the default)
on-failure	Restart the container if it exits due to an error, which manifests as a non-zero exit code.
always	Always restart the container if it stops. If it is manually stopped, it is restarted only when Docker daemon restarts or the container itself is manually restarted. (See the second bullet listed in restart policy details)
unless-stopped	Similar to always, except that when the container is stopped (manually or otherwise), it is not restarted even after Docker daemon restarts.

Docker Registries

Local Registry (Local to Host)

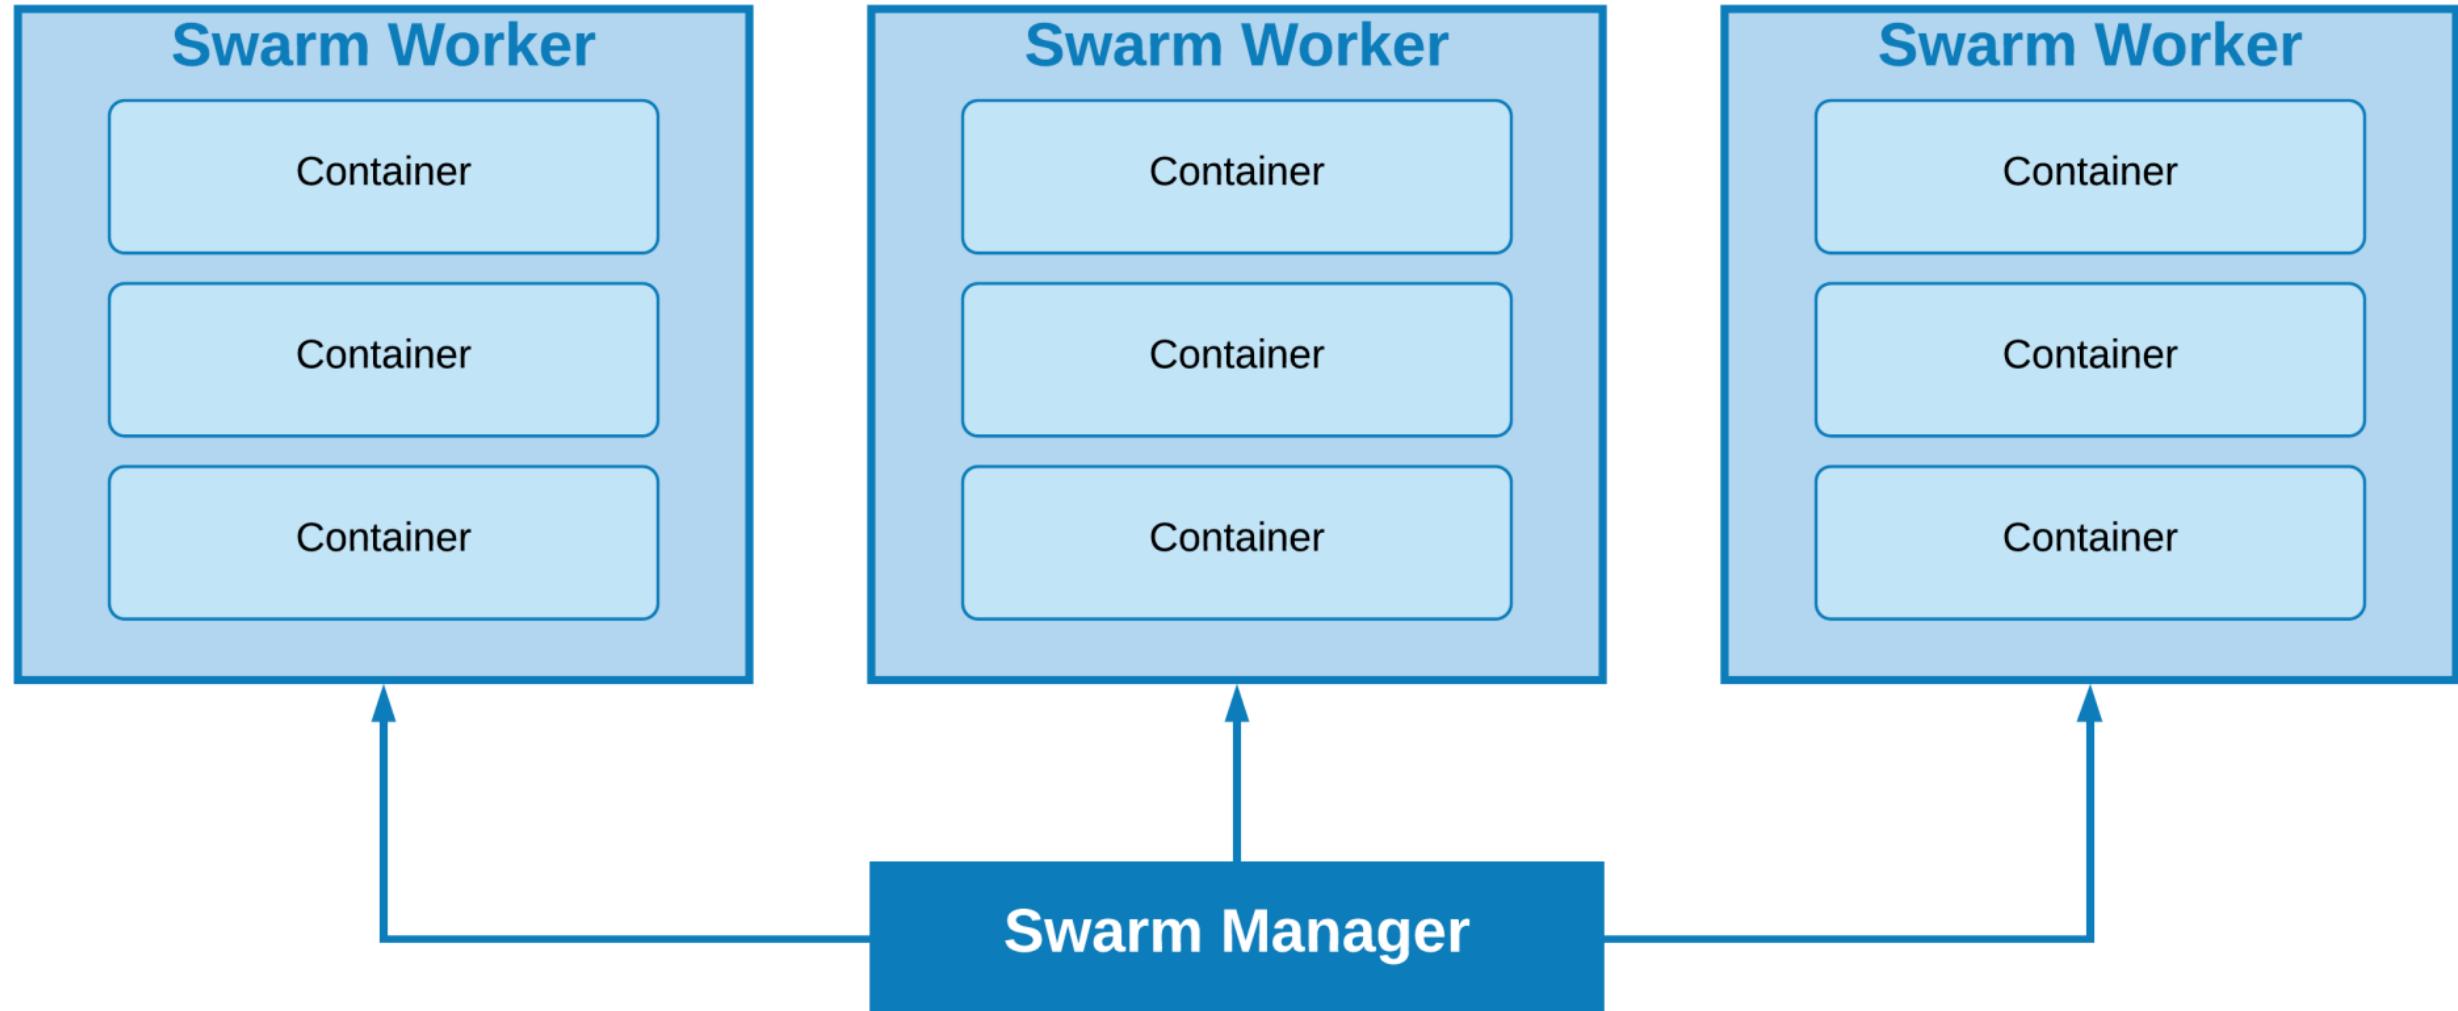
Remote Registry (Private)

Global Registry (Public)

Docker Trusted Registry

- Enterprise-grade image storage solution from Docker
- Highly Secure
- Image and job management with CICD
- HA Availability
- Efficiency with Near to user storage and bandwidth sharing
- Security Scanning
- Similar tools in market Sonatype Nexus (open source), AWS ECR (PaaS), azure Container Registry, GCP container Registry etc

Docker Swarm



Docker Limitations

- Hardware Issues? High Availability?
- How IP address will be managed for failover?
- Scaling?
- Auto Healing?
- Autoscaling?
- No Application Management - Only Containerization
- Updation of application/management

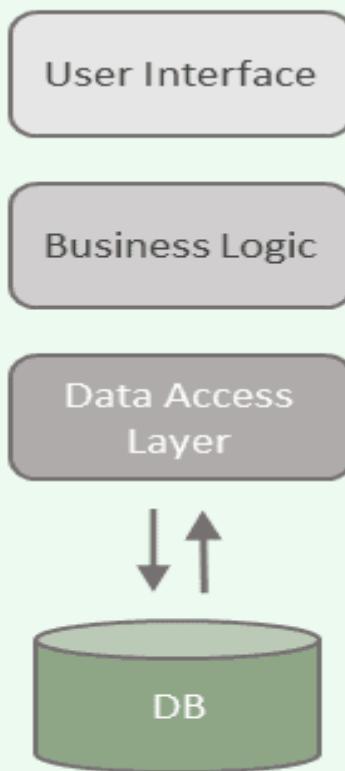
Why Kubernetes

- Kubernetes can schedule and run application containers on clusters of physical or virtual machines.
- **host-centric** infrastructure to a **container-centric** infrastructure.
- Orchestrator
- Load balancing
- Auto Scaling
- Application Health checks
- Rolling updates

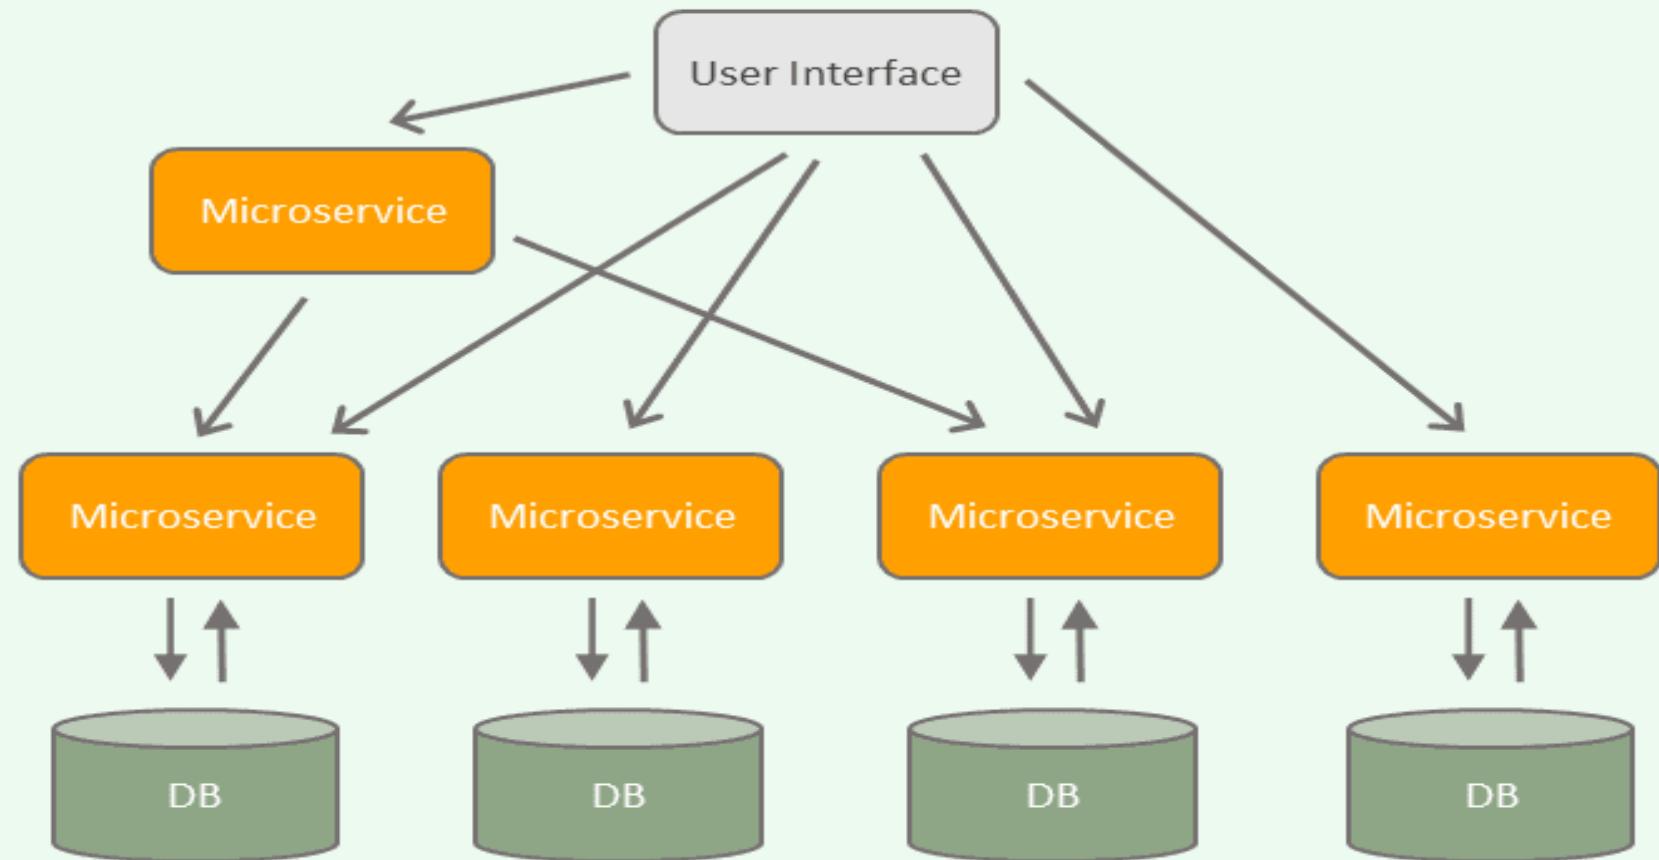
Questions & Answers



MONOLITHIC ARCHITECTURE



MICROSERVICES ARCHITECTURE



Container Orchestration

Raman Khanna

Containers Limitation?

High Availability?

Overlay Network?

Versioning of Application – Rollout, Rollback?

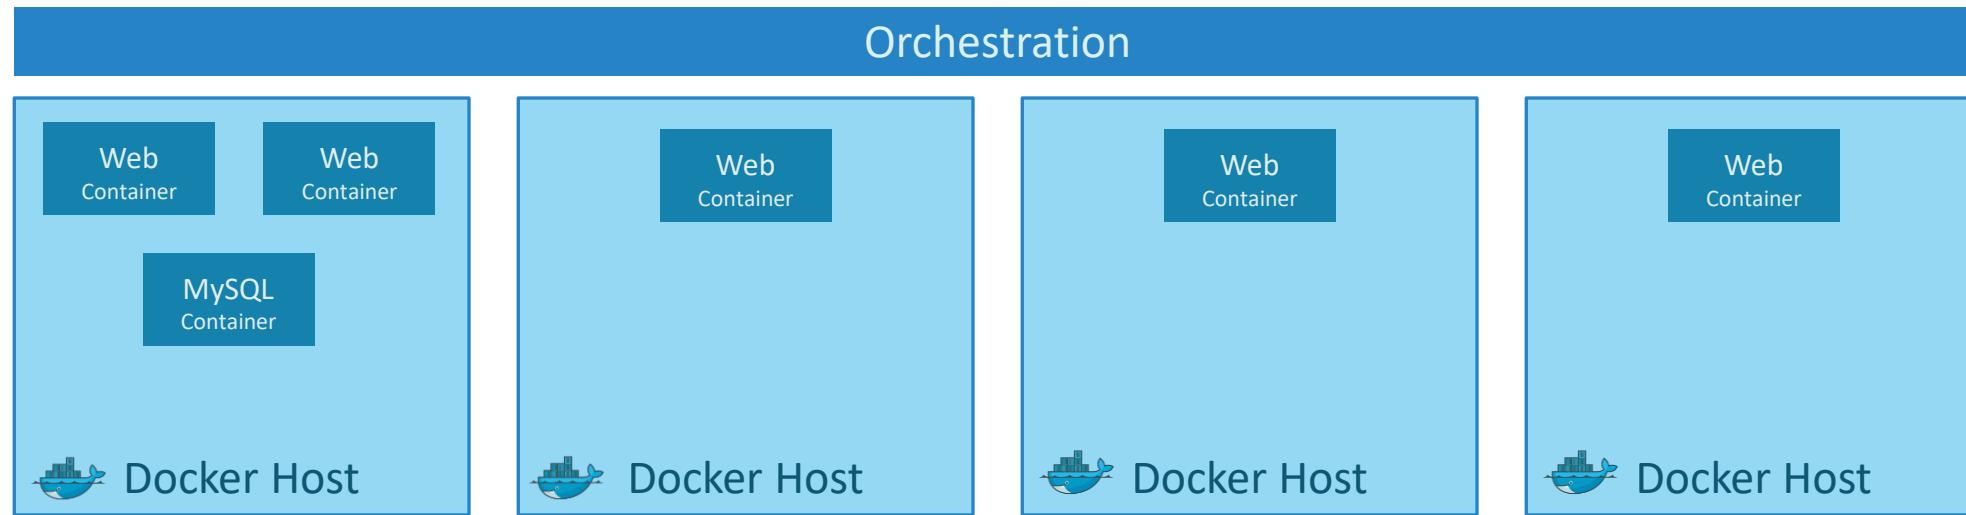
Scaling?

Autoscaling?

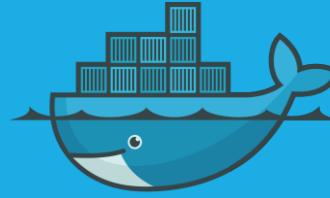
Monitoring?

Dependency between containers?

Container orchestration



Orchestration Technologies



Docker Swarm



kubernetes



MESOS

What is Kubernetes?

The Kubernetes project was started by Google in 2014.

Kubernetes builds upon a decade and a half of experience that Google has with running production workloads at scale.

Kubernetes can run on a range of platforms, from your laptop, to VMs on a cloud provider, to rack of bare metal servers.

Kubernetes is an open-source platform for automating deployment, scaling, and operations of application containers across clusters of hosts, providing container-centric infrastructure.

portable: with all public, private, hybrid, community cloud

self-healing: auto-placement, auto-restart, auto-replication, auto-scaling

Why Kubernetes

Kubernetes can schedule and run application containers on clusters of physical or virtual machines.

host-centric infrastructure to a **container-centric** infrastructure.

Orchestrator

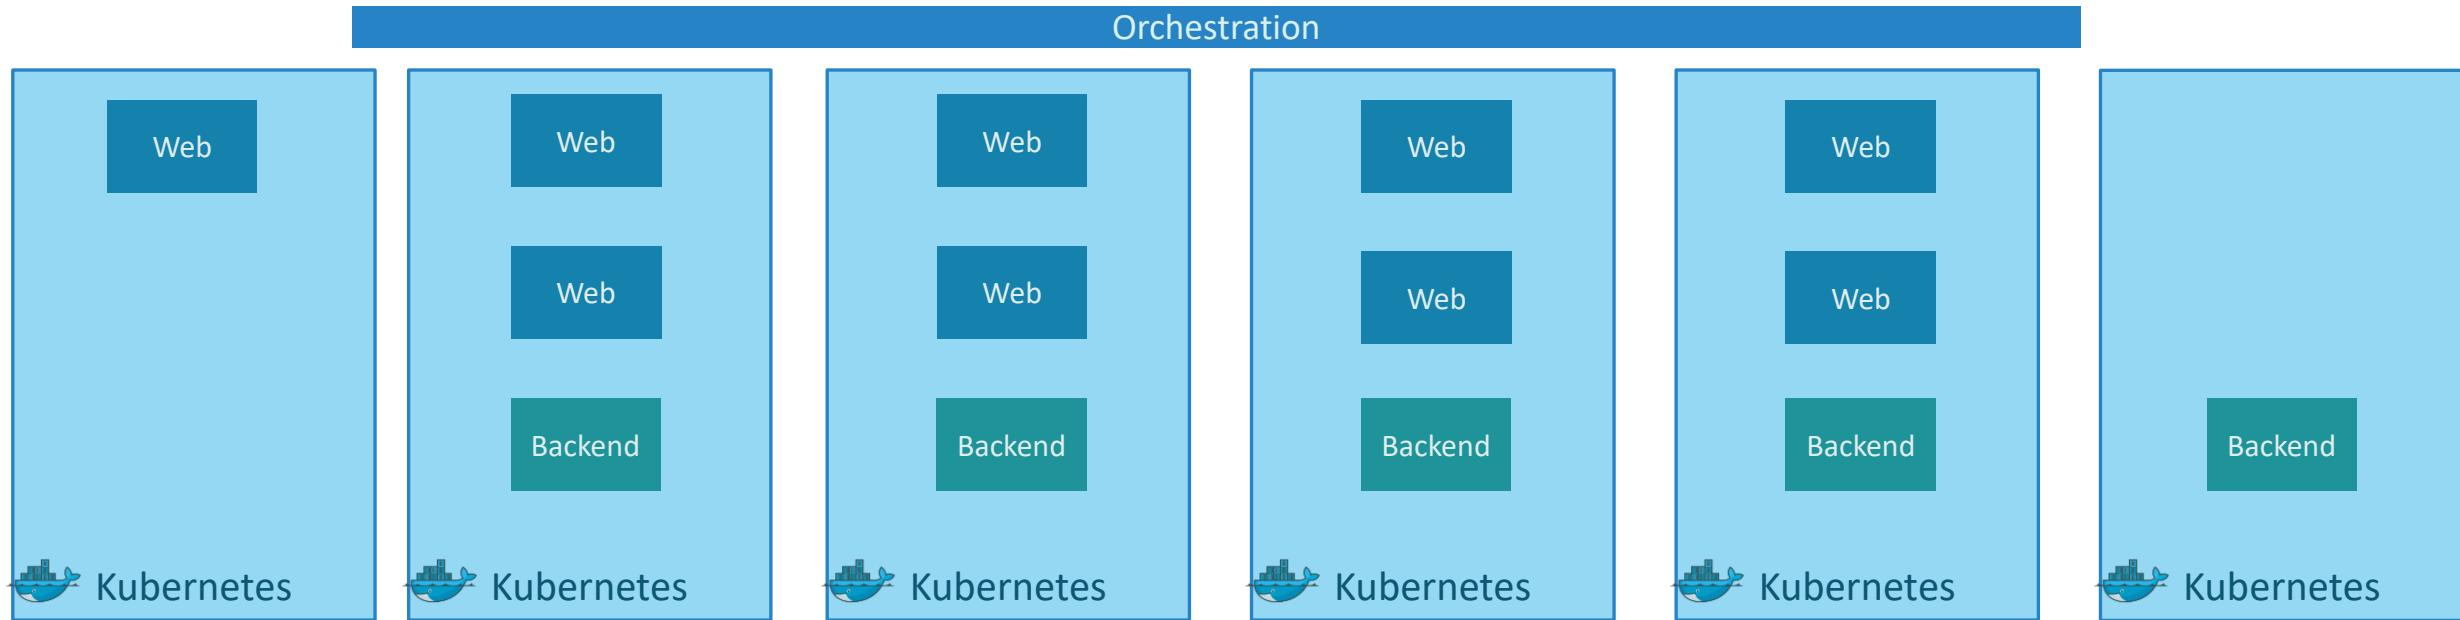
Load balancing

Auto Scaling

Application Health checks

Rolling updates

Kubernetes Advantage



And that is **kubernetes..**

Setup

Raman



Minikube



Kubeadm



Google Cloud Platform



Amazon Web Services

play-with-k8s.com

Setup Kubernetes

Setup - kubeadm

Raman

Kubernetes Cluster

A Kubernetes cluster consists of two types of resources:

Master: Which coordinates with the cluster

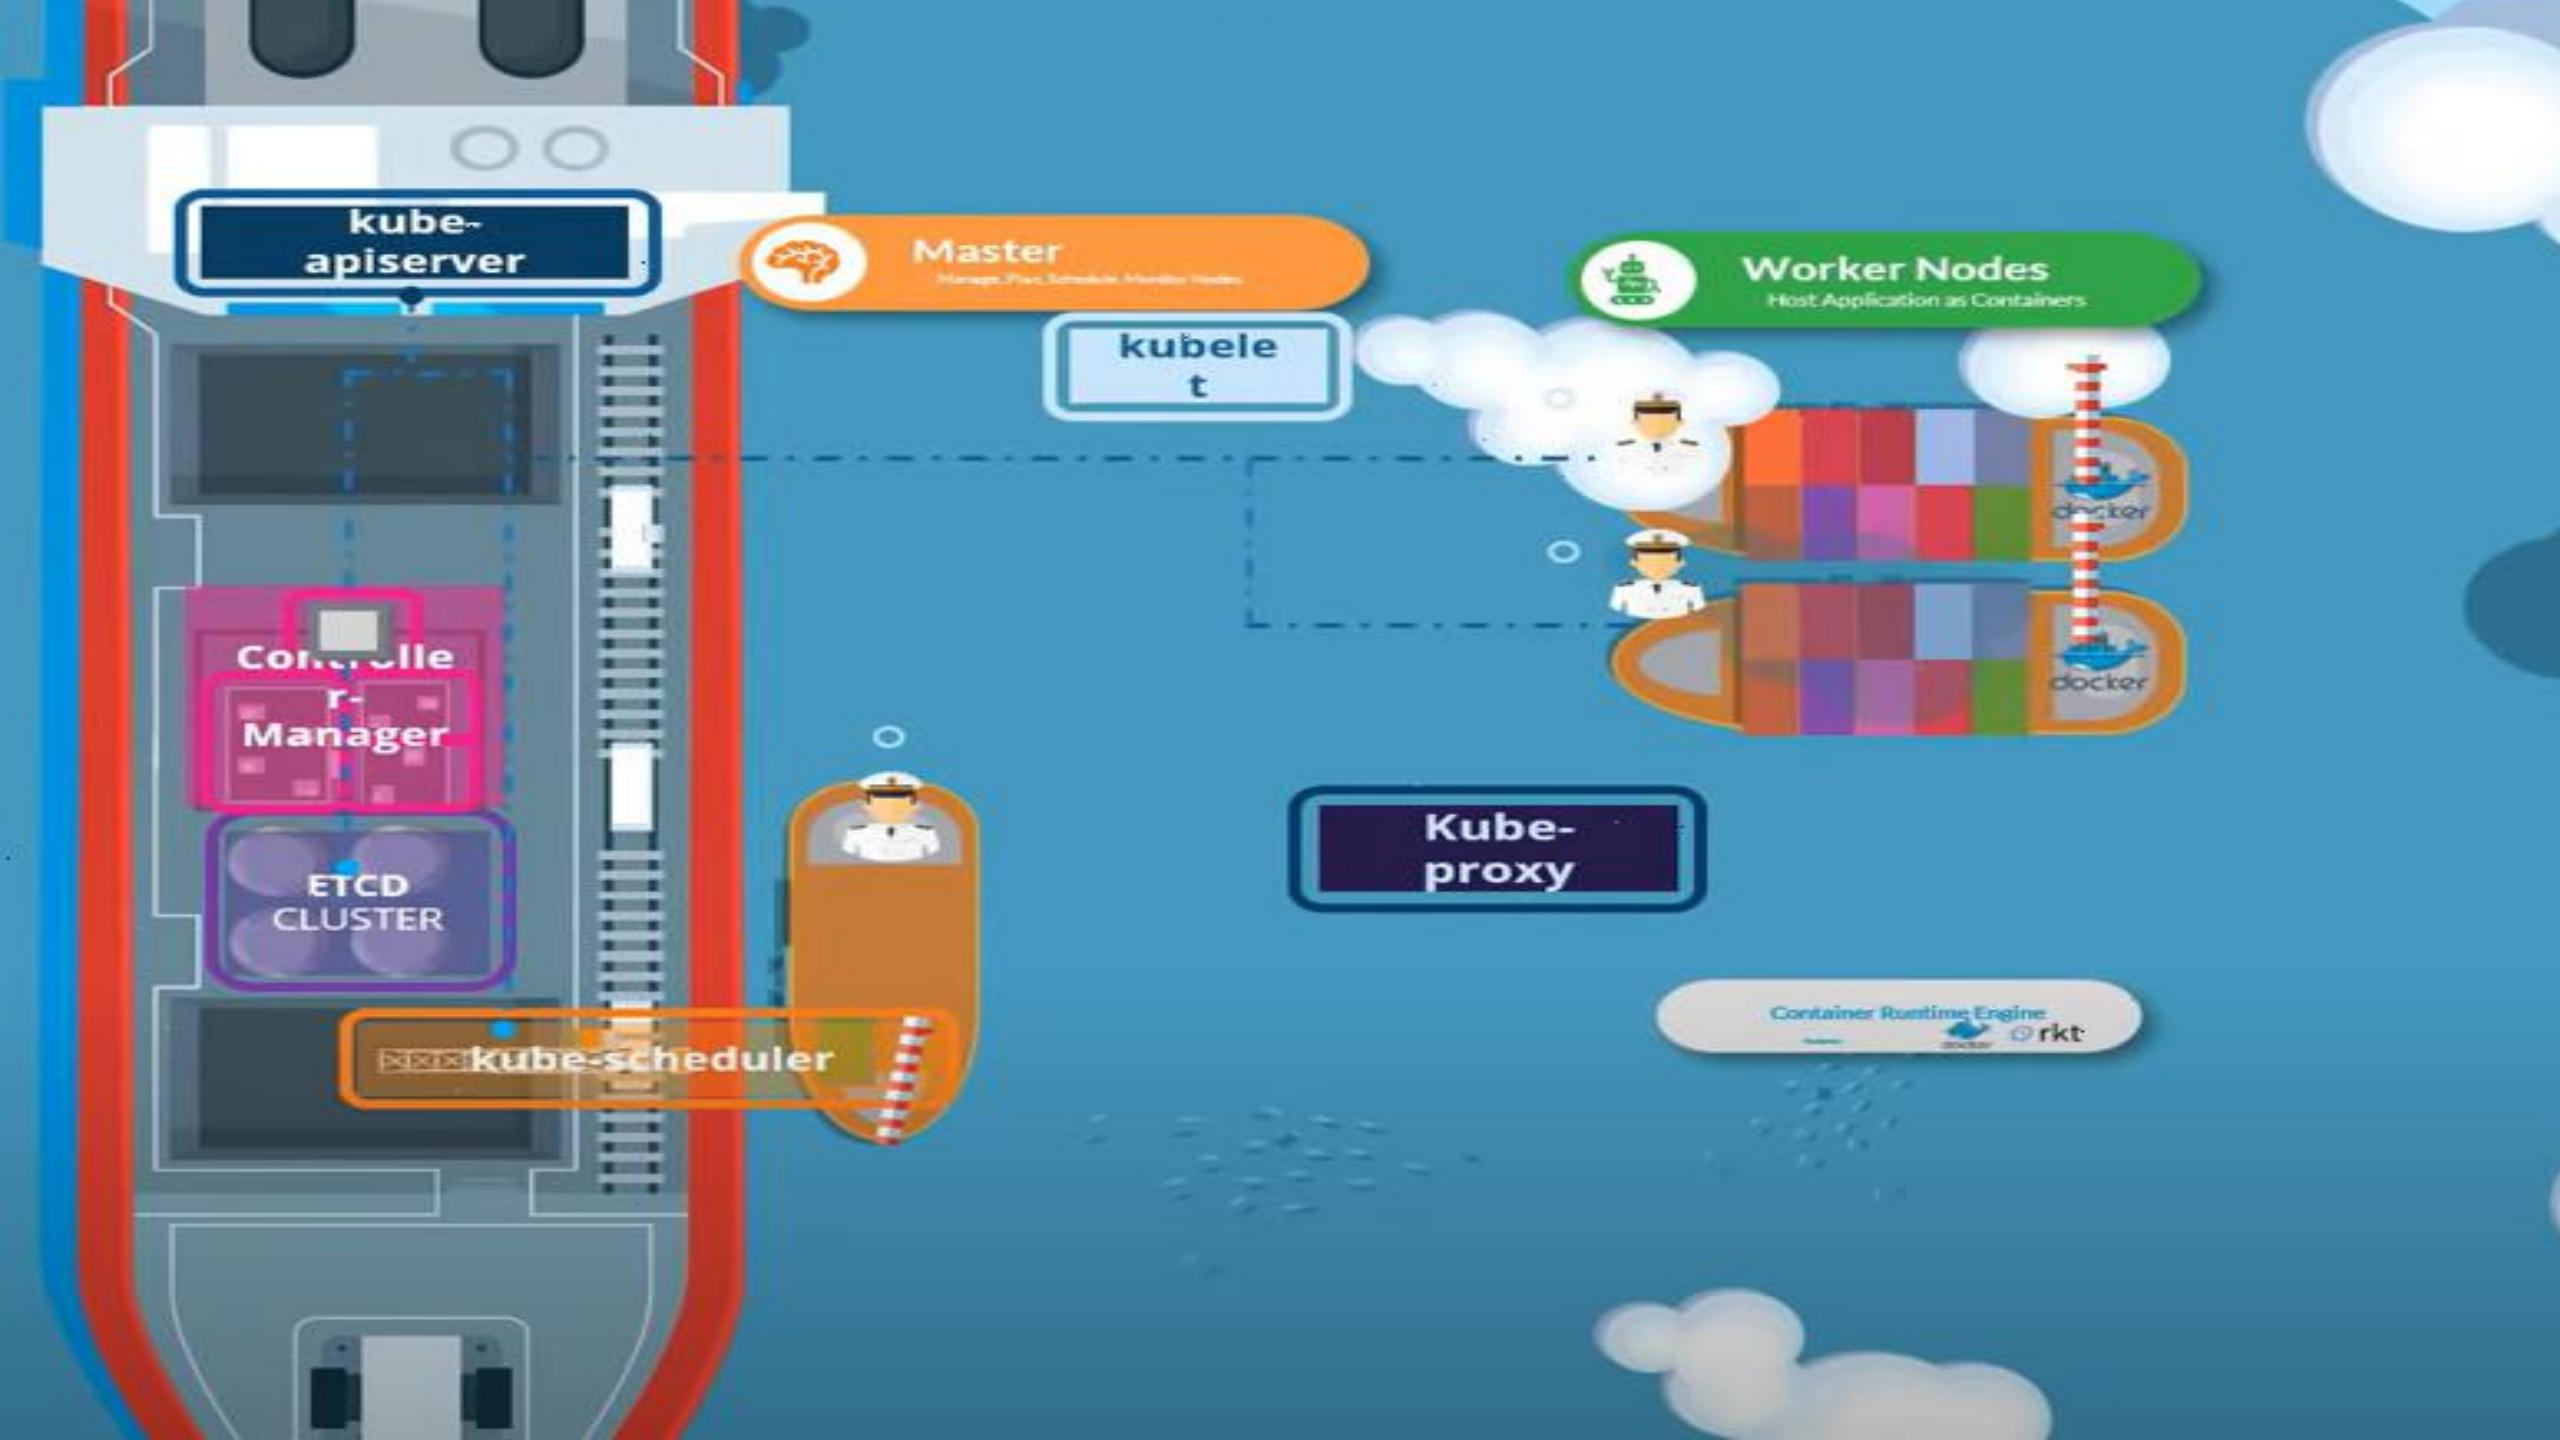
The Master is responsible for managing the cluster. The master coordinates all activities in your cluster, such as scheduling applications, maintaining applications' desired state, scaling applications, and rolling out new updates.

Nodes: Are the workers that run application

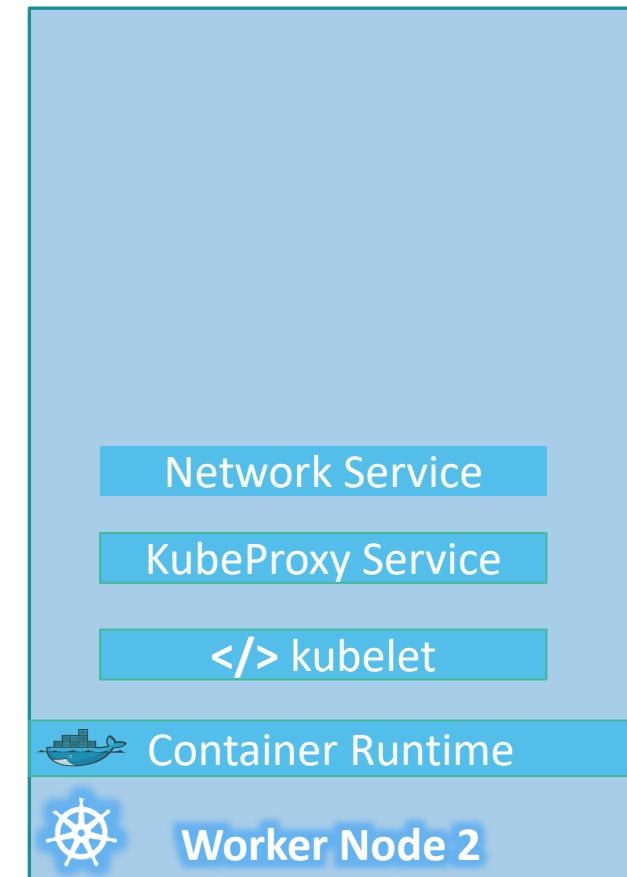
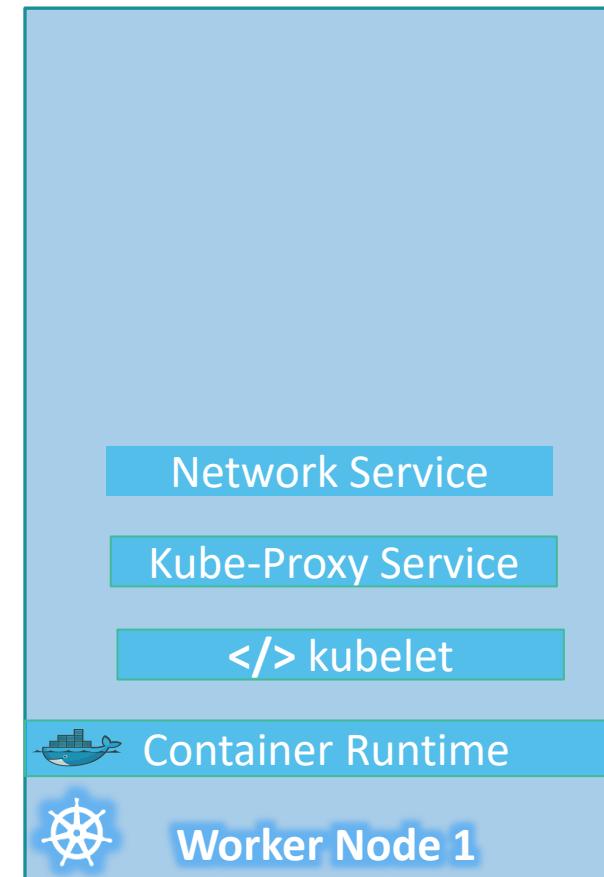
A node is a VM or a physical computer that serves as a worker machine in a Kubernetes cluster.

Masters manage the cluster and the nodes are used to host the running applications.

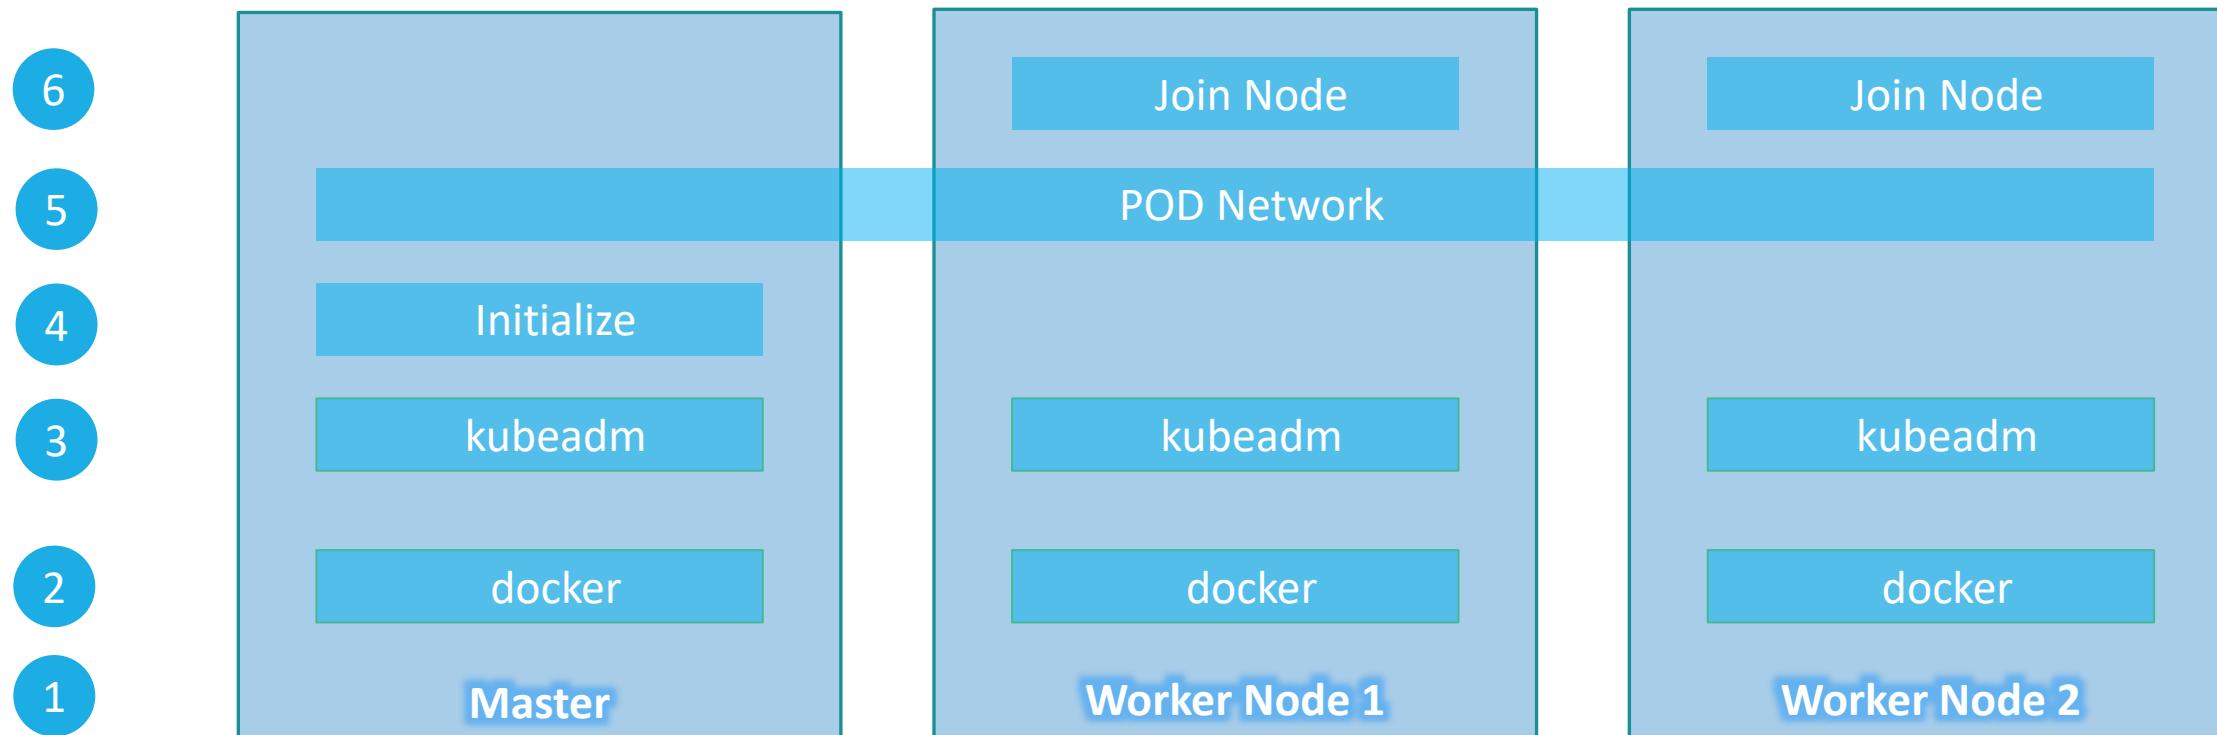
The nodes communicate with the master using the Kubernetes API, which the master exposes.



kubeadm



Steps



POD

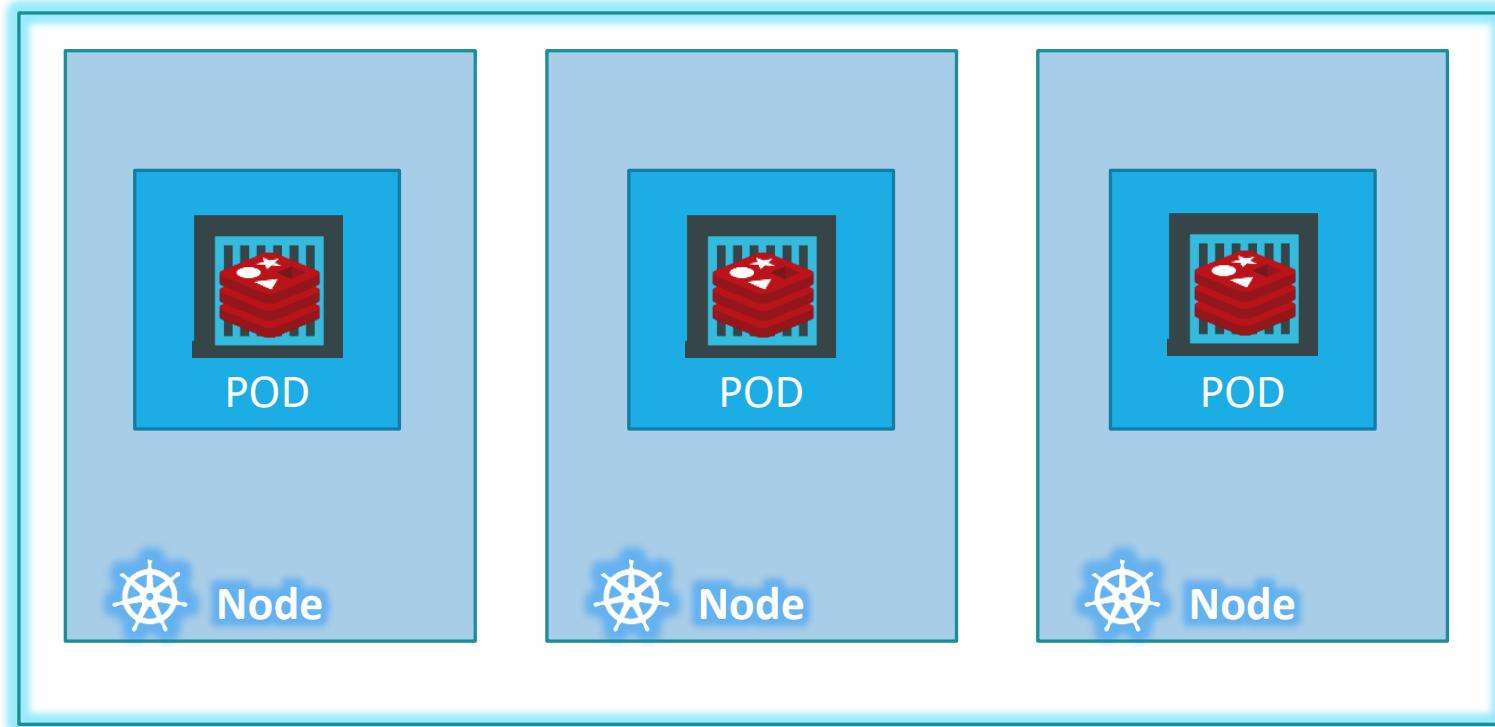
Raman

Assumptions

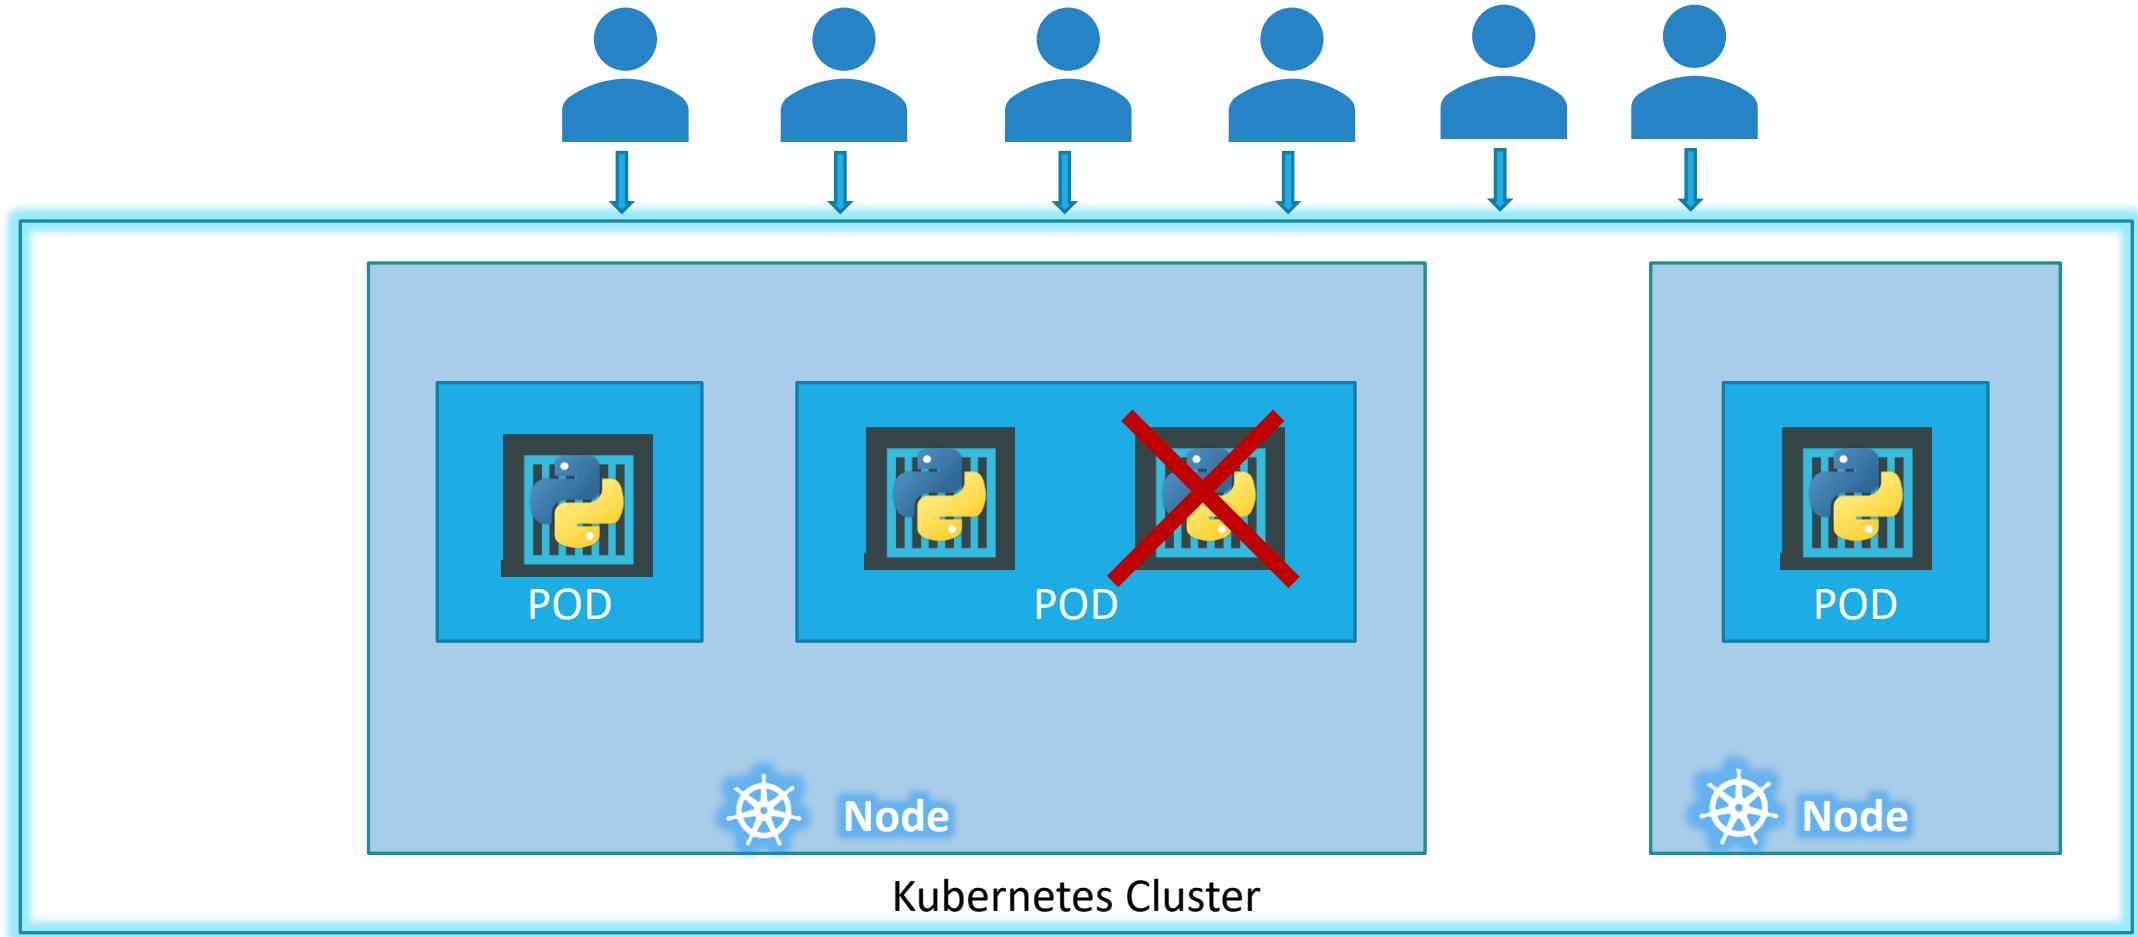
Docker Image

Kubernetes Cluster

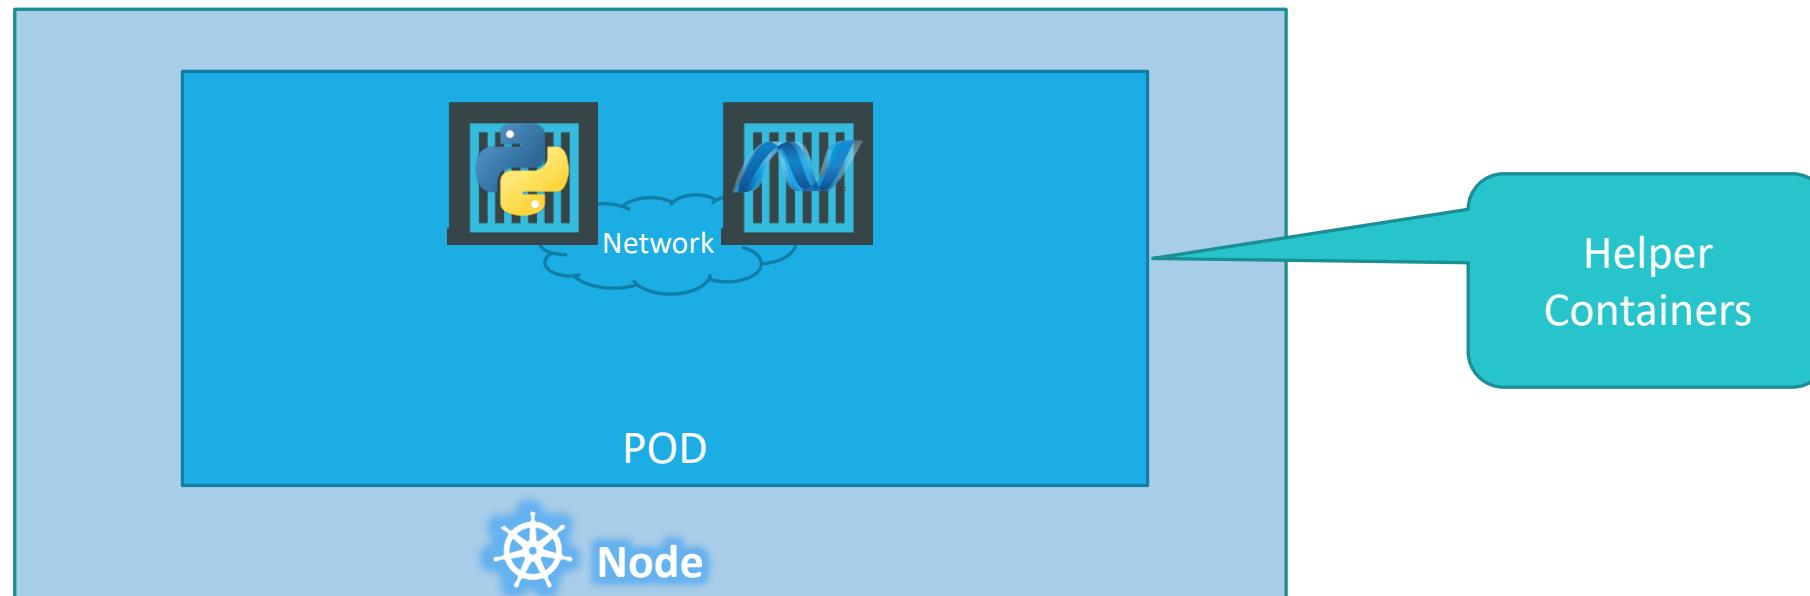
POD



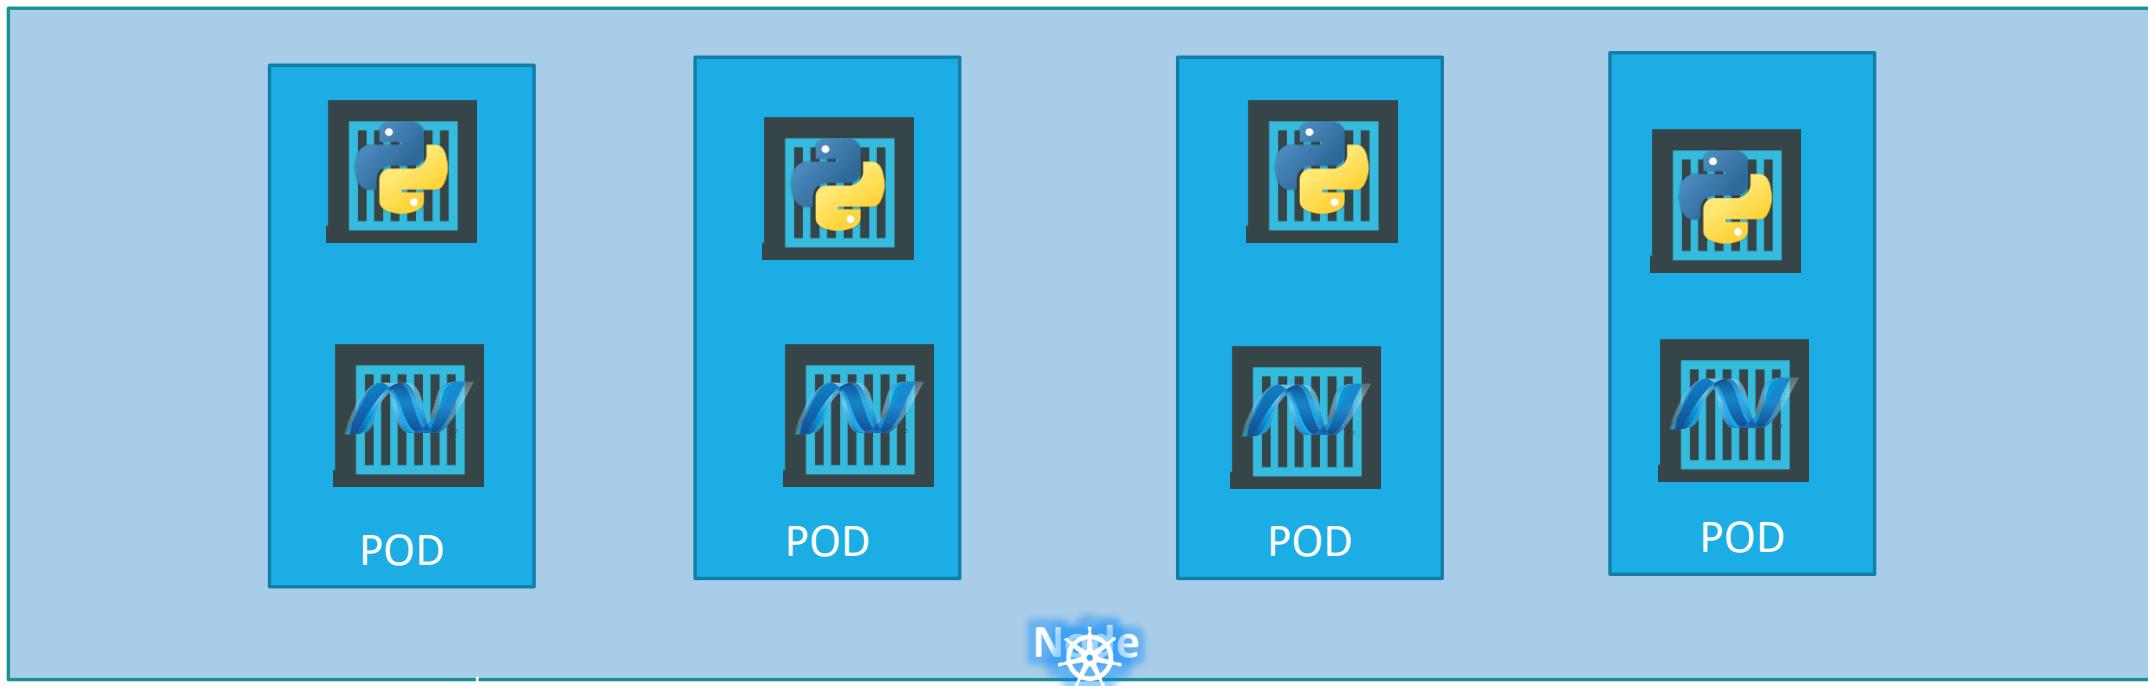
POD



Multi-Container PODs



PODs Again!



Note: I am avoiding networking and load balancing details to keep explanation simple.



kubectl

```
kubectl run nginx --image nginx
```

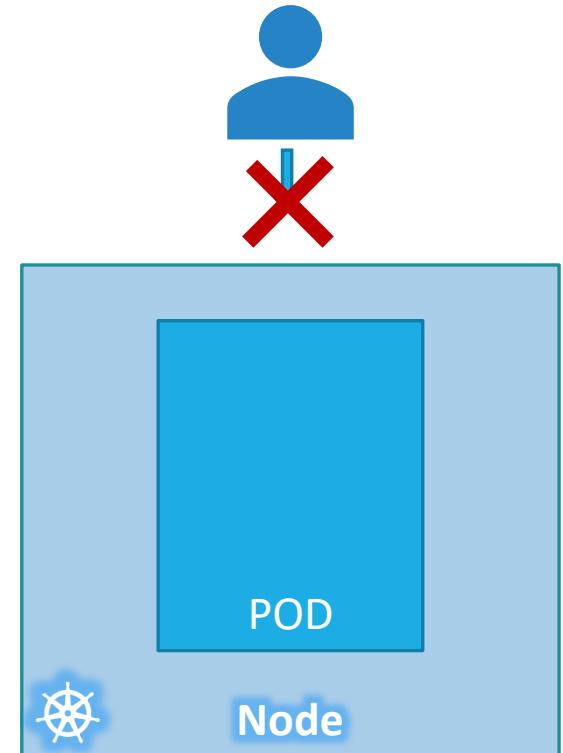
```
kubectl get pods
```

```
C:\Kubernetes>kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-8586cf59-whssr	0/1	ContainerCreating	0	3s

```
C:\Kubernetes>kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-8586cf59-whssr	1/1	Running	0	8s



YAML Introduction

Raman

POD

With YAML

YAML in Kubernetes

```
pod-definition.yml
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
spec:
  containers:
    - name: nginx-container
      image: nginx
```

1st Item in List

```
kubectl create -f pod-definition.yml
```



Kind	Version
POD	v1
Service	v1
ReplicaSet	apps/v1
Deployment	apps/v1

Commands

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myapp-pod	1/1	Running	0	20s

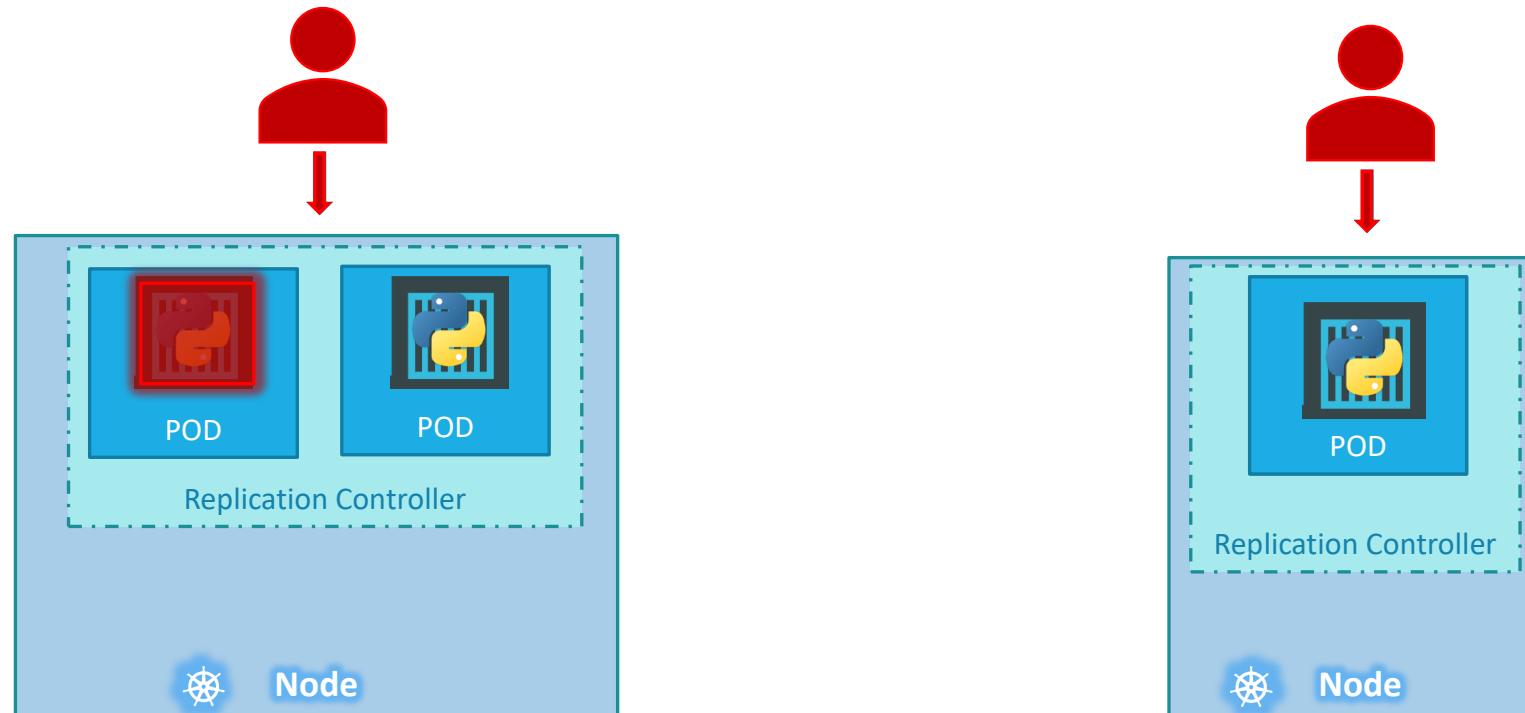
```
> kubectl describe pod myapp-pod
```

```
Name:           myapp-pod
Namespace:      default
Node:          minikube/192.168.99.100
Start Time:    Sat, 03 Mar 2018 14:26:14 +0800
Labels:         app=myapp
                name=myapp-pod
Annotations:   <none>
Status:        Running
IP:            10.244.0.24
Containers:
  nginx:
    Container ID:  docker://830bb56c8c42a86b4bb70e9c1488fae1bc38663e4918b6c2f5a783e7688b8c9d
    Image:          nginx
    Image ID:      docker-pullable://nginx@sha256:4771d09578c7c6a65299e110b3ee1c0a2592f5ea2618d23e4ffe7a4cab1ce5de
    Port:          <none>
    State:         Running
      Started:    Sat, 03 Mar 2018 14:26:21 +0800
    Ready:         True
    Restart Count: 0
    Environment:  <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from default-token-x95w7 (ro)
Conditions:
  Type      Status
  Initialized  True
  Ready       True
  PodScheduled  True
Events:
  Type  Reason          Age   From           Message
  ----  -----          --   --            --
  Normal Scheduled      34s   default-scheduler  Successfully assigned myapp-pod to minikube
  Normal SuccessfulMountVolume 33s   kubelet, minikube  MountVolume.SetUp succeeded for volume "default-token-x95w7"
  Normal Pulling        33s   kubelet, minikube  pulling image "nginx"
  Normal Pulled         27s   kubelet, minikube  Successfully pulled image "nginx"
  Normal Created        27s   kubelet, minikube  Created container
  Normal Started        27s   kubelet, minikube  Started container
```

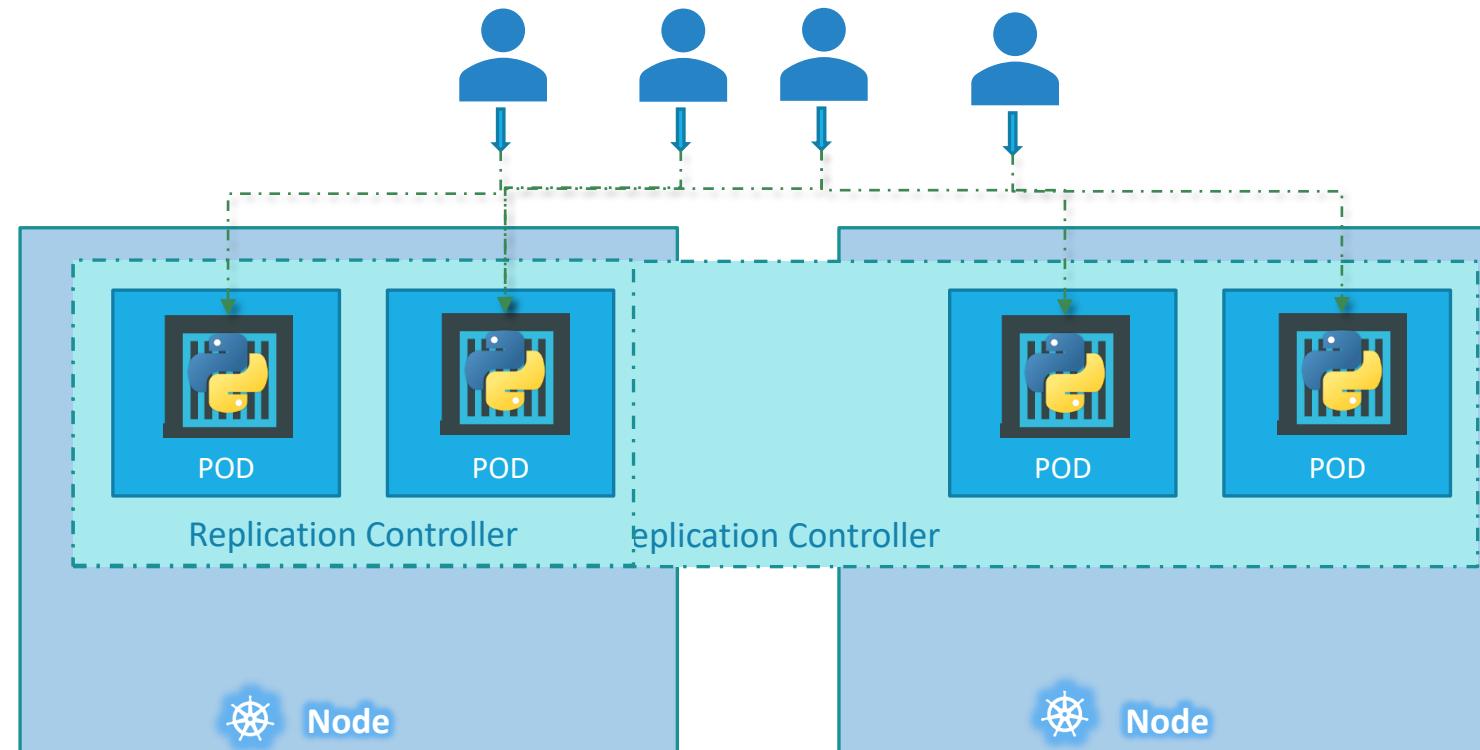
Replication Controller

Raman Khanna

High Availability



Load Balancing & Scaling



Replication Controller

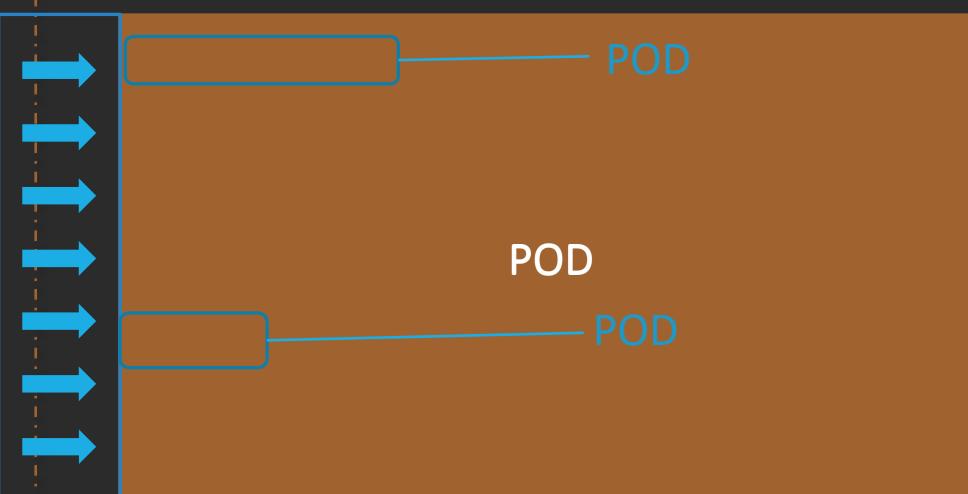
Replica Set

rc-definition.yml

```

apiVersion: v1
kind: ReplicationController
metadata:
  name: myapp-rc
  labels:
    app: myapp
    type: front-end
spec:
  template:
    - replicas: 3
      containers:
        - name: nginx-container
          image: nginx

```



pod-definition.yml

```

apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx

```

```

> kubectl create -f rc-definition.yml
replicationcontroller "myapp-rc" created

```

```

> kubectl get replicationcontroller
NAME      DESIRED   CURRENT   READY   AGE
myapp-rc   3         3         3       19s

```

```

> kubectl get pods

```

NAME	READY	STATUS	RESTARTS	AGE
myapp-rc-41vk9	1/1	Running	0	20s
myapp-rc-mc2mf	1/1	Running	0	20s
myapp-rc-px9pz	1/1	Running	0	20s

replicaset-definition.yml

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-repl
  labels:
    app: myapp
    type: front-end
spec:
  template:

```

POD

```

  replicas: 3
  selector:
    matchLabels:
      type: front-end

```

pod-definition.yml

```

apiVersion: v1
kind: Pod
error: unable to recognize "replicaset-
definition.yml": no matches for /, Kind=ReplicaSet

```

```

  labels:
    app: myapp
    type: front-end
  spec:
    containers:
      - name: nginx-container
        image: nginx

```

```

> kubectl create -f replicaset-definition.yml
replicaset "myapp-replicaset" created

```

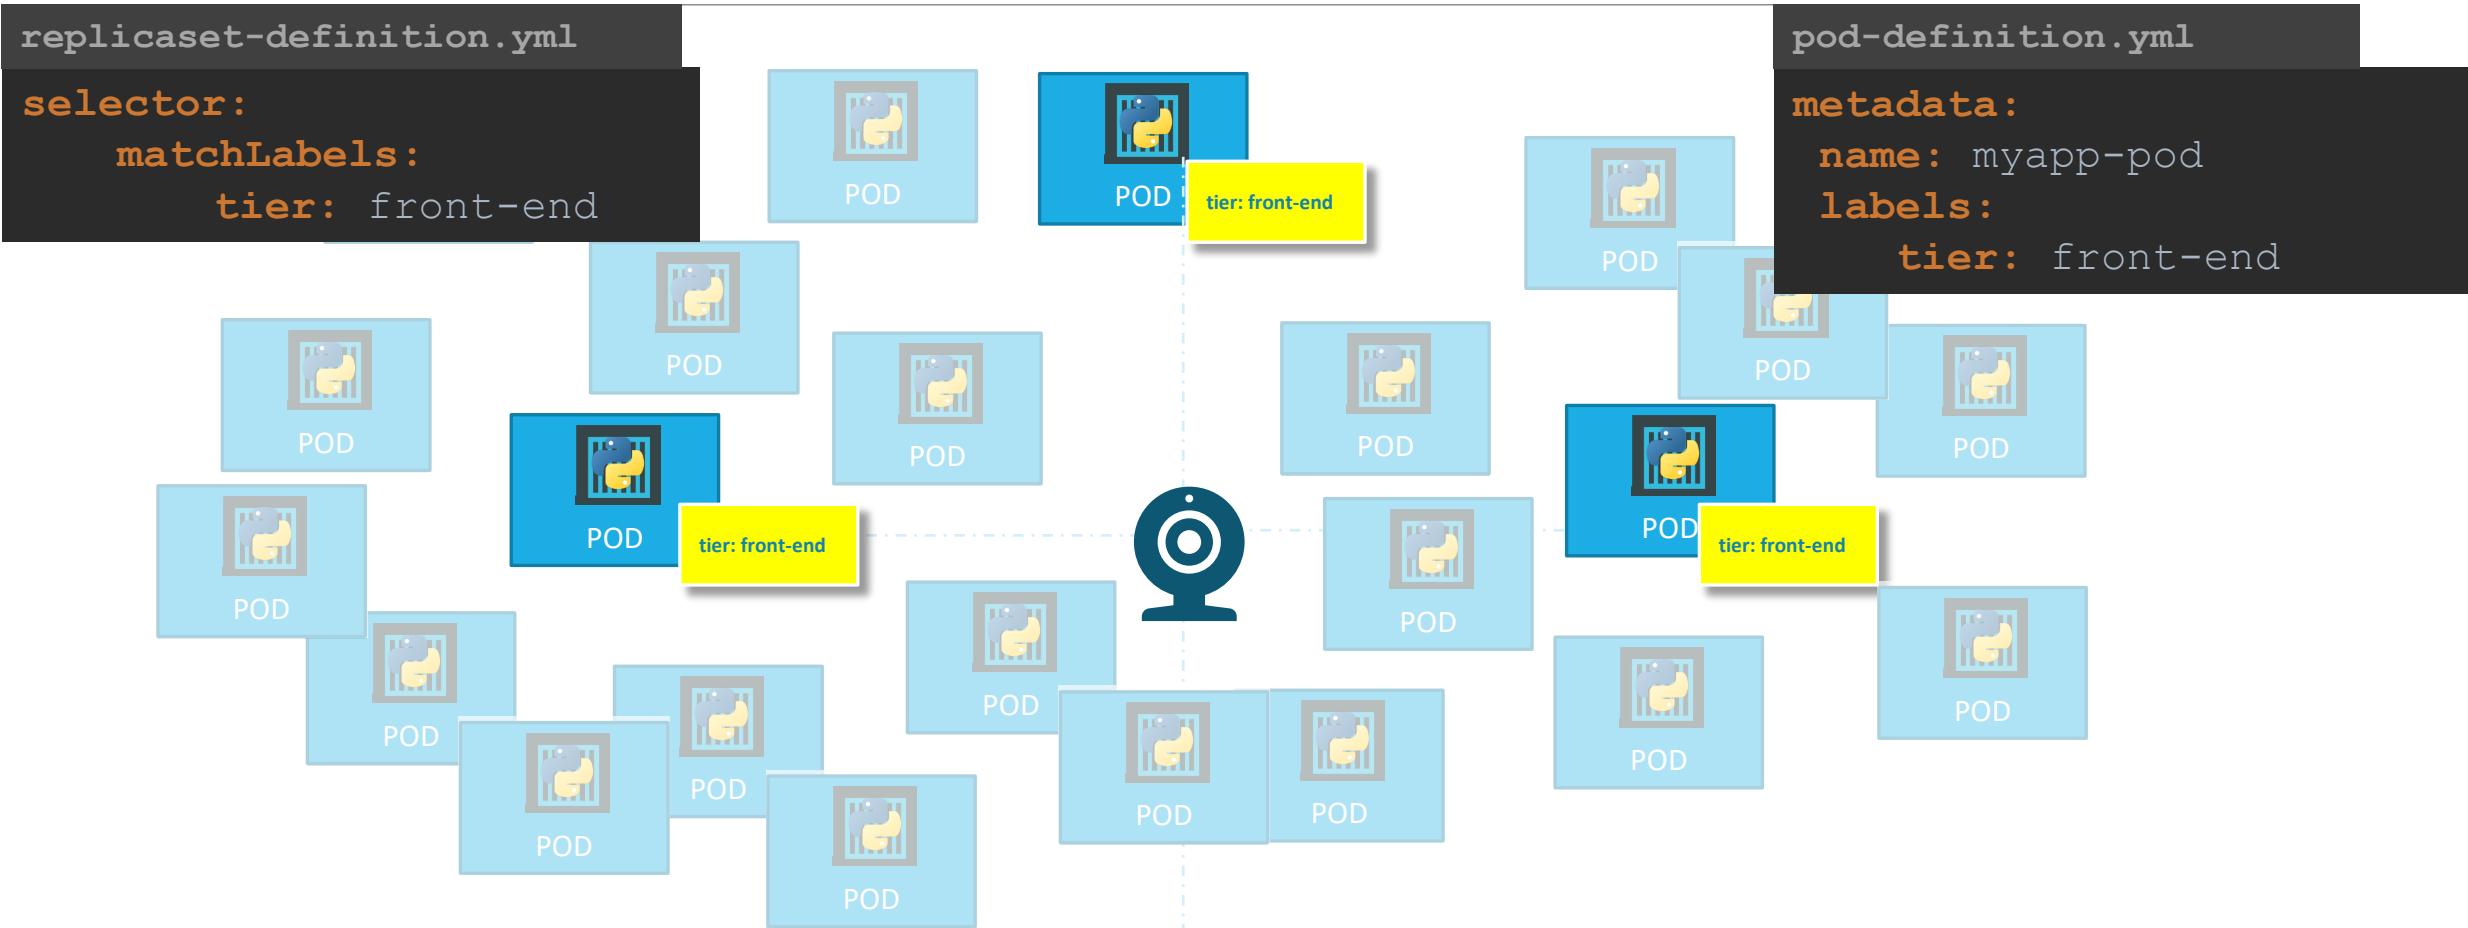
```
> kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
myapp-replicaset	3	3	3	19s

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myapp-replicaset-9dd19	1/1	Running	0	45s
myapp-replicaset-9jtpx	1/1	Running	0	45s
myapp-replicaset-hq84m	1/1	Running	0	45s

Labels and Selectors



replicaset-definition.yml

```

apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      type: front-end

```



Scale

```
> kubectl replace -f replicaset-definition.yml
```

```
> kubectl scale --replicas=6 -f replicaset-definition.yml
```

```
> kubectl scale --replicas=6 replicaset myapp-replicaset
```



```
replicaset-definition.yml
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
      - name: nginx-container
        image: nginx
  replicas: 6
  selector:
    matchLabels:
      type: front-end
```

commands

```
> kubectl create -f replicaset-definition.yml
```

```
> kubectl get replicaset
```

```
> kubectl delete replicaset myapp-replicaset
```

*Also deletes all underlying PODs

```
> kubectl replace -f replicaset-definition.yml
```

```
> kubectl scale -replicas=6 -f replicaset-definition.yml
```

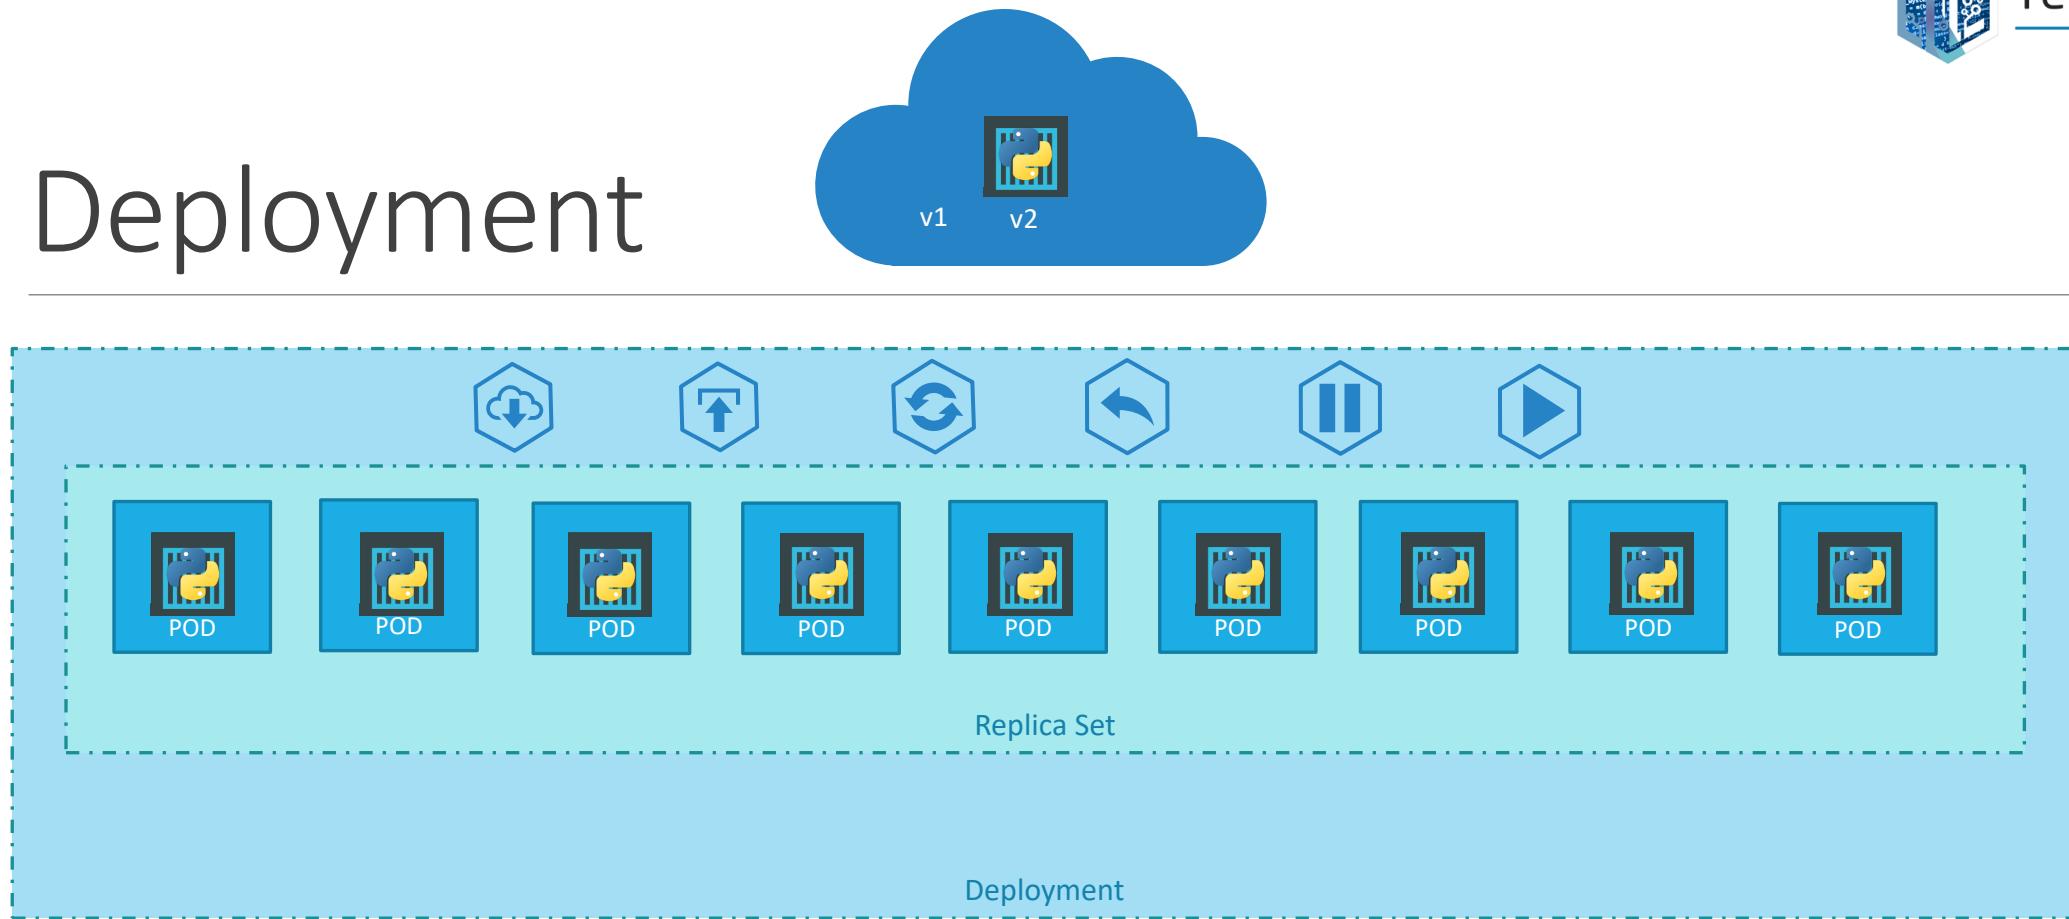
Demo

ReplicaSet

Deployment

Raman Khanna

Deployment



Definition

```
> kubectl create -f deployment-definition.yml
```

```
deployment "myapp-deployment" created
```

```
> kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
myapp-deployment	3	3	3	3	21s

```
> kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
myapp-deployment-6795844b58	3	3	3	2m

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myapp-deployment-6795844b58-5rbjl	1/1	Running	0	2m
myapp-deployment-6795844b58-h4w55	1/1	Running	0	2m
myapp-deployment-6795844b58-lfjhv	1/1	Running	0	2m

```
deployment-definition.yml
```

```
apiVersion: apps/v1
```

```
kind: Replicaset
```

```
metadata:
```

```
  name: myapp-deployment
```

```
  labels:
```

```
    app: myapp
```

```
    type: front-end
```

```
spec:
```

```
  template:
```

```
    metadata:
```

```
      name: myapp-pod
```

```
      labels:
```

```
        app: myapp
```

```
        type: front-end
```

```
  spec:
```

```
    containers:
```

- name: nginx-container

```
    image: nginx
```

```
  replicas: 3
```

```
  selector:
```

```
    matchLabels:
```

```
      type: front-end
```

commands

```
> kubectl get all
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
deploy/myapp-deployment	3	3	3	3	9h
NAME	DESIRED	CURRENT	READY	AGE	
rs/myapp-deployment-6795844b58	3	3	3	9h	
NAME	READY	STATUS	RESTARTS	AGE	
po/myapp-deployment-6795844b58-5rbjl	1/1	Running	0	9h	
po/myapp-deployment-6795844b58-h4w55	1/1	Running	0	9h	
po/myapp-deployment-6795844b58-1fjhv	1/1	Running	0	9h	

Demo

Deployment

Demo

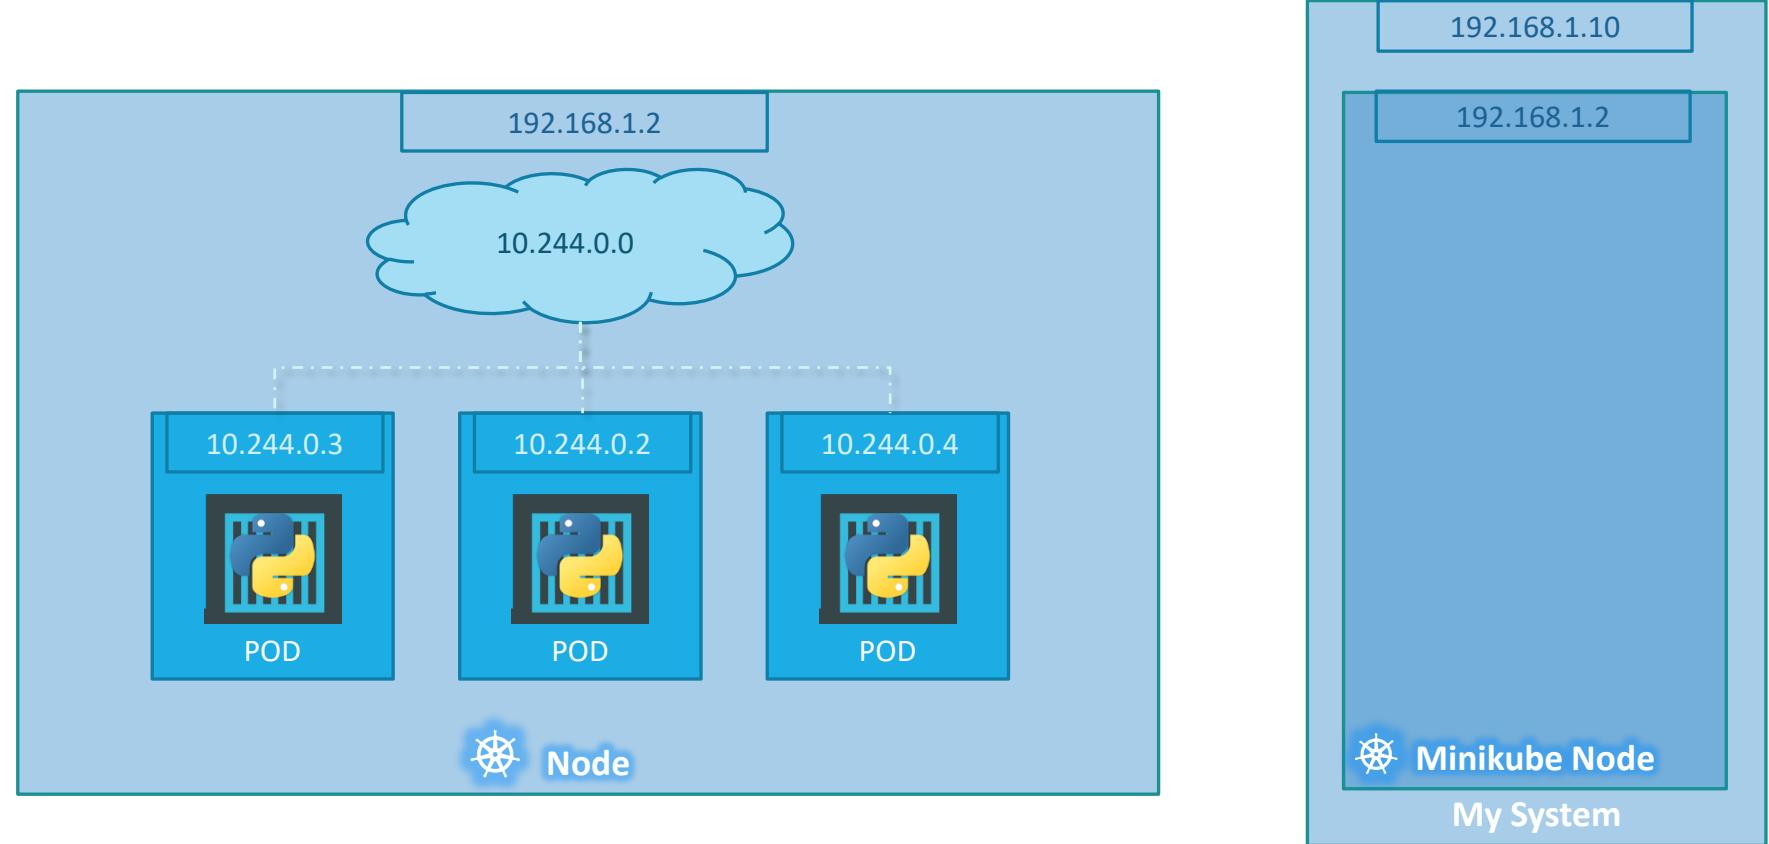
Deployment

Networking 101

Raman Khanna

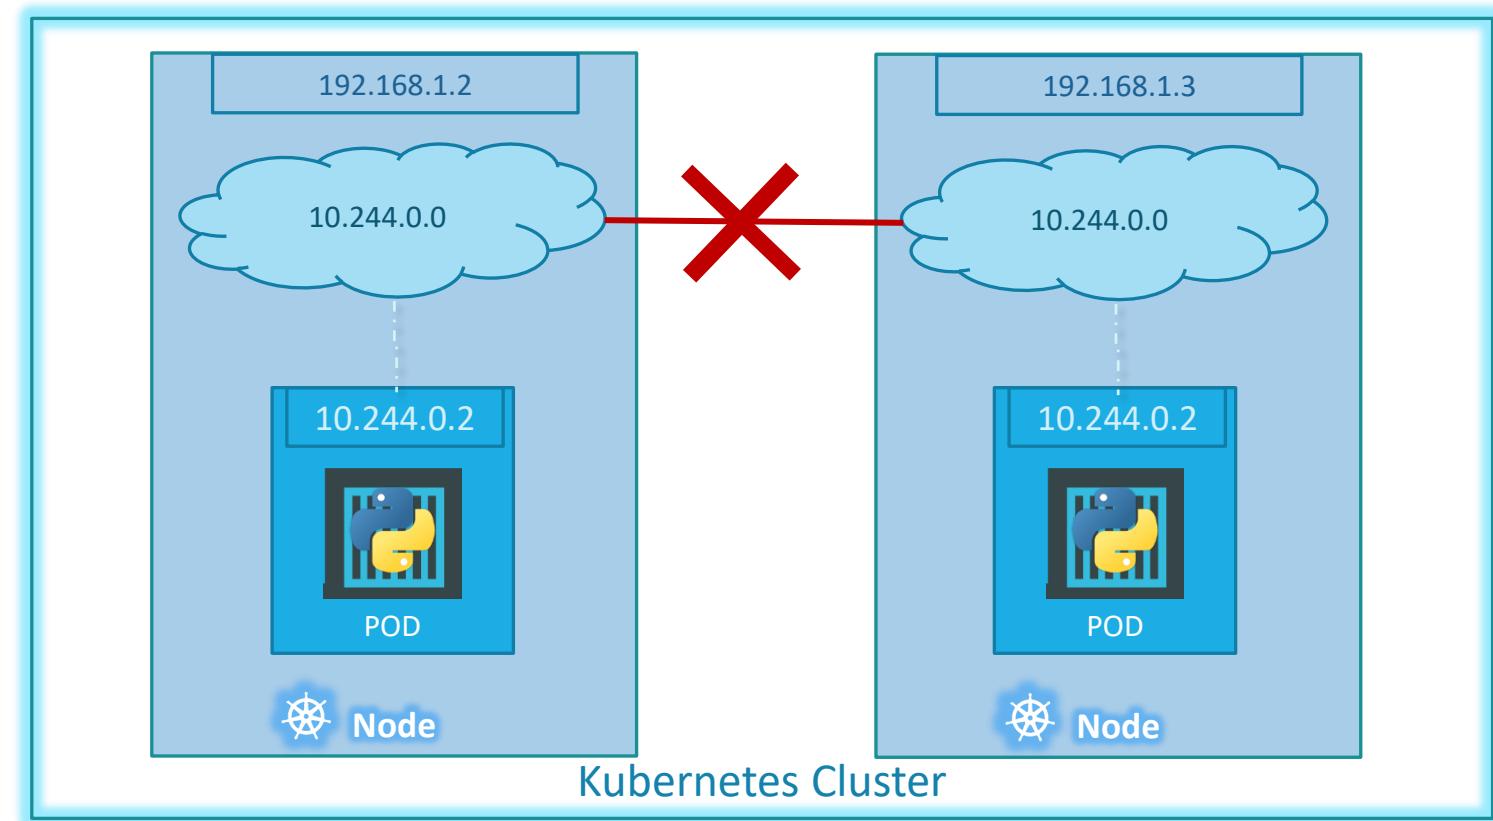
Kubernetes Networking - 101

- IP Address is assigned to a POD



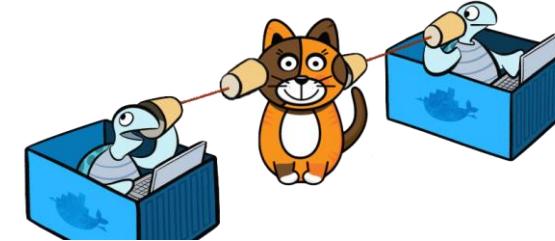
Cluster Networking

- All containers/PODs can communicate to one another without NAT
- All nodes can communicate with all containers and vice-versa without NAT





 flannel vmware
NSX



Cluster Networking Setup

(3/4) Installing a pod **network**

You **MUST** install a pod **network** add-on so that your pods can communicate with each other.

The **network** must be deployed before any applications. Also, kube-dns, an internal helper service, will not start up before a **network** is installed. kubeadm only supports Container Network Interface (CNI) based **networks** (and does not support kubenet).

Several projects provide Kubernetes pod **networks** using CNI, some of which also support [Network Policy](#). See the [add-ons page](#) for a complete list of available **network** add-ons. IPv6 support was added in [CNI v0.6.0](#). [CNI bridge](#) and [local-ipam](#) are the only supported IPv6 **network** plugins in 1.9.

Note: kubeadm sets up a more secure cluster by default and enforces use of [RBAC](#). Please make sure that the **network** manifest of choice supports RBAC.

You can install a pod **network** add-on with the following command:

```
kubectl apply -f <add-on.yaml>
```

NOTE: You can install **only one** pod **network** per cluster.

[Choose one...](#) [Calico](#) [Canal](#) [Flannel](#) [Kube-router](#) [Romana](#) [Weave Net](#)

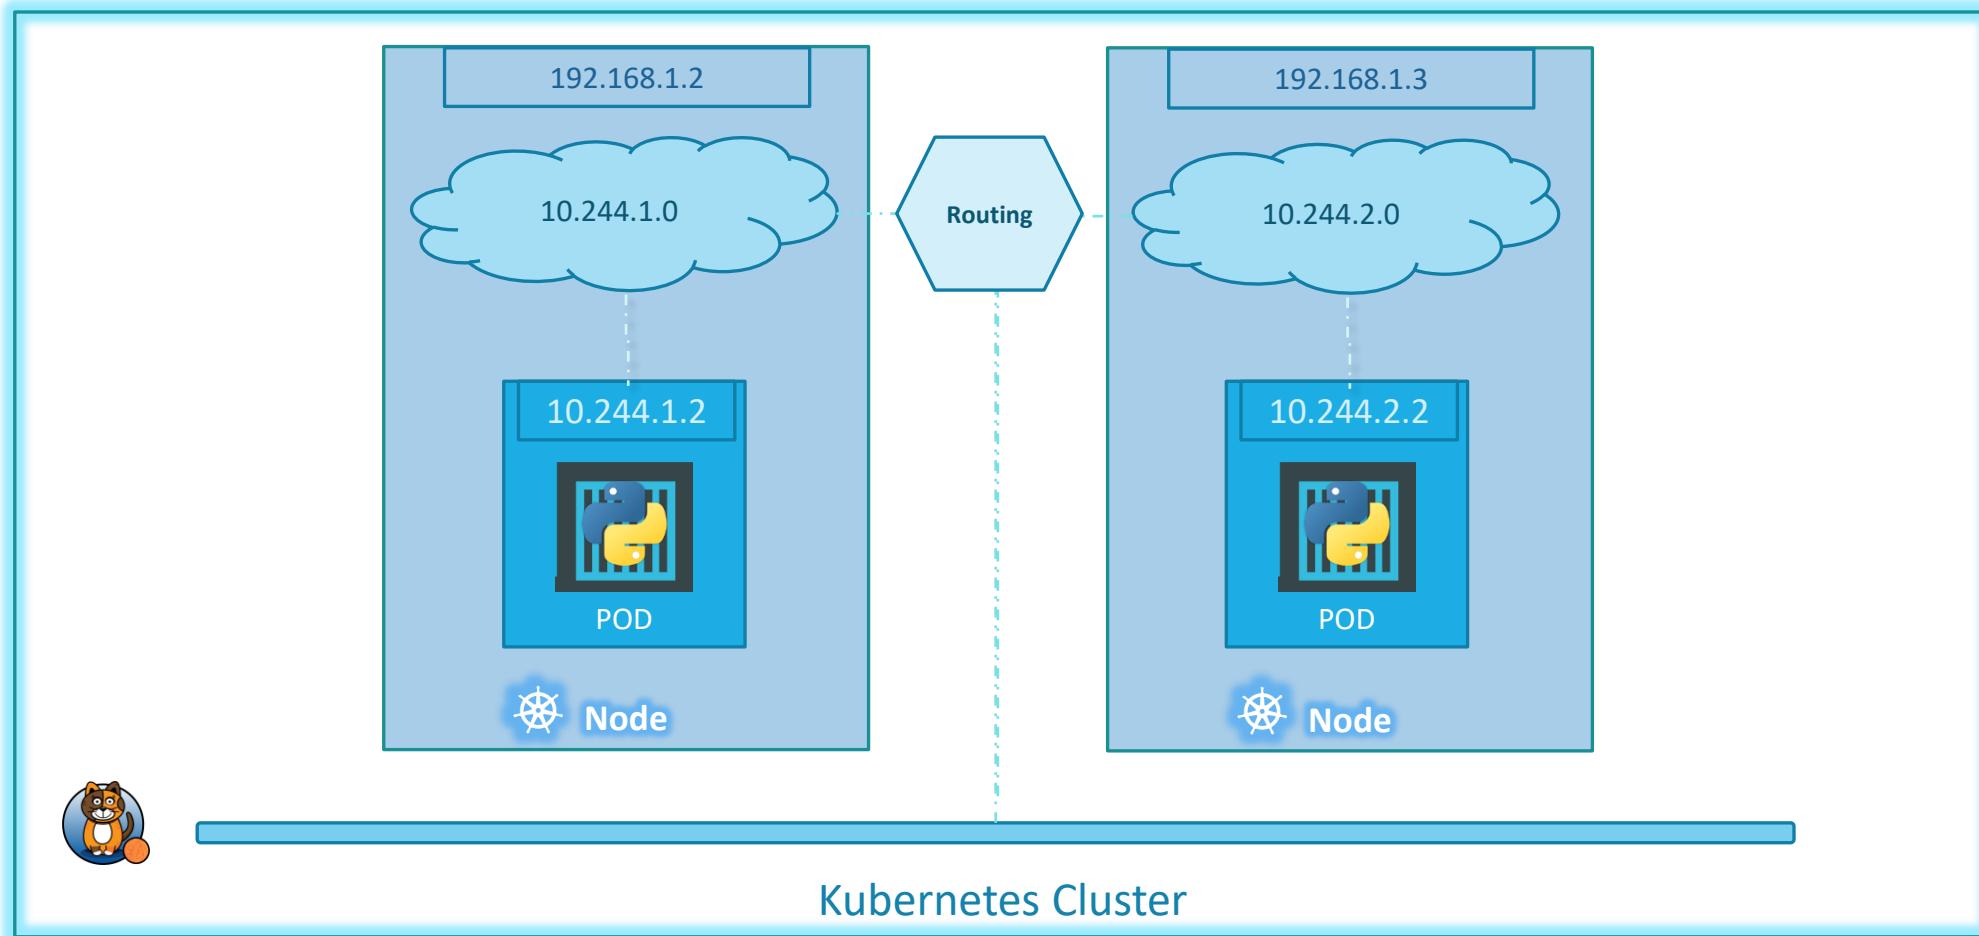
Refer to the Calico documentation for a [kubeadm quickstart](#), a [kubeadm installation guide](#), and other resources.

Note:

- In order for Network Policy to work correctly, you need to pass `--pod-network-cidr=192.168.0.0/16` to `kubeadm init`.
- Calico works on `amd64` only.

```
kubectl apply -f https://docs.projectcalico.org/v3.0/getting-started/kubernetes/installation/hosted/kubeadm/1.7/calico.yaml
```

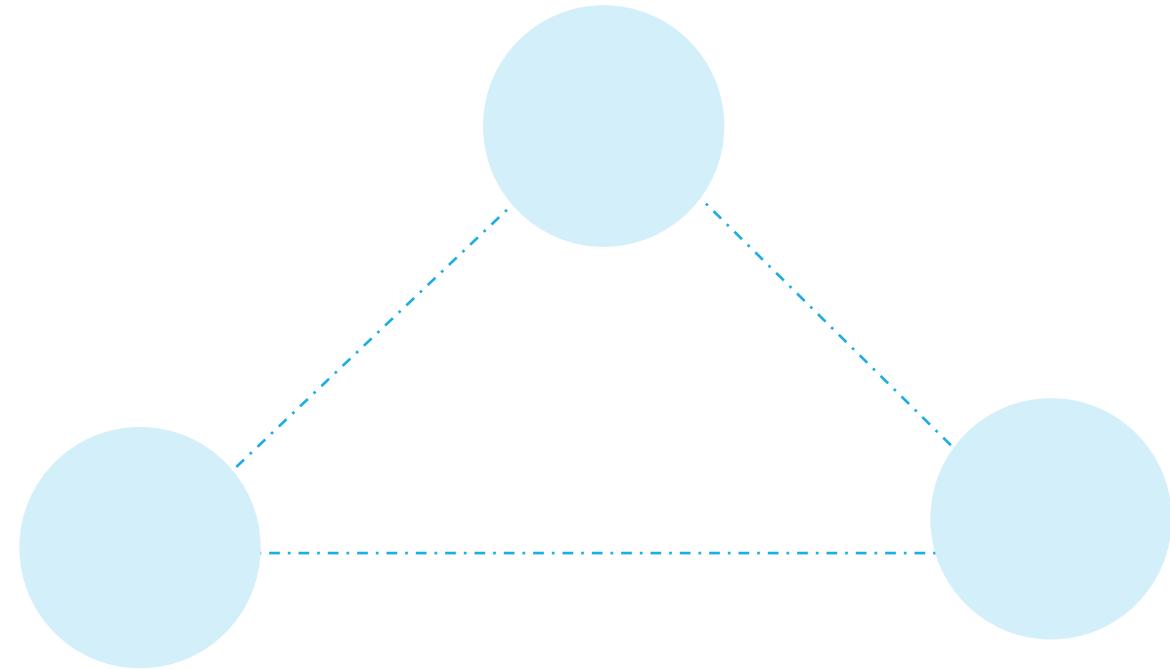
Cluster Networking



Demo

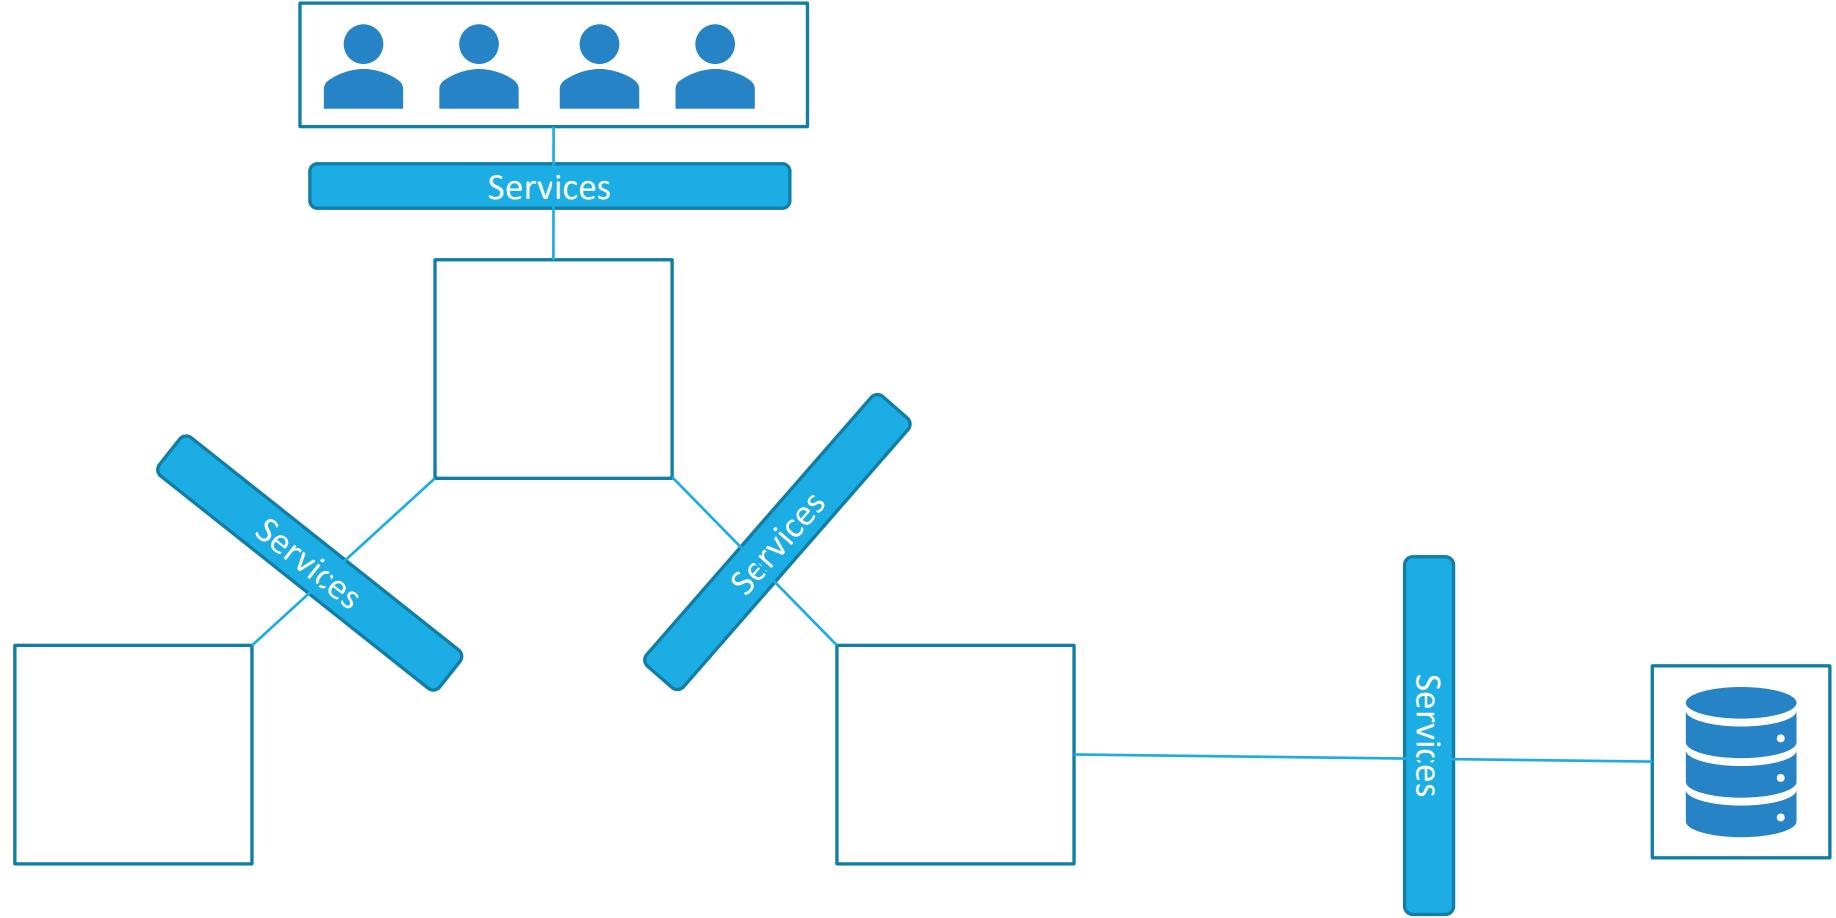
Networking

Services

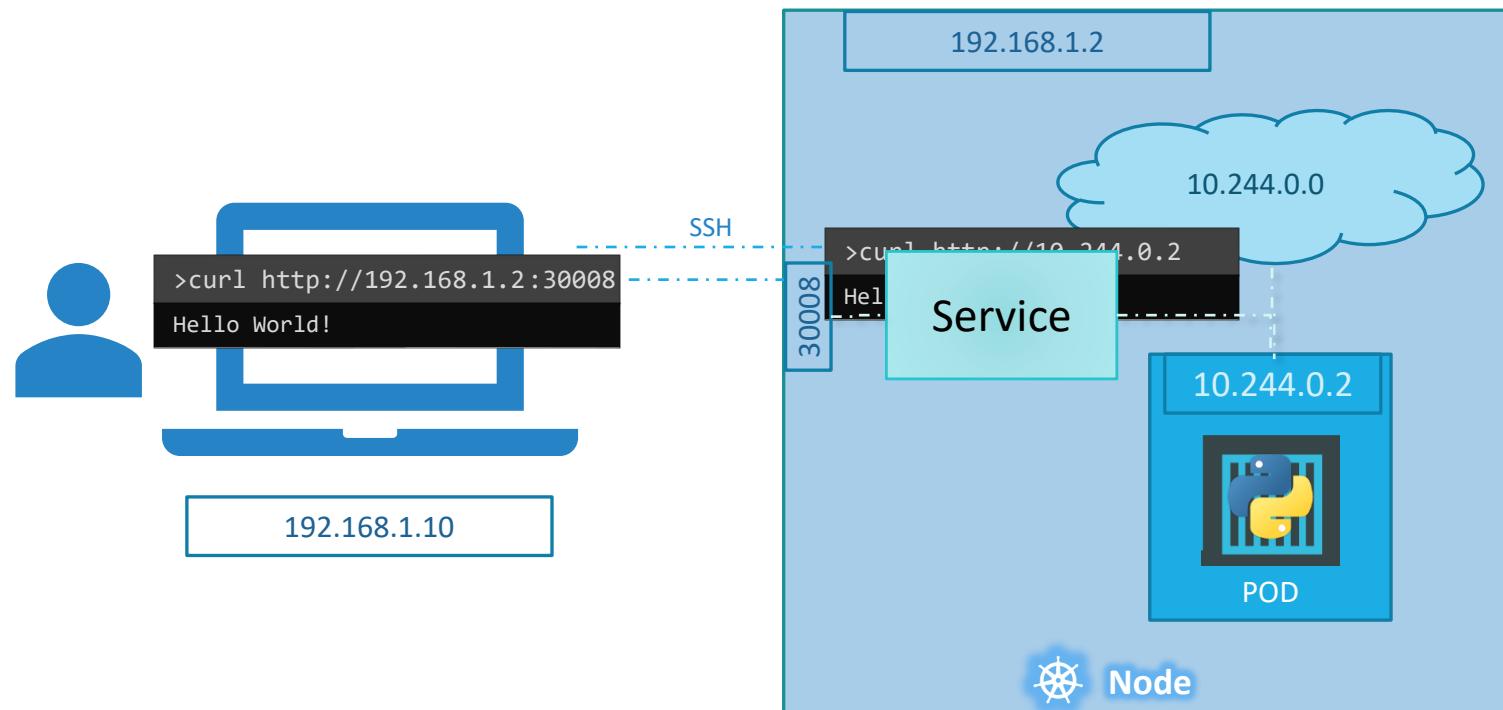


Raman Khanna

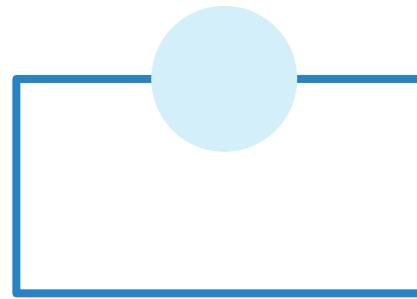
Services



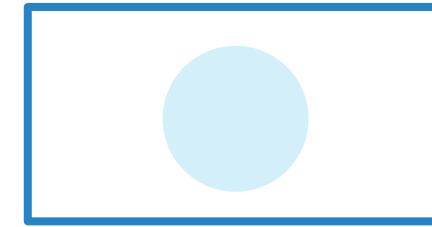
Service



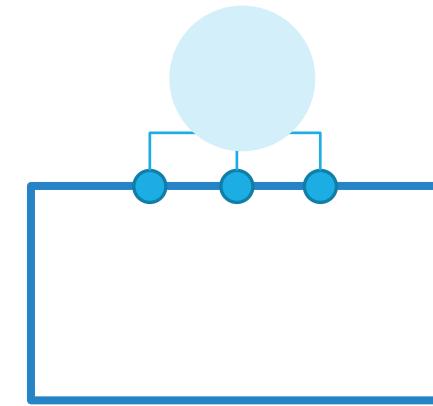
Services Types



NodePort

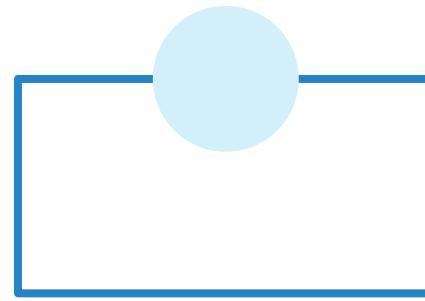


ClusterIP

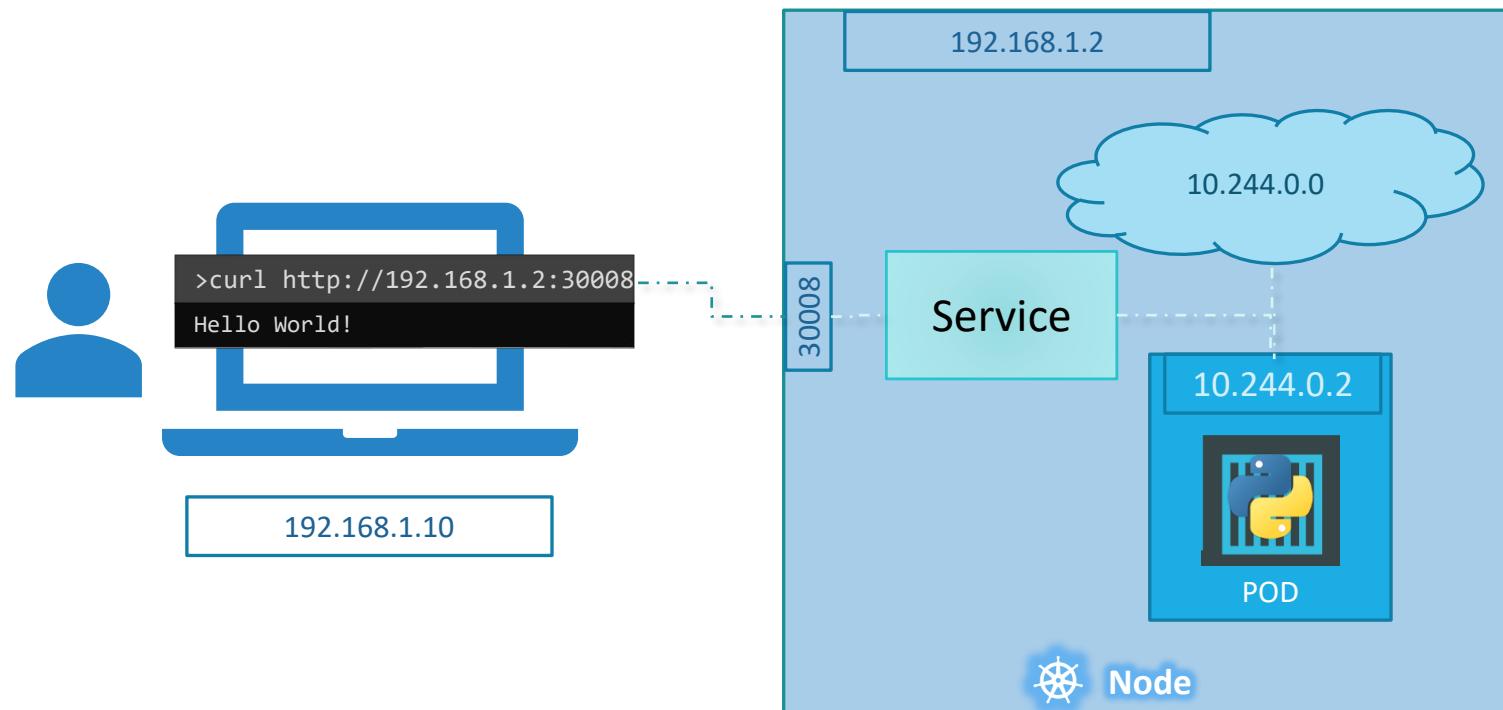


LoadBalancer

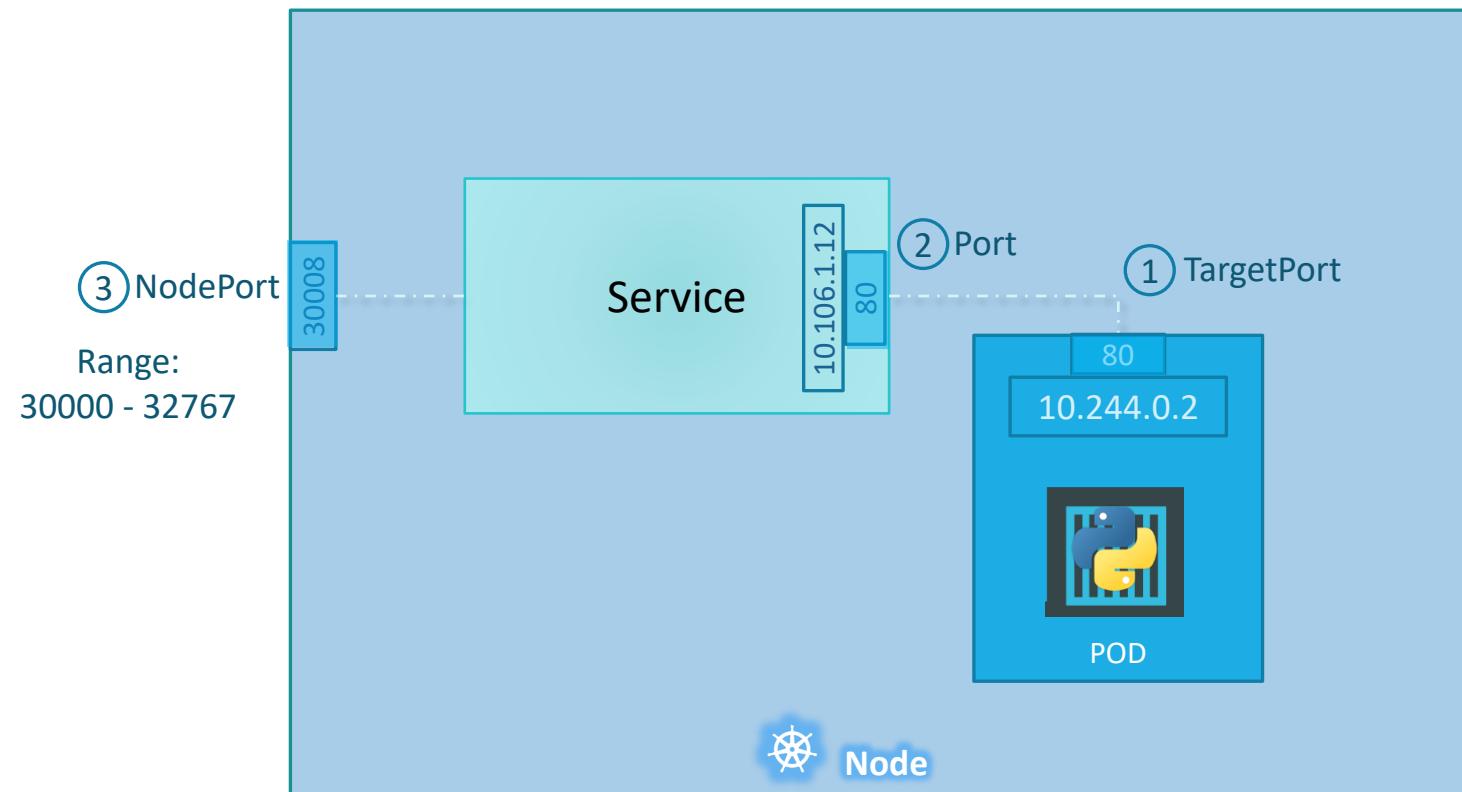
NodePort



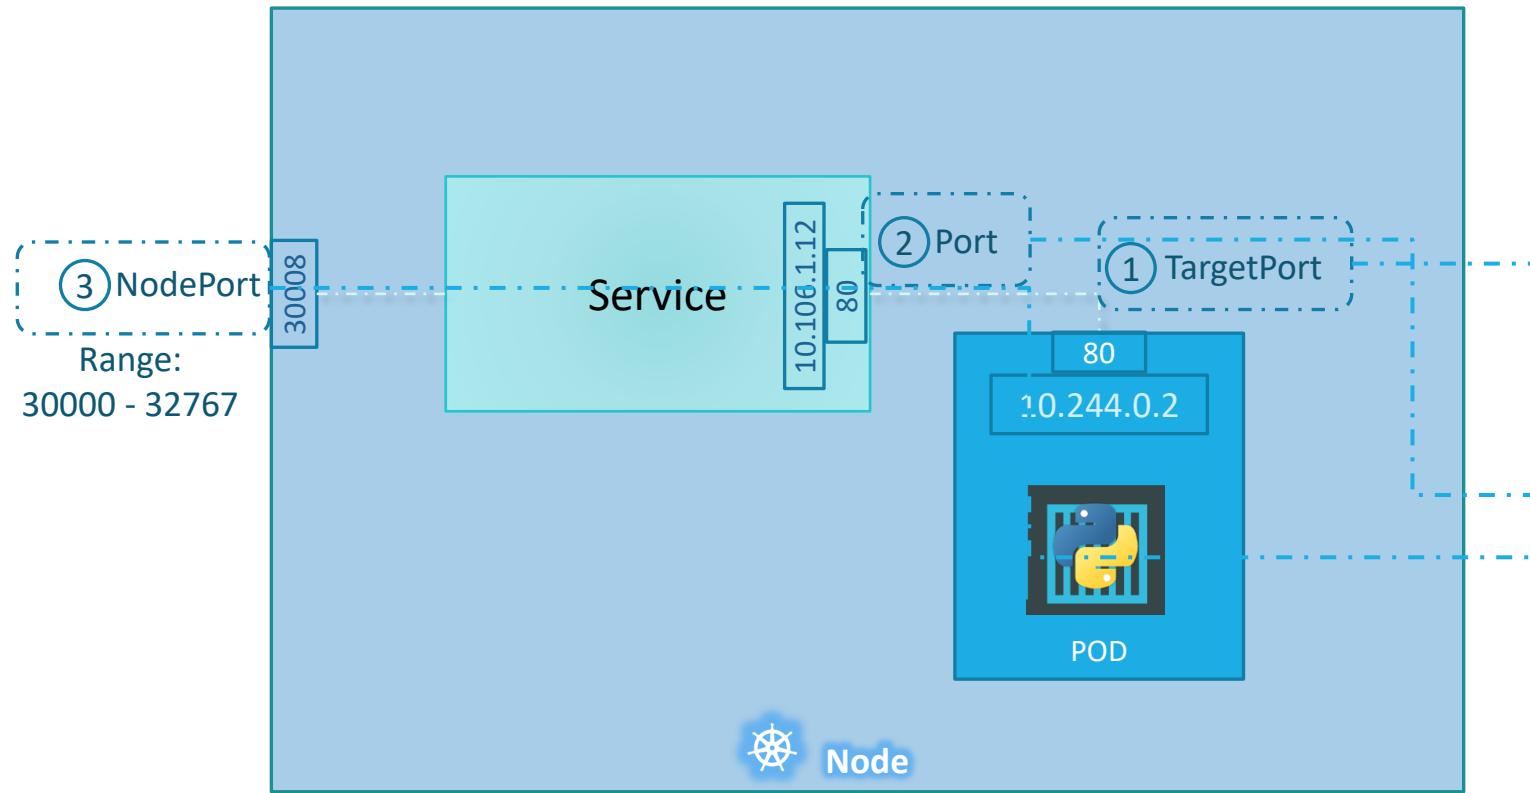
Service - NodePort



Service - NodePort



Service - NodePort



`service-definition.yml`

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```

Service - NodePort

service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
```

pod-definition.yml

```
> kubectl create -f service-definition.yml
service "myapp-service" created
```

```
> kubectl get services
```

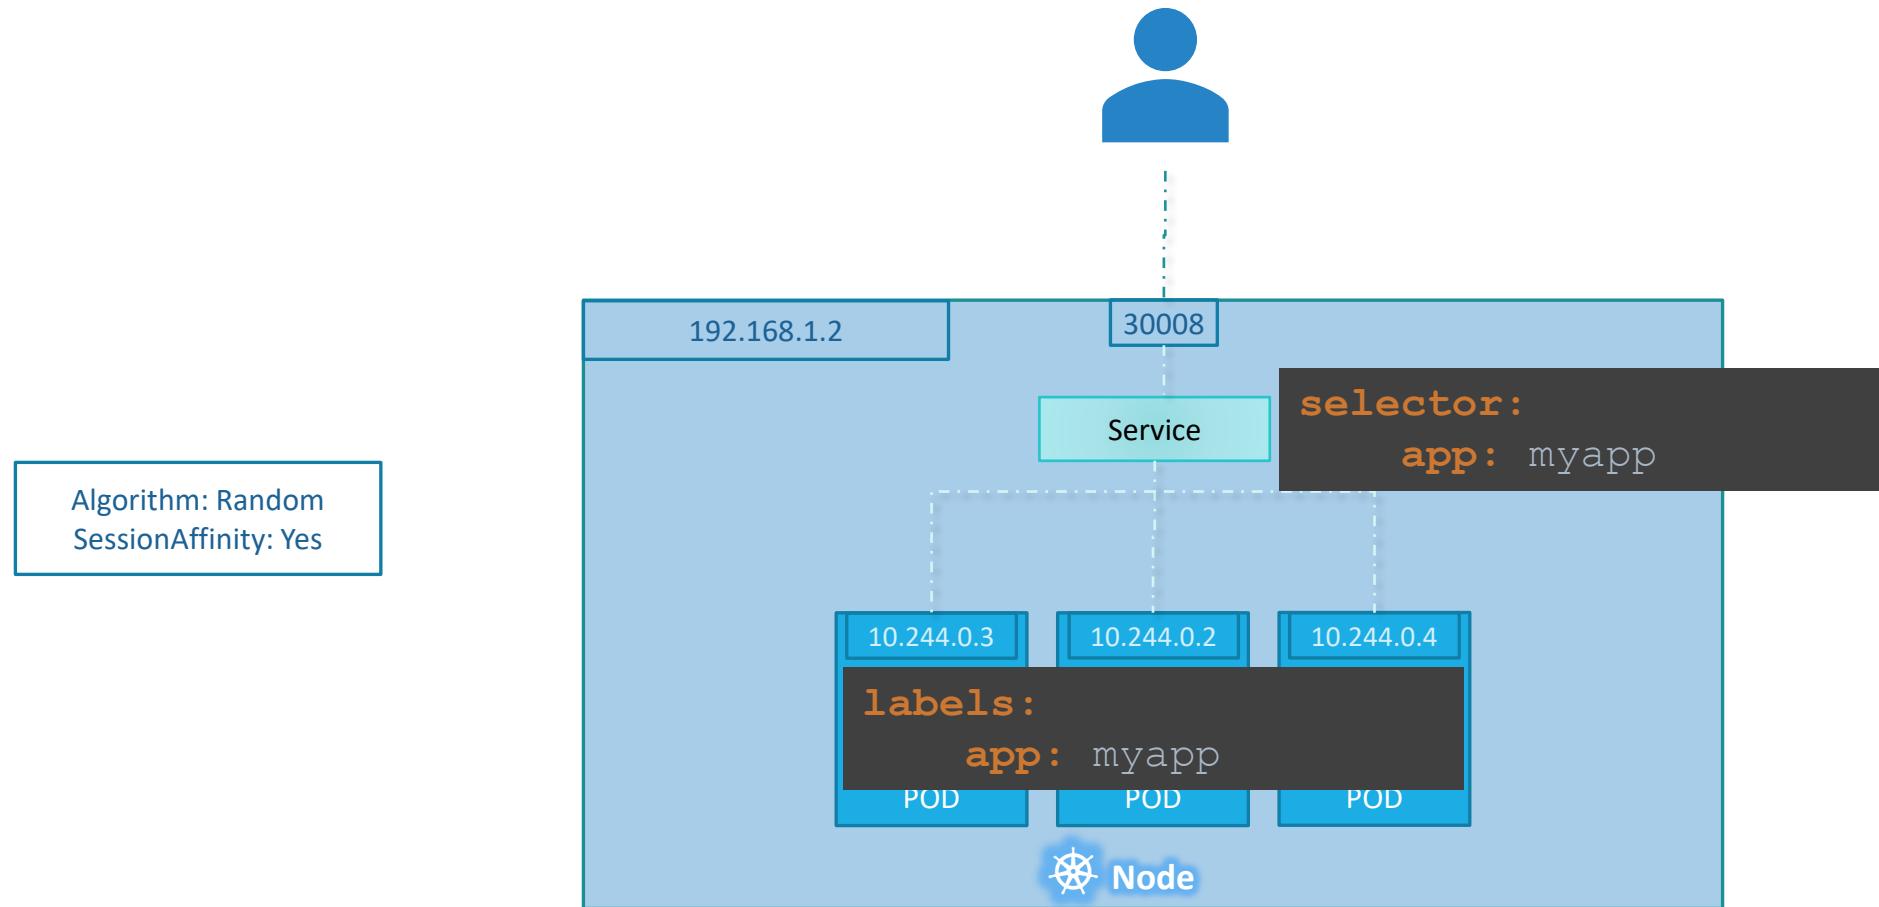
NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
myapp-service	NodePort	10.106.127.123	<none>	80:30008/TCP	5m

app: myapp

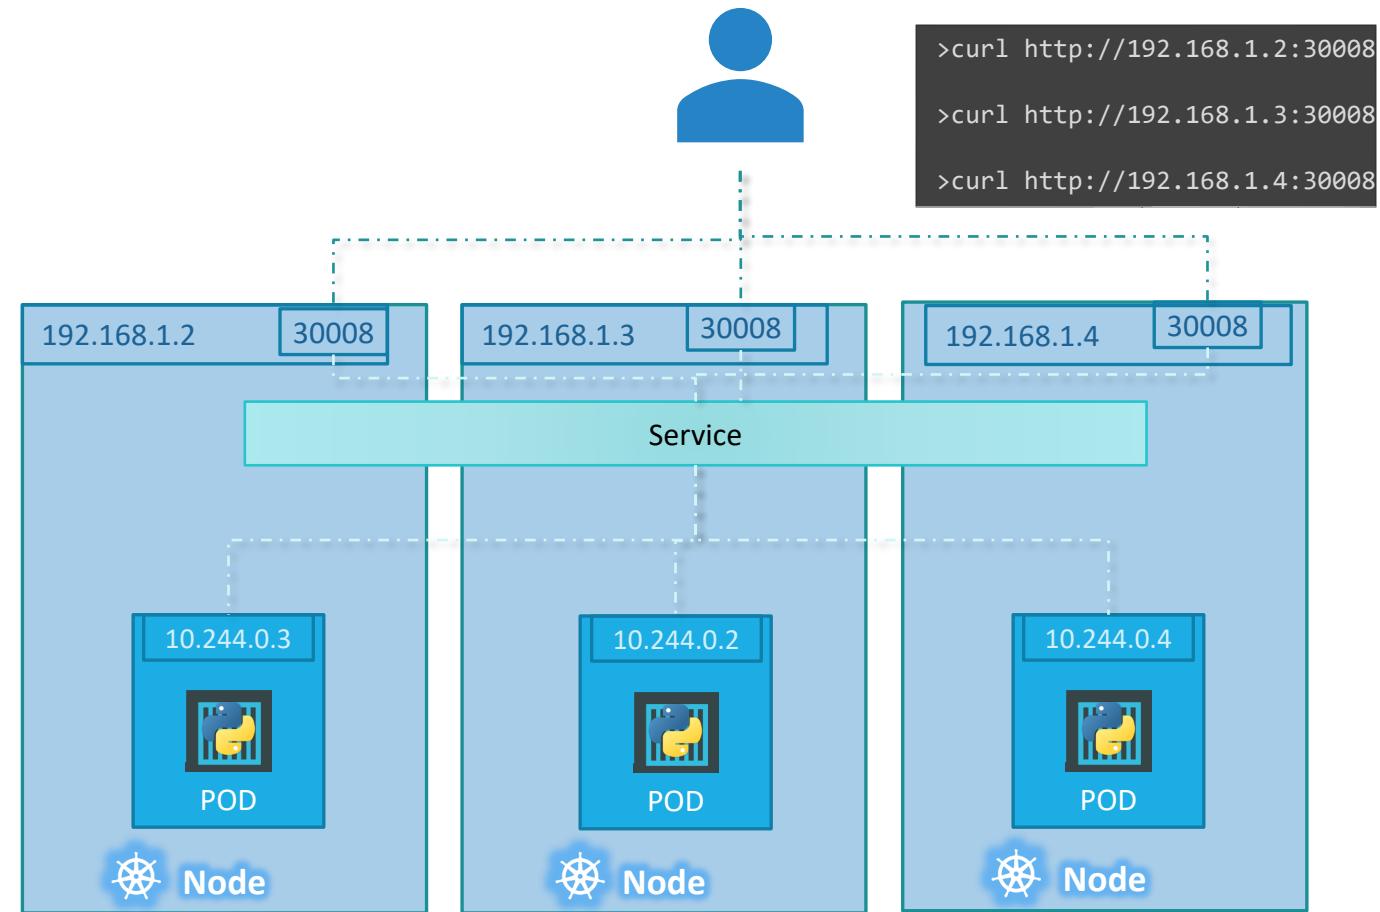
```
> curl http://192.168.1.2:30008
```

```
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
```

Service - NodePort



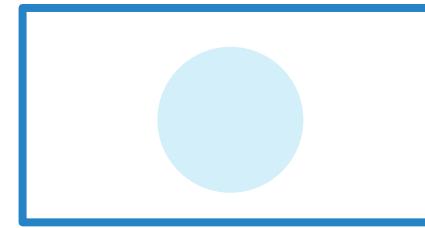
Service - NodePort



Demo

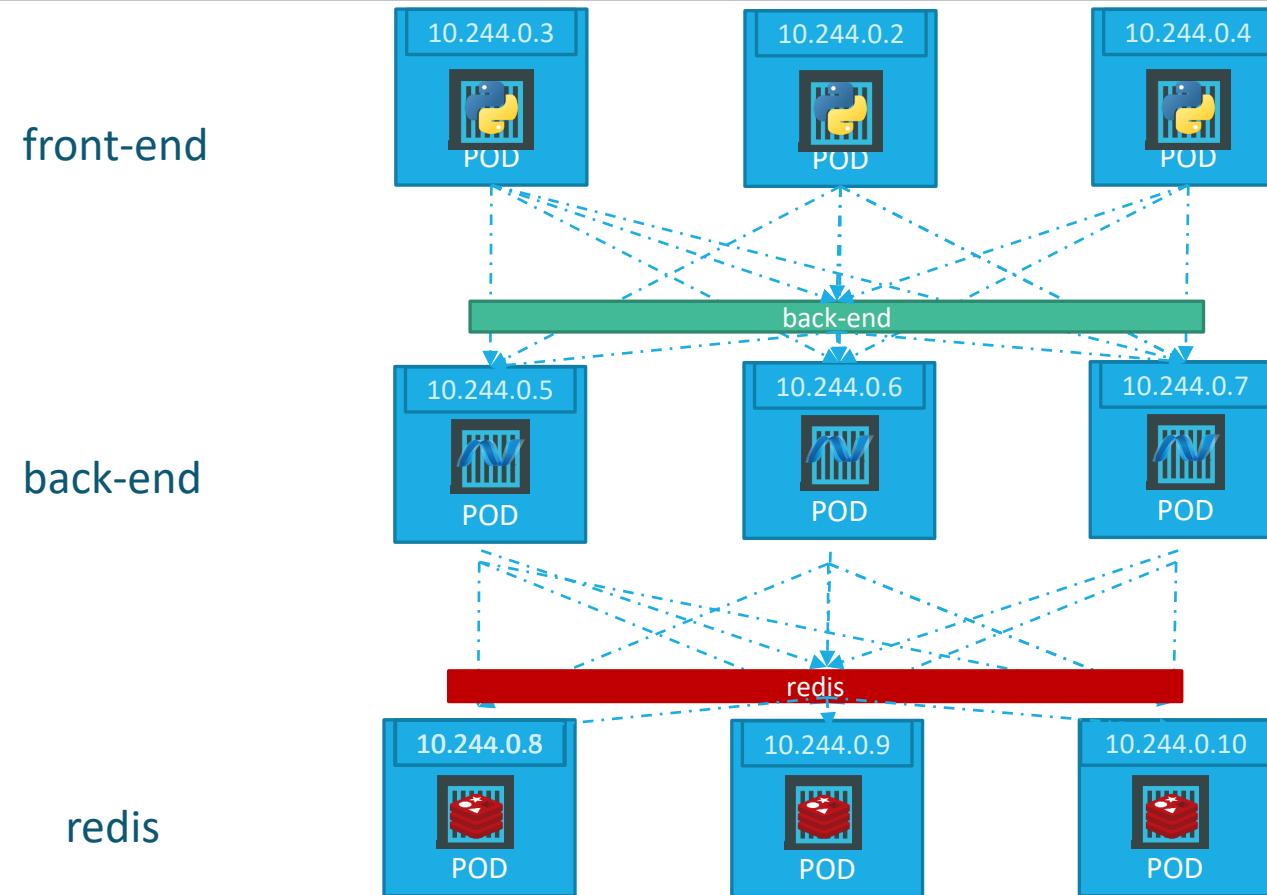
Service - NodePort

ClusterIP

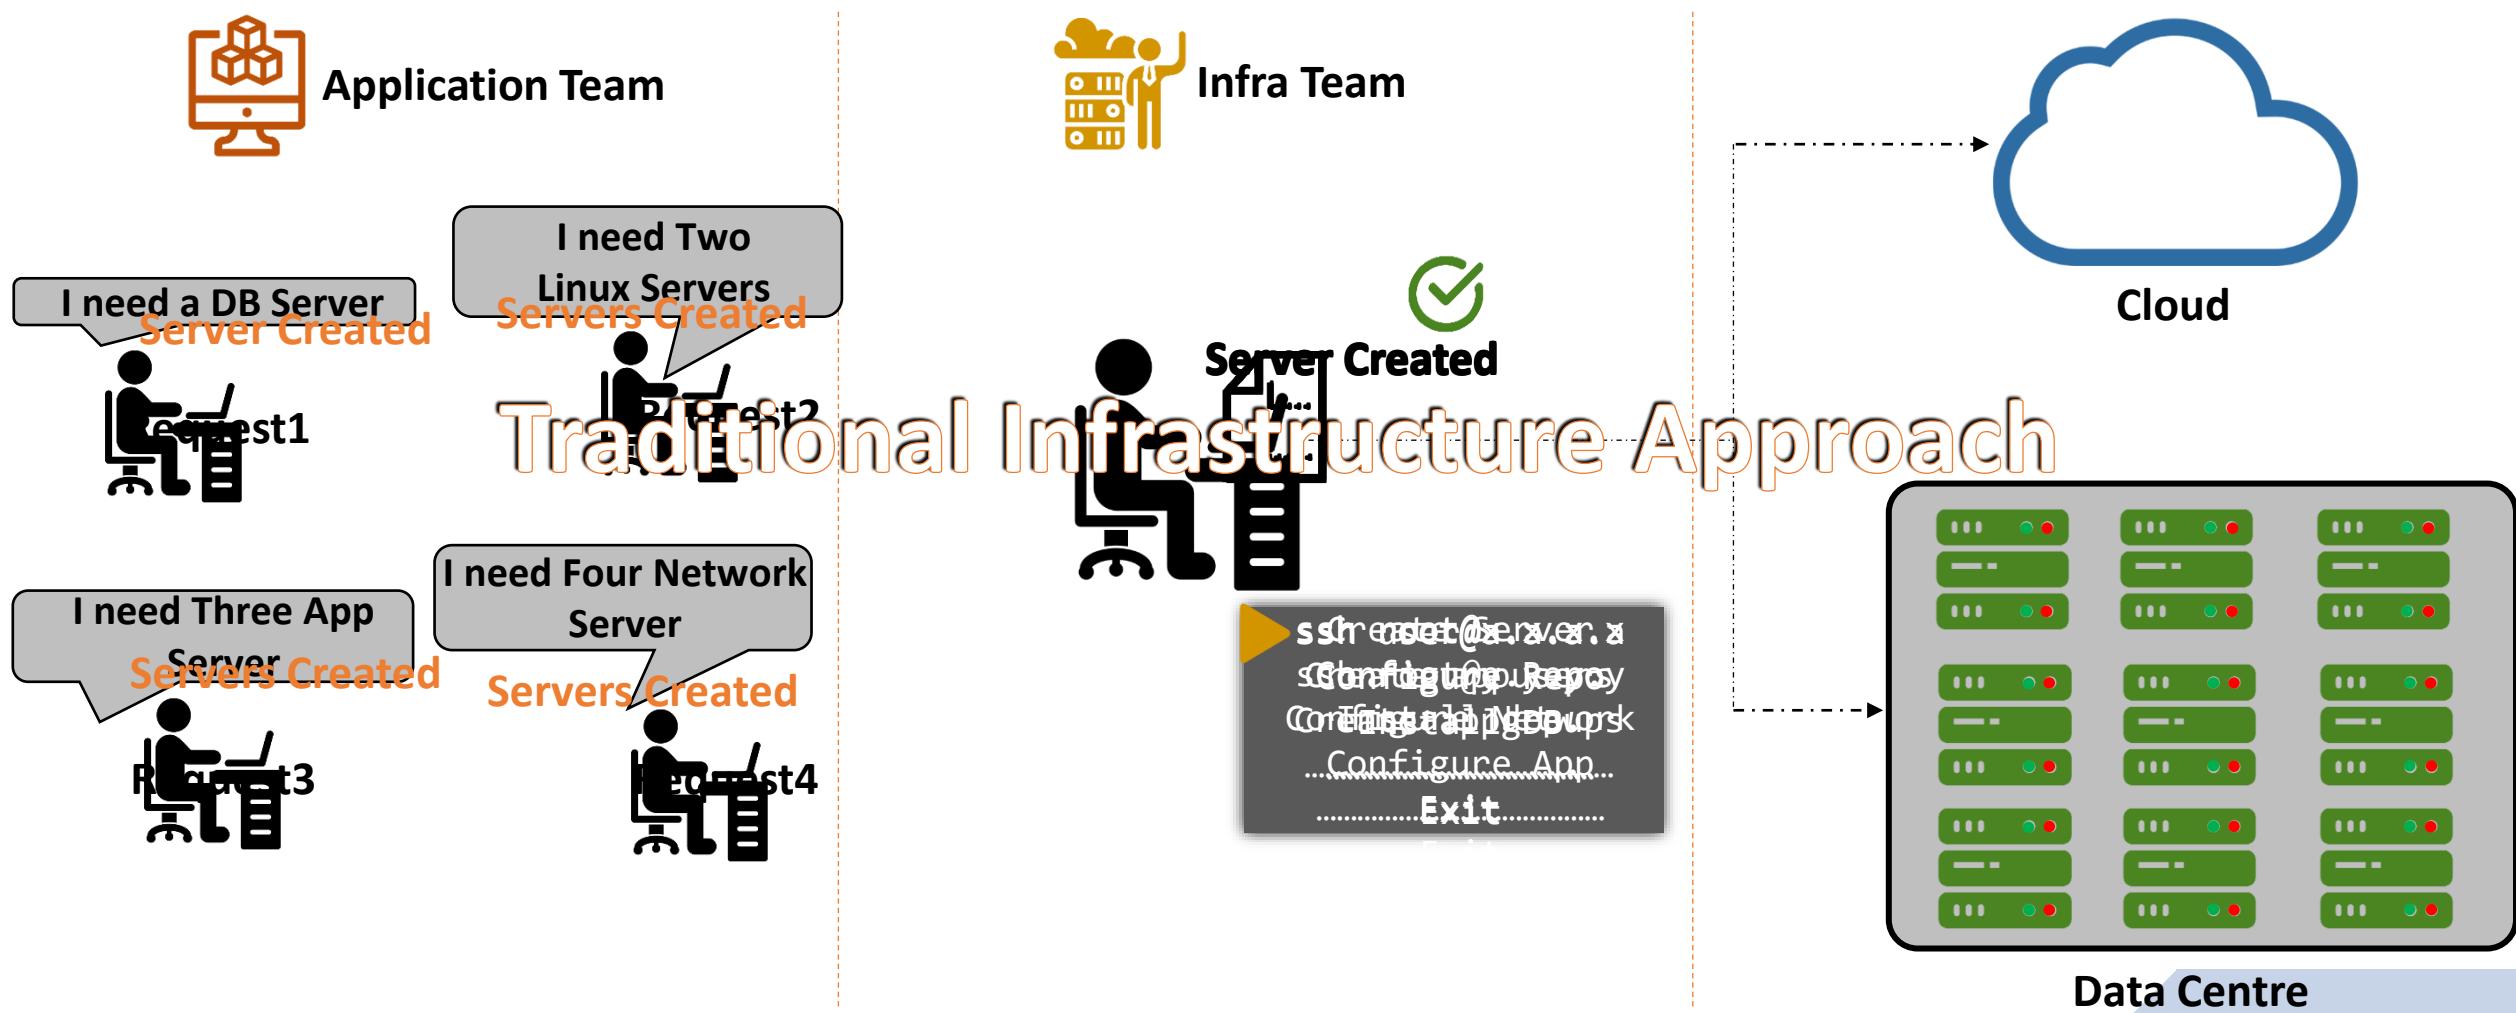


Raman Khanna

ClusterIP



Why DevOps IaC





Application Team

I need DB Server



need Two Linux Servers



need Three Linux Servers



need Four Linux Servers



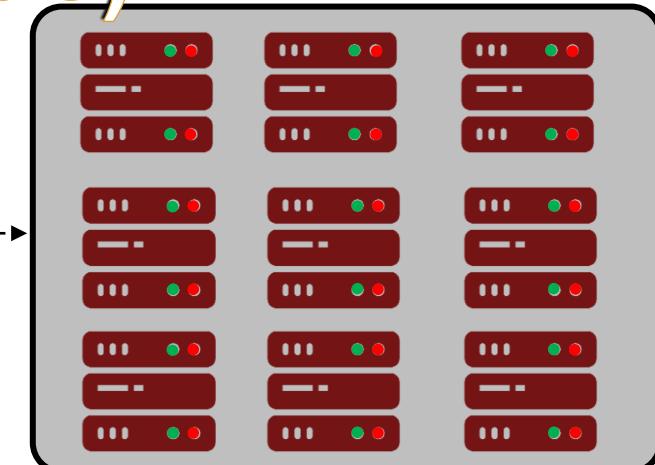
Infra Team

```
File is: main.tf
provider "aws" {
  region = "us-east-1"
}
resource "aws_instance" "requestfour" {
  count = "4"
  ami   = "ami-030f25abfb8b2"
  instance_type = "t2.micro"
  tags = {
    Name = "DevOpsInAction"
  }
}
output "myawsserver" {
  value =
"${aws_instance.requestfour[0].public_ip}"
}
```

IaC is Managing Infrastructure in files rather than
manually configuring resources in a user interface



Cloud



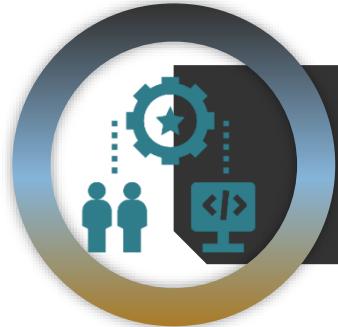
Data Centre

Terraform

Terraform is an easy-to-use IT Orchestration & Automation Software for System Administrators & DevOps Engineers.

- It is the infrastructure as code offering from Hashicorp.
- It is a tool for building, changing, and managing infrastructure in a safe, repeatable way.
- Configuration language called the HashiCorp Configuration Language (HCL) is used to configure the Infrastructure.
- Compatible with almost all major public and private Cloud service provider

Terraform



Infrastructure as
code (IAC)



July 2014, HashiCorp

What is Terraform?

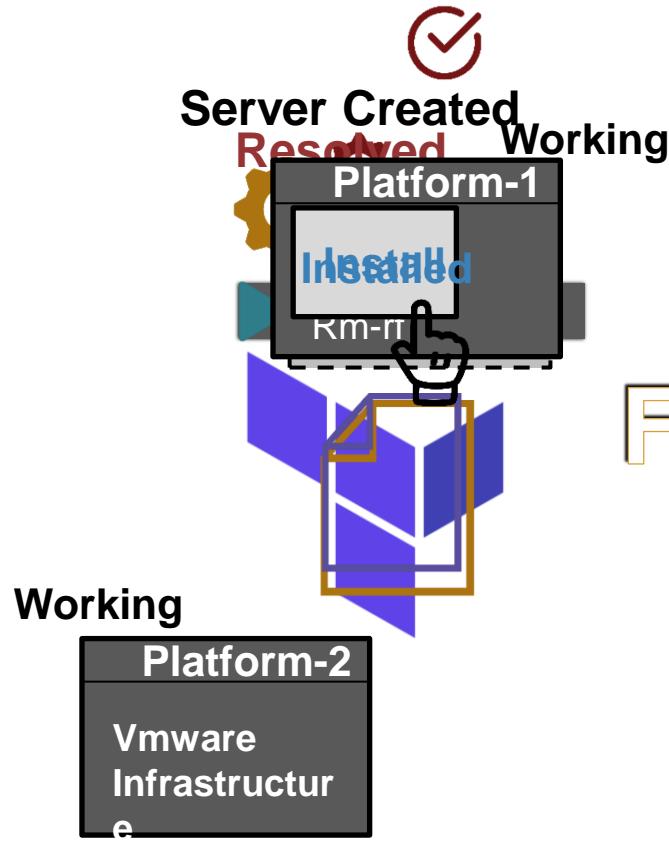


OpenSource /
Enterprise

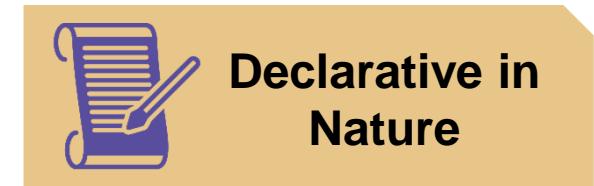


HCL (Hashicorp
Configuration
Language)

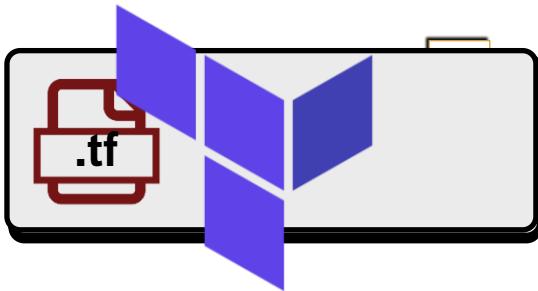
Terraform



Feature & Advantages



Terraform



Terraform Terminologies

Providers

Variables

Resources

Provisioners

DataSources

Outputs

Modules

File extension
.tf

Terraform

main.tf

```
provider "aws" {  
    region = "us-east-1"  
}
```

Provider Block

```
resource "aws_instance" "myserver" {  
    ami = "ami-030ff268bd7b4e8bE"  
    instance_type = "t2.micro"  
    tags = {  
        Name = "DevOpsInAction"  
    }  
}
```

Terraform File (Sample Code)

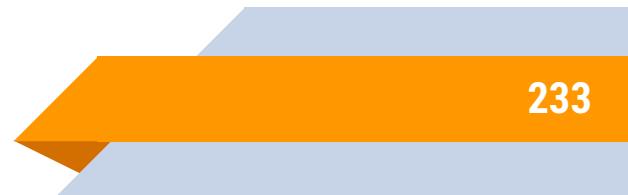
Resource Block

```
output "myserveroutputs" {  
    description = "Display Servers Public IP"  
    value = "${aws_instance.myserver.public_ip}"  
}
```

Output Block



Why Terraform?

- Infrastructure as Code – Write stuff in files, Version it, share it and collaborate with team on same.
 - Declarative in Nature
 - Automated provisioning
 - Clearly mapped Resource Dependencies
 - Can plan before you apply
 - Consistent
 - Compatible with multiple providers and infra can be combined on multiple providers
 - 50+ list of official and verified providers
 - Approx. 2500+ Modules readily available to work with
 - Both Community and Enterprise versions available
 - A best fit in DevOps IaC model
- 



Why Terraform?

- **Platform Agnostic** – Manage Heterogeneous Environment
- **Perfect State Management** – Maintains the state and Refreshes the state before each apply action.

Terraform state is the source of truth. If a change is made or a resource is appended to a configuration, Terraform compares those changes with the state file to determine what changes result in a new resource or resource modifications.

- **Confidence:** Due to easily repeatable operations and a planning phase to allow users to ensure the actions taken by Terraform will not cause disruption in their environment.

Terraform and its Peers

- Chef
- Puppet
- SaltStack
- Ansible
- CloudFormation
- Terraform
- Kubernetes

Terraform and its Peers

Many tools available in Market. Few things to consider, before selecting any tool:

- Configuration Management vs Orchestration
- Mutable Infrastructure vs Immutable Infrastructure
- Procedural vs Declarative

Terraform and its Peers

	Chef	Puppet	Ansible	SaltStack	CloudFormation	Terraform
Code	Open source	Open source	Open source	Open source	Closed source	Open source
Cloud	All	All	All	All	AWS only	All
Type	Config Mgmt	Config Mgmt	Config Mgmt	Config Mgmt	Orchestration	Orchestration
Infrastructure	Mutable	Mutable	Mutable	Mutable	Immutable	Immutable
Language	Procedural	Declarative	Declarative	Declarative	Declarative	Declarative
Architecture	Client/Server	Client/Server	Client-Only	Client/Server	Client-Only	Client-Only



Knowledge Checks

- What is Configuration Management?
- What is Orchestration?
- List a few available configuration Management tools.
- What are the Advantages of Terraform?

Summary: Terraform

Terraform is an easy-to-use IT Orchestration & Automation, Software for System Administrators & DevOps Engineers.

- Terraform is a tool for building, changing, and versioning infrastructure safely and efficiently.
- Terraform can manage existing and popular service providers as well as custom in-house solutions.
- Maintain Desired State
- Highly scalable and can create a complete datacenters in minutes
- Agentless solution
- Declaration in nature than Procedural
- Uses Providers API to provision the Infrastructure
- Terraform creates a dependency graph to determine the correct order of operations.

Providers

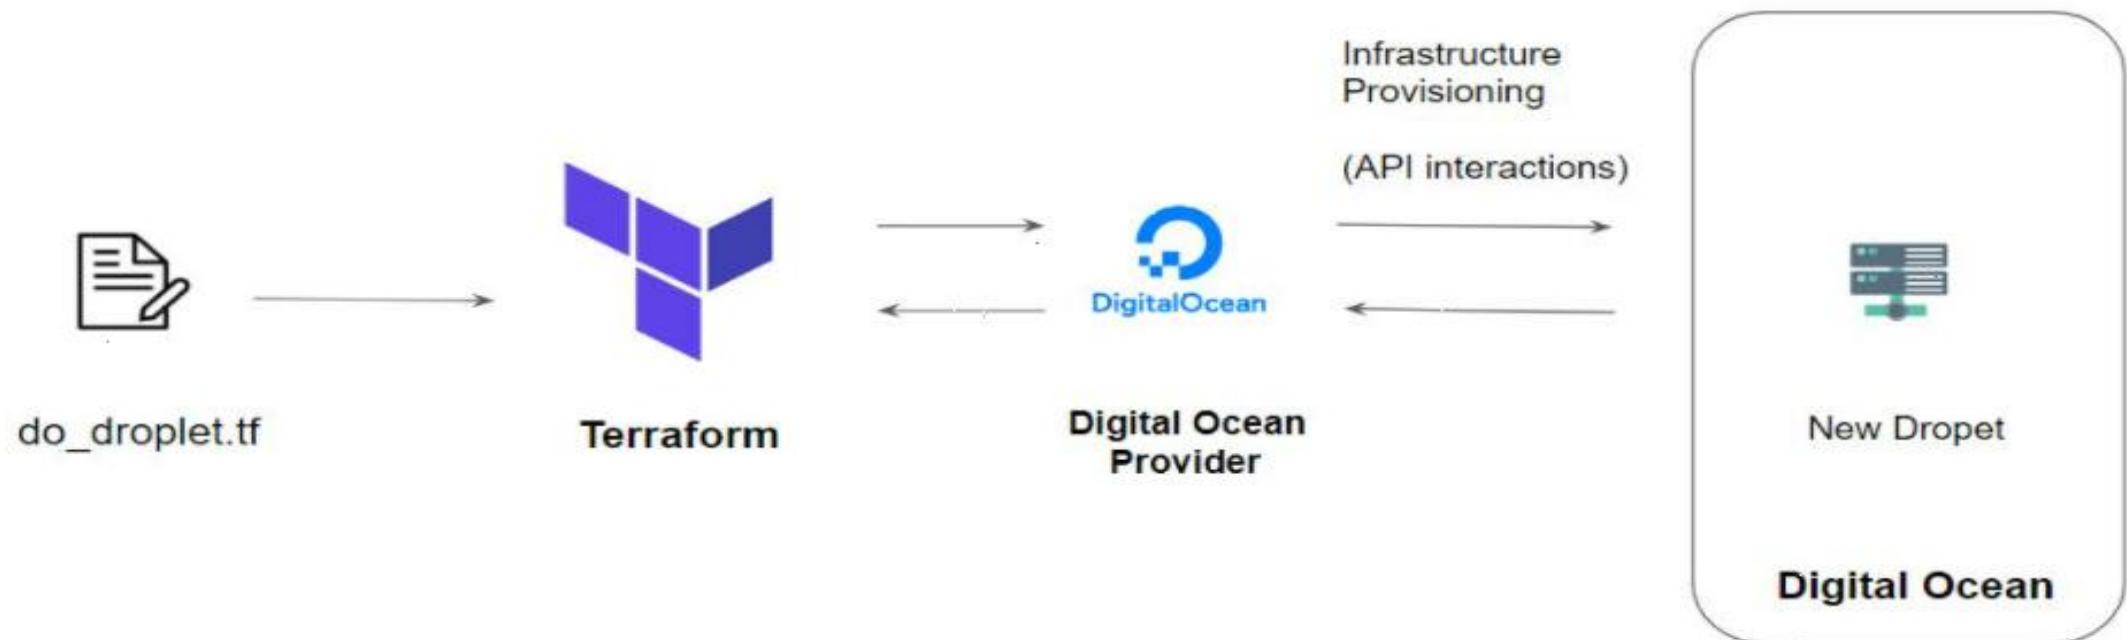
A provider is responsible for understanding API interactions and exposing resources over to a particular cloud service provider. Most providers configure a specific infrastructure platform (either cloud or self-hosted).

```
provider "aws" {  
    region    = "us-east-2"  
    access_key = "PUT-YOUR-ACCESS-KEY-HERE"  
    secret_key = "PUT-YOUR-SECRET-KEY-HERE"  
}
```

A provider is responsible for creating and managing resources.

<https://registry.terraform.io/browse/providers>

Overview of Provider Architecture :





AWS



Azure



Google Cloud Platform



Kubernetes



Oracle Cloud Infrastructure



Alibaba Cloud

Resources

- Resources are the most important element in the Terraform language. Each resource block describes one or more infrastructure objects, such as virtual networks, compute instances, etc
- ```
resource "aws_instance" "web" {
 ami = "ami-a1b2c3d4"
 instance_type = "t2.micro"
}
```

A resource block declares a resource of a given type ("aws\_instance") with a given local name ("web"). The name is used to refer to this resource from elsewhere in the same Terraform module but has no significance outside that module's scope.

The resource type and name together serve as an identifier for a given resource and so must be unique within a module.

Resource names must start with a letter or underscore, and may contain only letters, digits, underscores, and dashes.

# LAB 3: Creating first ec2 instance

..

■ <https://registry.terraform.io/providers/hashicorp/aws/latest/docs/resources/instance>

# Configuration files

- Whatever you want to achieve(deploy) using terraform will be achieved with configuration files.
- Configuration files ends with .tf extension (tf.json for json version).
- Terraform uses its own configuration language, designed to allow concise descriptions of infrastructure.
- The Terraform language is declarative, describing an intended goal rather than the steps to reach that goal.
- A group of resources can be gathered into a module, which creates a larger unit of configuration.
- As Terraform's configuration language is declarative, the ordering of blocks is generally not significant. Terraform automatically processes resources in the correct order based on relationships defined between them in configuration

# Example

- You can write up the terraform code in hashicorp Language – HCL.
- Your configuration file will always endup with .tf extension

```
provider "aws" {
 region = "us-east-2"
 access_key = "PUT-YOUR-ACCESS-KEY-HERE"
 secret_key = "PUT-YOUR-SECRET-KEY-HERE"
}

resource "aws_instance" "myec2" {
 ami = "ami-082b5a644766e0e6f"
 instance_type = "t2.micro"
}
tags = {
 Name = "Techlanders-aws-ec2-instance"
}
```

# Terraform Workflow

## Few Steps to work with terraform:

- 1) Set the Scope - Confirm what resources need to be created for a given project.
- 2) Author - Create the configuration file in HCL based on the scoped parameters
- 3) Run **terraform init** to initialize the plugins and modules
- 4) Run **terraform validate** to validate the template
- 5) Do **terraform plan**
- 6) Run **terraform apply** to apply the changes



# Terraform validate

- Terraform validate will validate the terraform configuration file
- It'll throw error for syntax issues:

```
[root@TechLanders aws]# terraform validate
Success! The configuration is valid.
```

```
[root@TechLanders aws]#
```

# Terraform init

- Terraform init will initialize the modules and plugins.
- If you ever set or change modules or backend configuration for Terraform, rerun this command to reinitialize your working directory.
- If you forget running init, terraform plan/apply will remind you about initialization.
- Terraform init will download the connection plugins from Repository “registry.terraform.io” under your current working directory/.terraform:  

```
[root@TechLanders plugins]# pwd
/root/aws/.terraform/plugins
[root@TechLanders plugins]# ls -l
total 4
drwxr-xr-x. 3 root root 23 Aug 15 07:06 registry.terraform.io
-rw-r--r--. 1 root root 136 Aug 15 07:06 selections.json
[root@TechLanders plugins]#
```
- Important concept:
  - Always make a best practice to initialize the terraform modules with versions. i.e.  
hashicorp/aws: version = "~~> 3.2.0"

# Example

- Perform Terraform Init:

```
[root@TechLanders aws]# terraform init
```

Initializing the backend...

Initializing provider plugins...

- Finding latest version of hashicorp/aws...
- Installing hashicorp/aws v3.2.0...
- Installed hashicorp/aws v3.2.0 (signed by HashiCorp)

The following providers do not have any version constraints in configuration,  
so the latest version was installed.

To prevent automatic upgrades to new major versions that may contain breaking  
changes, we recommend adding version constraints in a required\_providers block  
in your configuration, with the constraint strings suggested below.

\* hashicorp/aws: version = "~> 3.2.0"

Terraform has been successfully initialized!

If you ever set or change modules or backend configuration for Terraform,  
rerun this command to reinitialize your working directory. If you forget, other  
commands will detect it and remind you to do so if necessary.

```
[root@TechLanders aws]#
```

# Terraform plan

- terraform plan will create an execution plan and will update you what changes it going to make.
- It'll update you upfront what its gonna add, change or destroy.
- Terraform will automatically resolve the dependency between components- which to be created first and which in last.

```
[root@TechLanders aws]# terraform plan
```

Refreshing Terraform state in-memory prior to plan...

The refreshed state will be used to calculate this plan but will not be persisted to local or remote state storage.

An execution plan has been generated and is shown below. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
aws_instance.myserver will be created
+ resource "aws_instance" "myserver" {
 + ami = "ami-06b35f67f1340a795"
 + arn = (known after apply)
```

Plan: 1 to add, 0 to change, 0 to destroy.

# Terraform apply

- Terraform apply will apply the changes.
- Before it applies changes, it'll showcase changes again and will ask to confirm to move ahead:

```
[root@TechLanders aws]# terraform apply
```

An execution plan has been generated and is shown below. Resource actions are indicated with the following symbols:

+ create

Do you want to perform these actions? Terraform will perform the actions described above. Only 'yes' will be accepted to approve.

Enter a value: yes

```
aws_instance.myserver: Creating...
```

```
aws_instance.myserver: Still creating... [10s elapsed]
```

```
aws_instance.myserver: Still creating... [20s elapsed]
```

```
aws_instance.myserver: Creation complete after 21s [id=i-0a63756c96d338801]
```

```
Apply complete! Resources: 1 added, 0 changed, 0 destroyed.
```

```
[root@TechLanders aws]#
```

# Terraform apply

- Terraform apply will create **tfstate** file to maintain the desired state:

```
[root@TechLanders aws]# ls -l
total 8
-rw-r--r--. 1 root root 234 Aug 15 07:06 myinfra.tf
-rw-r--r--. 1 root root 3209 Aug 15 08:02 terraform.tfstate
[root@TechLanders aws]# cat terraform.tfstate
{
 "version": 4,
 "terraform_version": "0.13.0",
 "serial": 1,
 "lineage": "7f7e0e15-95ef-d8fa-b1cd-12024aed5fa6",
 "outputs": {},
 "resources": [
 "provider": "provider[\"registry.terraform.io/hashicorp/aws\"]",
 "instances": [
 {
 "schema_version": 1,
 "attributes": {
 "ami": "ami-06b35f67f1340a795",
 "arn": "arn:aws:ec2:us-east-2:677729060277:instance/i-0a63756c96d338801",
 }
 }
]
]
}
```

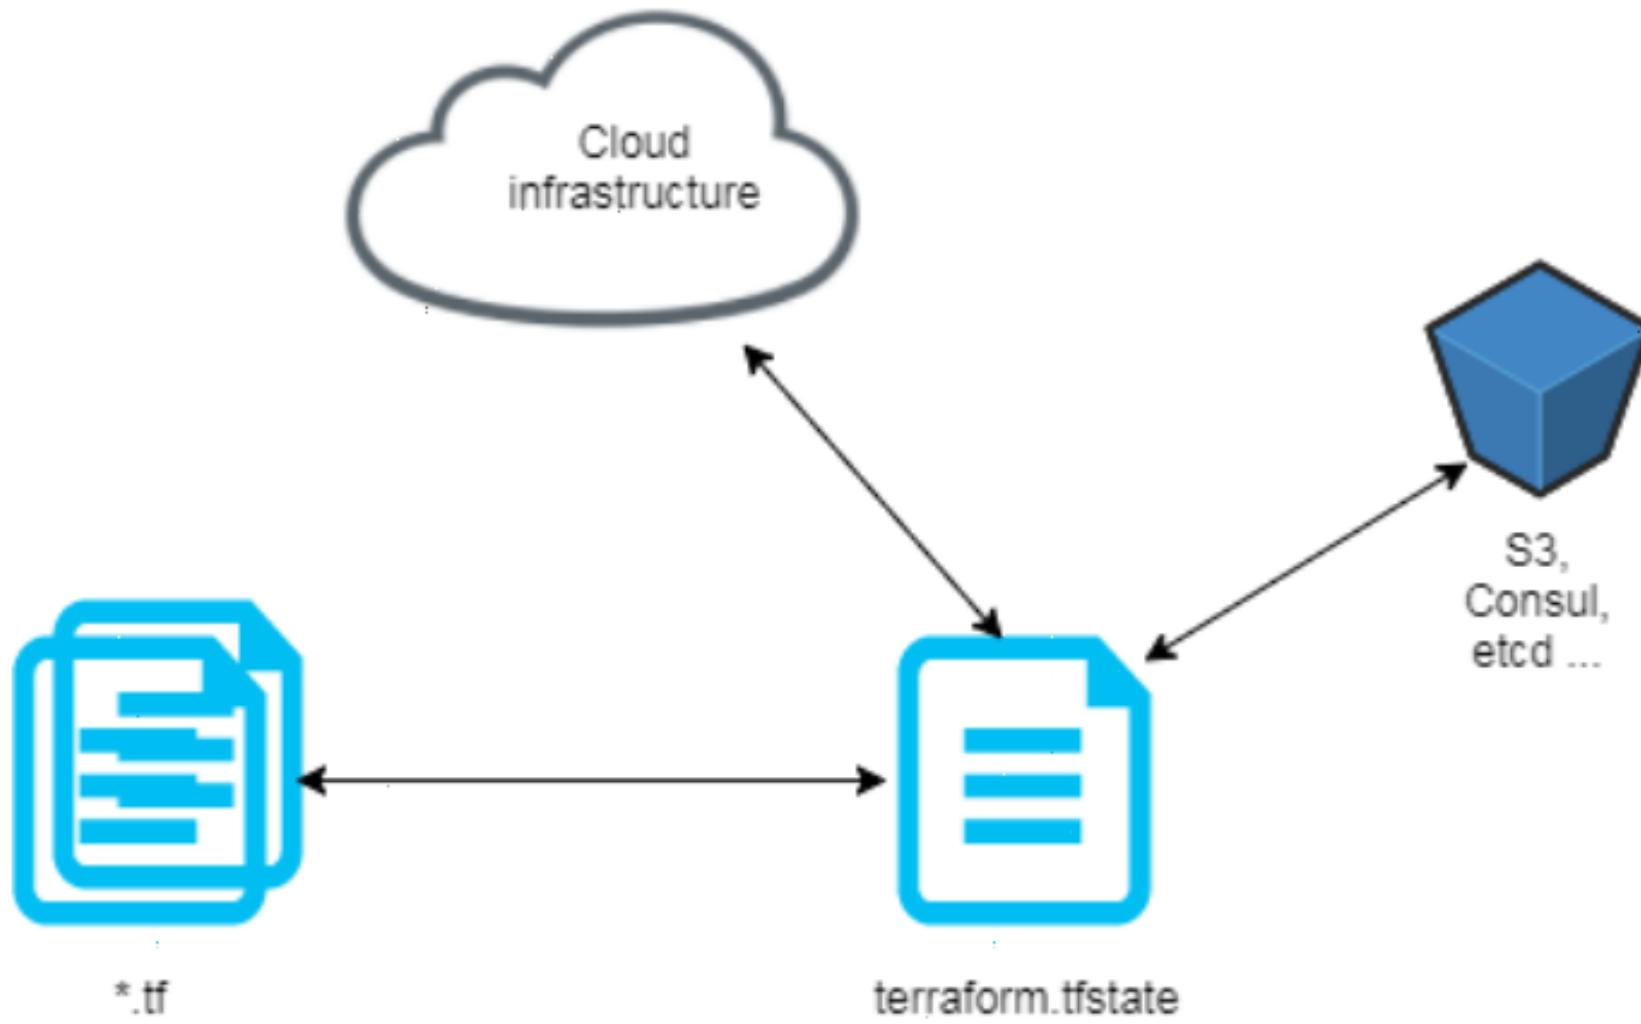
- Note: -auto-approve option can be given alongwith terraform apply to avoid the human intervention.

# Terraform show

- Terraform show will show the current state of the environment been created by your config file:

```
[root@ip-172-31-6-233 aws]# terraform show
aws_instance.myserver:
resource "aws_instance" "myserver" {
 ami = "ami-06b35f67f1340a795"
 arn = "arn:aws:ec2:us-east-2:677729060277:instance/i-0a63756c96d338801"
 associate_public_ip_address = true
 availability_zone = "us-east-2a"
 cpu_core_count = 1
 cpu_threads_per_core = 1

```



# Desired State Maintenance (DSC)

- Delete the newly created server and then check for the terraform plan

```
[root@TechLanders aws]# terraform plan
```

Refreshing Terraform state in-memory prior to plan...

The refreshed state will be used to calculate this plan, but will not be persisted to local or remote state storage.

```
aws_instance.myserver: Refreshing state... [id=i-0a63756c96d338801]
```

An execution plan has been generated and is shown below.

Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
aws_instance.myserver will be created
```

```
+ resource "aws_instance" "myserver" {
```

- Run terraform apply command again and witness the provisioning of new server on console.

```
[root@TechLanders aws]# terraform apply
```

```
aws_instance.myserver: Refreshing state... [id=i-0a63756c96d338801]
```

An execution plan has been generated and is shown below.

Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
aws_instance.myserver will be created
```

# Infrastructure as Code

- Modify your template file to change the instance size from t2.micro to t2.small and plan/apply the changes:

```
[root@TechLanders aws]# cat myinfra.tf
resource "aws_instance" "myserver" {
 ami = "ami-06b35f67f1340a795"
 instance_type = "t2.small"
}
[root@TechLanders aws]#
```

- Run terraform plan and apply again to check the differences

```
[root@TechLanders aws]# terraform apply
aws_instance.myserver: Refreshing state... [id=i-0a1f8a600cb968c7c]
```

An execution plan has been generated and is shown below.  
Resource actions are indicated with the following symbols:

~ update in-place

Plan: 0 to add, 1 to change, 0 to destroy.

Do you want to perform these actions?

Terraform will perform the actions described above.

Only 'yes' will be accepted to approve.

Enter a value: yes

```
aws_instance.myserver: Modifying... [id=i-0a1f8a600cb968c7c]
```

# Refreshing the state

- In case the requirement is to just check for any updates been done in the running environment, we can run terraform refresh command:

```
C:\Users\gagandeep\Desktop\terraform>terraform refresh
```

```
google_compute_network.vpc_network: Refreshing state... [id=projects/accenture-286519/global/networks/terraform-net3]
```

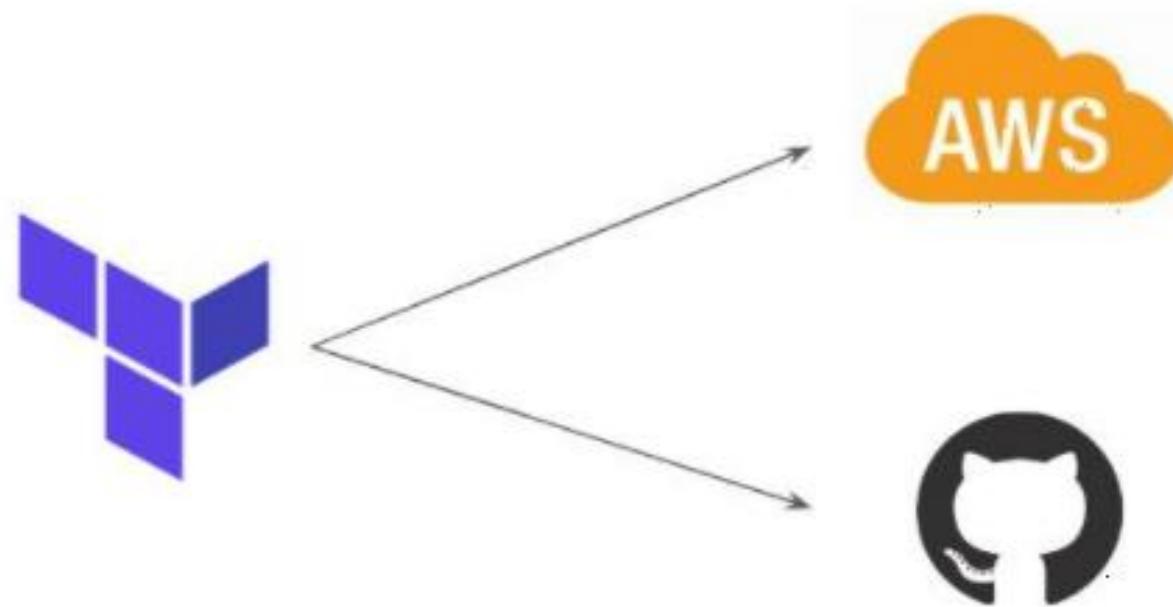
```
google_compute_address.vm_static_ip: Refreshing state... [id=projects/accenture-286519/regions/us-central1/addresses/terraform-static-ip1]
```

```
google_compute_instance.vm_instance1: Refreshing state... [id=projects/accenture-286519/zones/us-central1-b/instances/terraform-instance1]
```

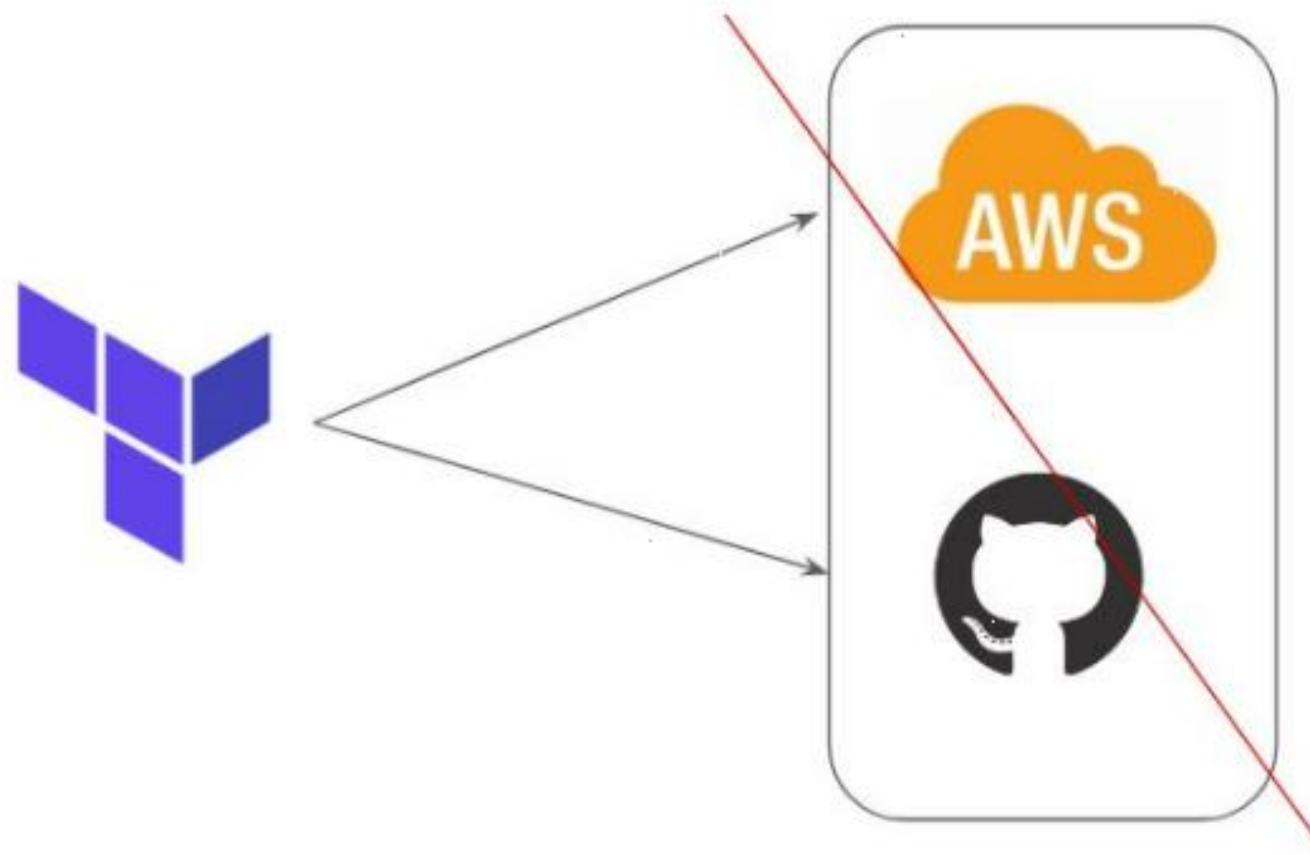
```
C:\Users\gagandeep\Desktop\terraform>
```

# Lab 4: Working with other providers

...



# Destroying Infra in one go :



# Destroying Infra in one go

- Terraform destroy will destroy the infrastructure in one go by using your tfstate file.

```
[root@TechLanders aws]# terraform destroy
aws_instance.myserver: Refreshing state... [id=i-0a1f8a600cb968c7c]
```

An execution plan has been generated and is shown below.

Resource actions are indicated with the following symbols:

- destroy

Terraform will perform the following actions:

```
aws_instance.myserver will be destroyed
- resource "aws_instance" "myserver" {
 - ami = "ami-06b35f67f1340a795"
```

Enter a value: yes

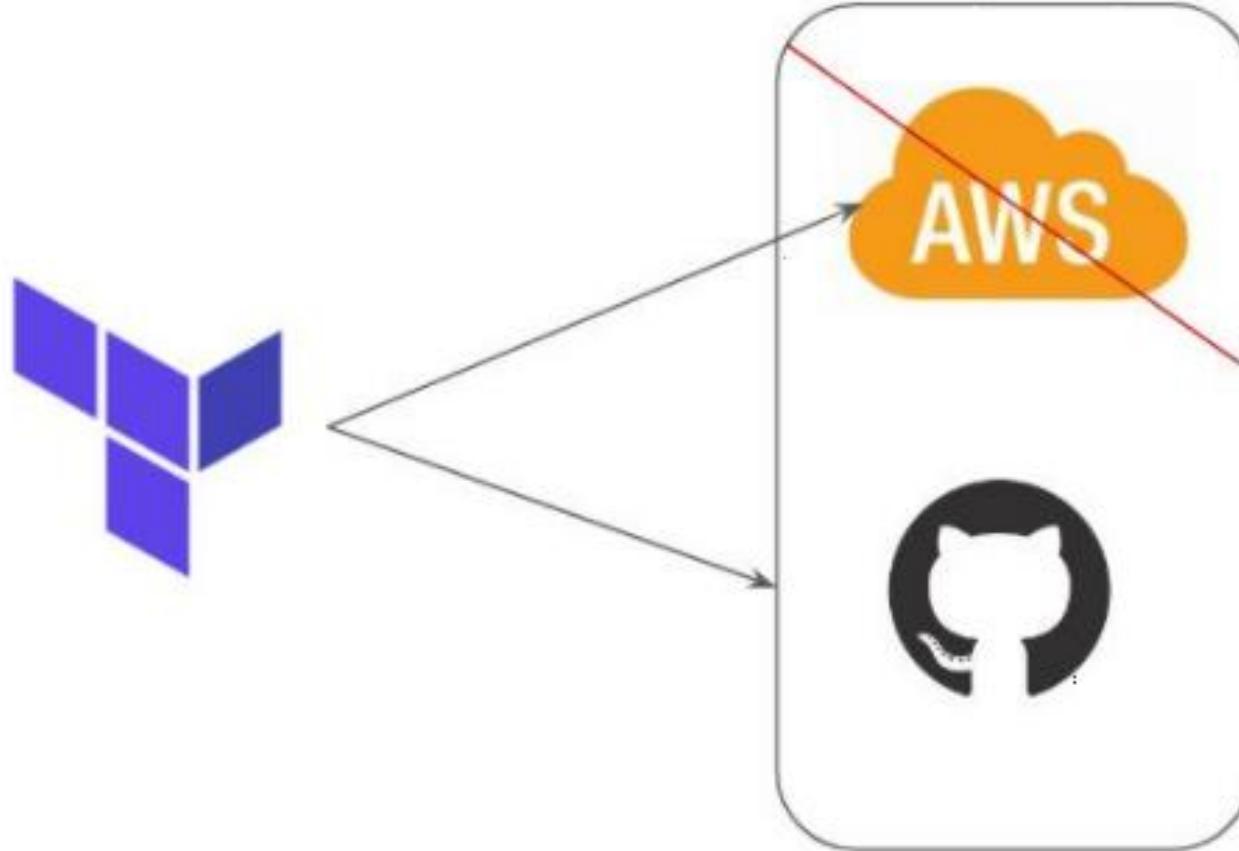
```
aws_instance.myserver: Destroying... [id=i-0a1f8a600cb968c7c]
```

```
aws_instance.myserver: Still destroying... [id=i-0a1f8a600cb968c7c, 10s elapsed]
```

```
aws_instance.myserver: Still destroying... [id=i-0a1f8a600cb968c7c, 20s elapsed]
```

```
aws_instance.myserver: Destruction complete after 29s
```

```
Destroy complete! Resources: 1 destroyed.
```



# Destroying Infra

- Terraform destroy can also delete selected resources given with -target option and can also be auto-approved with -auto-approve option. But it is always recommended to modify the configuration file instead of -target.

```
terraform destroy -target github_repository.repo
```

github\_repository.repo: Refreshing state... [id=terraform-repo]

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

- destroy

Terraform will perform the following actions:

```
github_repository.repo will be destroyed
- resource "github_repository" "repo" {
 - allow_auto_merge = false -> null
```

, which means that the result of this plan may not represent all of the changes requested by the current configuration.

The -target option is not for routine use and is provided only for exceptional situations such as recovering from errors or mistakes, or when Terraform specifically suggests to use it as part of an error message.

Note: Multiple -target options are supported as well.

# Lab 5 : Desired ,current state and last known configuration

..

Terraform tries to ensure that the deployed infrastructure is based on the desired state.

If there is a difference between the two, terraform plan presents a description of the changes necessary to achieve the desired state.



# LAB 6:CHALLENGE WITH DESIRED AND CURRENT STATE

..

Provider plugins are released separately from Terraform itself.

They have a different set of version numbers.



Version 1



Version 2

# PROVIDER VERSIONING :

- Different Version Parameters :
- version = "2.7"
- version = ">= 2.8"
- version = "~> 2.x"
- version = "<= 2.8"
- version = ">=2.10,<=2.30"

# LAB 7: PROVIDER VERSIONING

..

# Output from a run

Terraform provides output for every run and same can be used to list the resources details which are created using help of Terraform:

```
output "myawsserver-ip" {
 value = [aws_instance.myawsserver.public_ip]
}
```

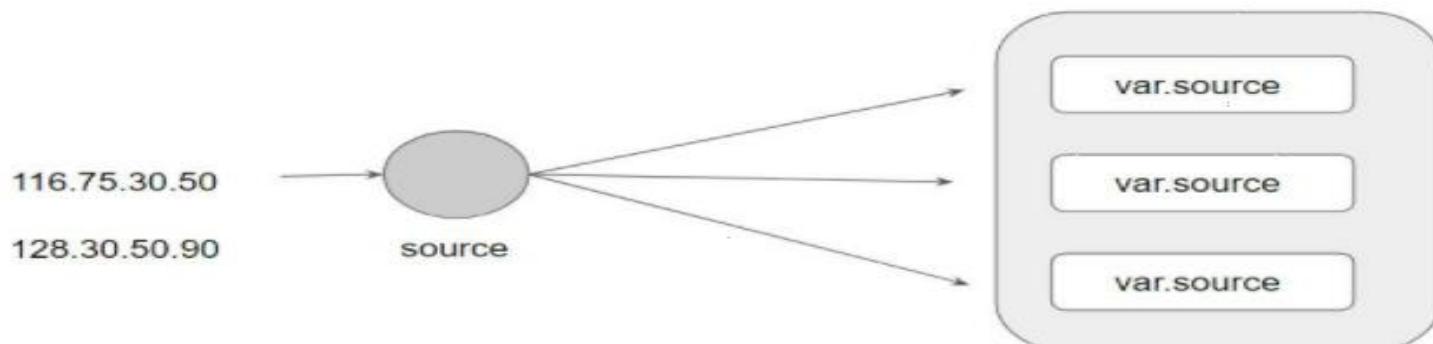
# LAB 8 : EC2 instance with output value ..

# Variables

Repeated static values can create more work in the future.



Terraform Variables allows us to centrally define the values that can be used in multiple terraform configuration blocks.



## “ □ *Lab 9: Variables*

..

# Variables

- Variables can be of different types, based on terraform versions:

- Strings

```
variable "project" {
 type = string }
```

- Numbers

```
variable "web_instance_count" {
 type = number
 default = 1 }
```

- Lists

```
variable "cidrs" { default = ["10.0.0.0/16"] }
```

- Maps

```
variable "machine_types" {
 type = map
 default = {
 dev = "f1-micro"
 test = "n1-highcpu-32"
 prod = "n1-highcpu-32"
 }
```

**“ ■ LAB 10: Different Approaches  
for Variable Assignment ..**

# Variables

- Variables can be assigned via different ways:
  - Via UI (if no default value is set in variable.tf)

\*Via variable.tf default value

\*Via .tfvars file (terraform.tfvars or custom.tfvars)

- Via command line flags:

```
terraform plan -var="instancetype=t2.small"
```

```
terraform plan -var-file="custom.tfvars"
```

# Variables Definition Precedence

Terraform loads variables in the following order, with later sources taking precedence over earlier ones:

- Environment variables
- The `terraform.tfvars` file, if present.
- The `terraform.tfvars.json` file, if present.
- Any `*.auto.tfvars` or `*.auto.tfvars.json` files, processed in lexical order of their filenames.
- Any `-var` and `-var-file` options on the command line, in the order they are provided. (This includes variables set by a Terraform Cloud workspace.)

# Lab 11: Variables

- 1) Declare AMI as variable and use same in your aws\_instance resource
- 2) Define the value of AMI( with AMIID) inside terraform.tfvars file
- 3) terraform plan
- 4) Rename terraform.tfvars with abc.tfvars
- 5) Run terraform plan again
- 6) Run terraform plan with -var-file=abc.tfvars and see the outputs
- 7) Run terraform plan with -var ami="AMI\_ID"

# Executions

```
[root@main-tf app1]# terraform plan -var-file="abc.tfvars"
```

Terraform used the selected providers to generate the following execution plan. Resource actions are indicated with the following symbols:

+ create

Terraform will perform the following actions:

```
aws_instance.myawsserver will be created
+ resource "aws_instance" "myawsserver" {
 + ami = "ami-064ff912f78e3e561"
```

# **LAB 12: Fetching data from map and list in variable ..**



# Working with change

Changes are of two types:

- Up-date In-place
- Disruptive

So always be careful with what you are adding/modifying



# Update in-place

Update in-place will ensure your existing resources intact and modify the existing resources only. Here also based on what configuration is required to be changed, server may or may-not shutdown.

Making an update in your infra so that your resource state does not get affected.

(1 change)

**“ ■ LAB 13 :UPDATE IN PLACE ..**



# Update - Disruptive

Disruptive updates require a resource to be deleted and recreated.

For example, modifying the image type for an instance will require instance to be deleted and re-created.

Making an update in your infra so that your resource state does not get affected.

(1 change)

update Disruptive:

Making an update in your infra so that old resource gets terminated/destroyed and a new resource gets created/deployed

( 1 add, 1 destroy)

## “ ■ LAB 14 :UPDATE DISRUPTIVE ..

# LAB 15 : Changes outside of terraform

Changes which occurred outside of terraform are unwanted changes and if anything which is modified outside of terraform is detected, same will be marked in state files and will be corrected at next apply.

- Create a terraform code creating a server (code on next page)
- Run terraform show command to check current required state of infrastructure.
- Modify the tag through management console .
- Run terraform plan to check the behavior of terraform against the changes
- Check the terraform show command to view state file
- Check terraform refresh command to update the state frontend
- Run terraform apply to revert the changes
- Check the terraform refresh/show command as well as console again to validate the reversion of changes.

```
[root@main-tf app1]# cat fetch.tf
provider "aws" {
 region = "us-east-2"
}

resource "aws_instance" "myec2" {
 ami = "ami-064ff912f78e3e561"
 instance_type = "t2.micro"
 tags = {
 Name = "raman-server"
 }
}
```

# Terraform in DevOps

